

Module 13: Disaster Recovery and High Availability

Topics

- How High Availability and Disaster Recovery work together
- Building highly available systems on AWS
- Best practices for high availability and disaster recovery
- Common patterns of disaster recovery on AWS

Topics

- How High Availability and Disaster Recovery work together
- Building highly available systems on AWS
- Best practices for high availability and disaster recovery
- Common patterns of disaster recovery on AWS

How Availability and Disaster Recovery

- Think of it as a spectrum
- It's part of a business continuity plan
- It's not an all or nothing proposition
- In the face of internal or external events, how to you....
 - Keep your application running 24x7
 - Make sure your data is safe
 - Get an application back up after a major disaster



How Availability

- HA is one end of the spectrum
 - Rather than recovery with defined RTO/RPO, design for continuous availability
- Goal: application never goes down
 - Eliminate single points of failure
 - “Self Healing” app recovers automatically from component failures
 - Use graceful degradation if necessary
- Traditional IT model: HA is very expensive, suitable only for absolutely mission-critical apps

How Availability Terms

- Uptime: period when system is availability for use
- Downtime (or Outage): period when system is unavailable for normal function or offline
- Graceful Degradation: system continues to be available, but at a reduced level of service or function
- Availability: percentage uptime in a given period (nines)
 - 99,999% (“five nines”) availability= no more than about 5 minutes downtime per year

Disaster Recovery on the spectrum

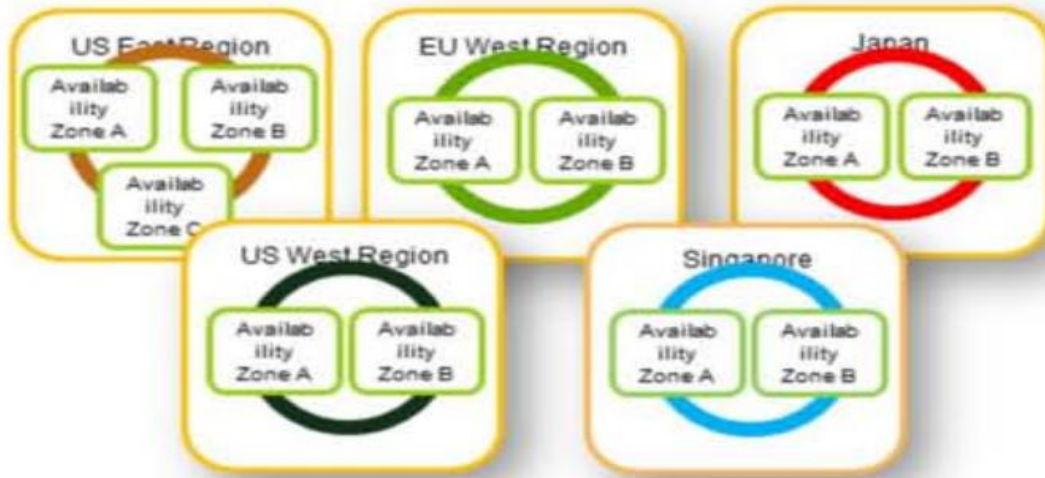
- Recover from any event
- Recover Time Object (RTO)
 - Acceptable time period within which normal operation (or degraded operation) needs to be restored after event
- Recovery Point Object (RPO)
 - Acceptable data loss measured in time
- Traditional IT model has DR in a second physical site
 - Low end DR: off-site backups
 - High end DR: hot site active-active architecture

Topics

- How High Availability and Disaster Recovery work together
- Building highly available systems on AWS
- Best practices for high availability and disaster recovery
- Common patterns of disaster recovery on AWS

AWS: Regions and Availability Zones

- Regions are completely separate clouds
- Multi network-connected Availability Zones in each Region





AWS Solutions Architect Associate

High Availability (HA)


High Availability Introduction

High Availability (HA)

The ability for a system to remain available

Think about what could cause a service to become **unavailable**:

1. When an AZ becomes unavailable eg. data-center flooded
2. When a Region becomes unavailable eg. meteor strike
3. When an web-application becomes unresponsive eg. too much traffic
4. When an instance becomes unavailable eg. instance failure
5. When a web application becomes unresponsive due to distance in geographic location



The solution we need to implement in order to ensure **High Availability**:

1. We should run our instances in Multi-AZ, an **Elastic Load Balancer** can route traffic to operational AZs.
2. We should run instances in another region. We can route traffic to another Region via **Route53**
3. We should use **Auto Scaling Groups** to increase the amount of instances to meet the demand of traffic
4. We should use **Auto Scaling Groups** to ensure a minimum amount of instances are running and have **ELB** route traffic to healthy instances
5. We should use **CloudFront** to cache static content for faster delivery in nearby regions. We can also run our instances in nearby regions and route traffic using a geolocation policy in **Route53**



AWS Solutions Architect Associate

High Availability (HA)

Scale Up and Scale Out



Scale Up vs Scale Out

When utilization increases and we are reaching capacity we can:

Scale up (Vertical Scaling)

Increasing the size of instances



- Simpler to manage.
- Lower availability (if a single instance fails service becomes unavailable)



Scale out (Horizontal Scaling)

Adding more of the same

- More complexity to manage.
- Higher availability (if a single instance fails it doesn't matter)



You will generally want to **scale out** and **then up** to balance complexity vs availability

Take advantage of multiple availability zones

- No cost difference between servers running in a single AZ versus multiple Azs
- Most services are designed with multiple AZ use in mind

Topics

- How High Availability and Disaster Recovery work together
- Building highly available systems on AWS
- Best practices for high availability and disaster recovery
- Common patterns of disaster recovery on AWS

Best Practices for HA

- Build loosely coupled systems
 - Queues, load-balance tiers
- Implement elasticity
 - Bootstrapping, load balancing, Auto Scaling, etc...
 - Instance asks: "Who am I and what is my role"?
- Use abstract machine and system representations
 - Build images from recipes, stacks from CloudFormation

Design for failure: Basic principles

- Goal: Applications should continue to function event if the underlying physical hardware fails or is removed or replaced
 - Avoid single points of failure
 - Assume everything fails, and design backwards
 - Design your recovery process

Topic

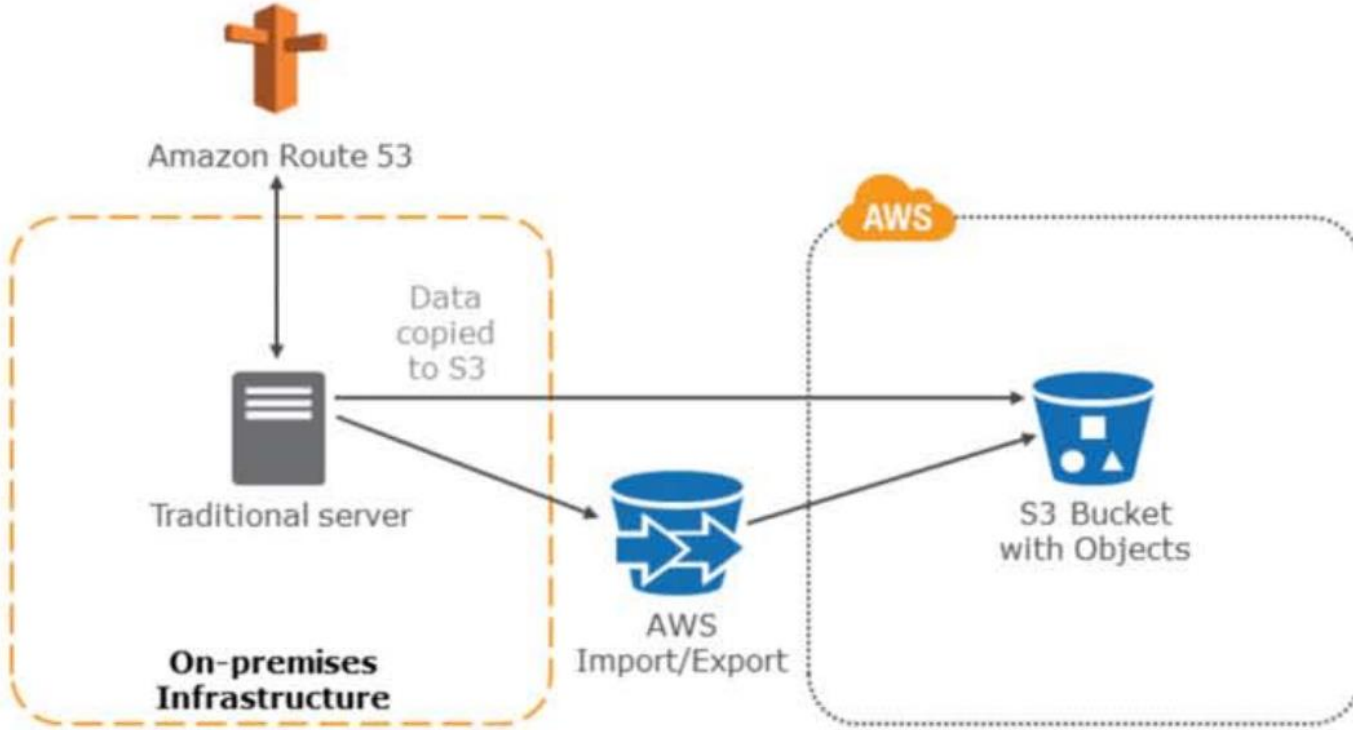
- How High Availability and Disaster Recovery work together
- Building highly available systems on AWS
- Best practices for high availability and disaster recovery
- Common patterns of disaster recovery on AWS

Common Practices of Disaster Recovery on AWS

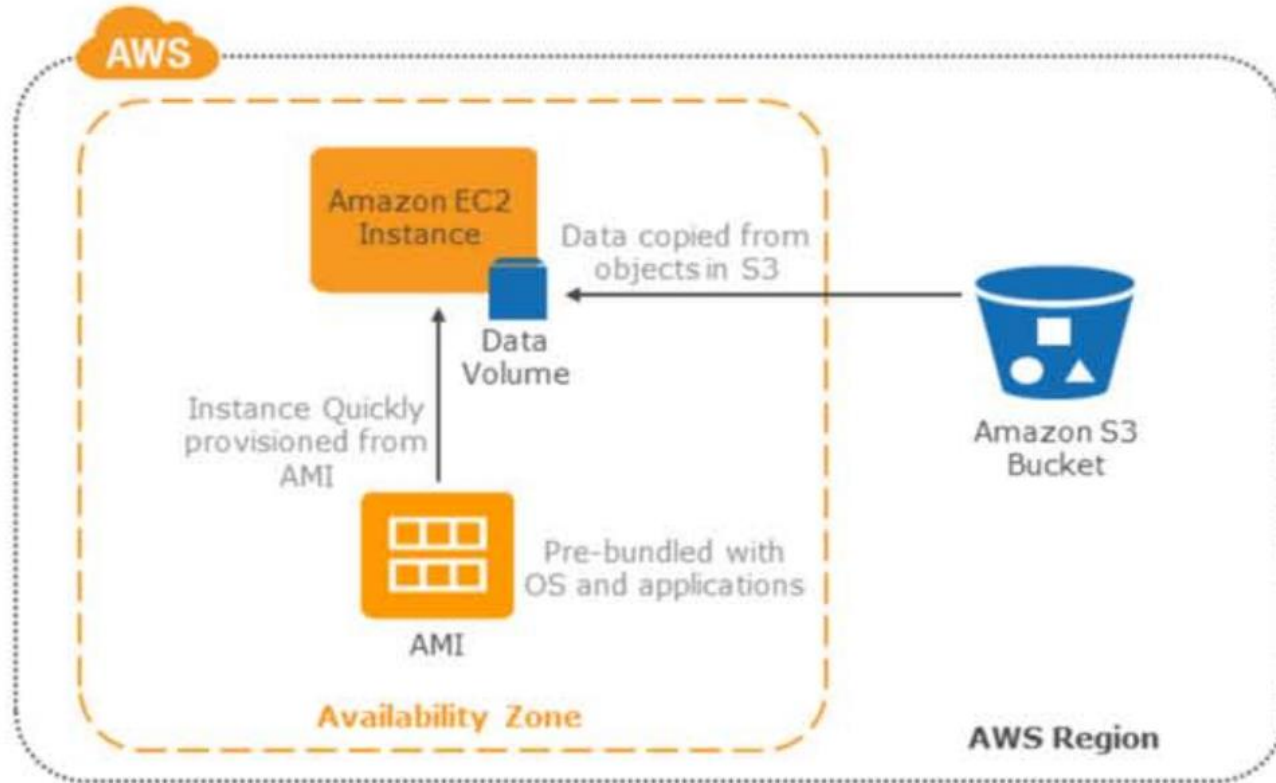
- Example Architectural Patterns (sorted by increasingly optimal RTO/RPO)

- Backup and Restore
- Pilot Light
- Fully Working Low Capacity Standby
- Multi-Site Hot Standby

Backup and Restore



Backup and Restore



Backup and Restore

- Advantages

- Simple to get started
- Extremely cost effective (mostly backup storage)

- Preparation Phase

- Take backups of current systems
- Store backups in S3
- Describe procedure to restore from backups on AWS
 - Know which AMI to use, build your own as needed
 - Know how to restore system from backups
 - Know how to switch to new system
 - Know how to configure the deployment

Common Practices of Disaster Recovery on AWS

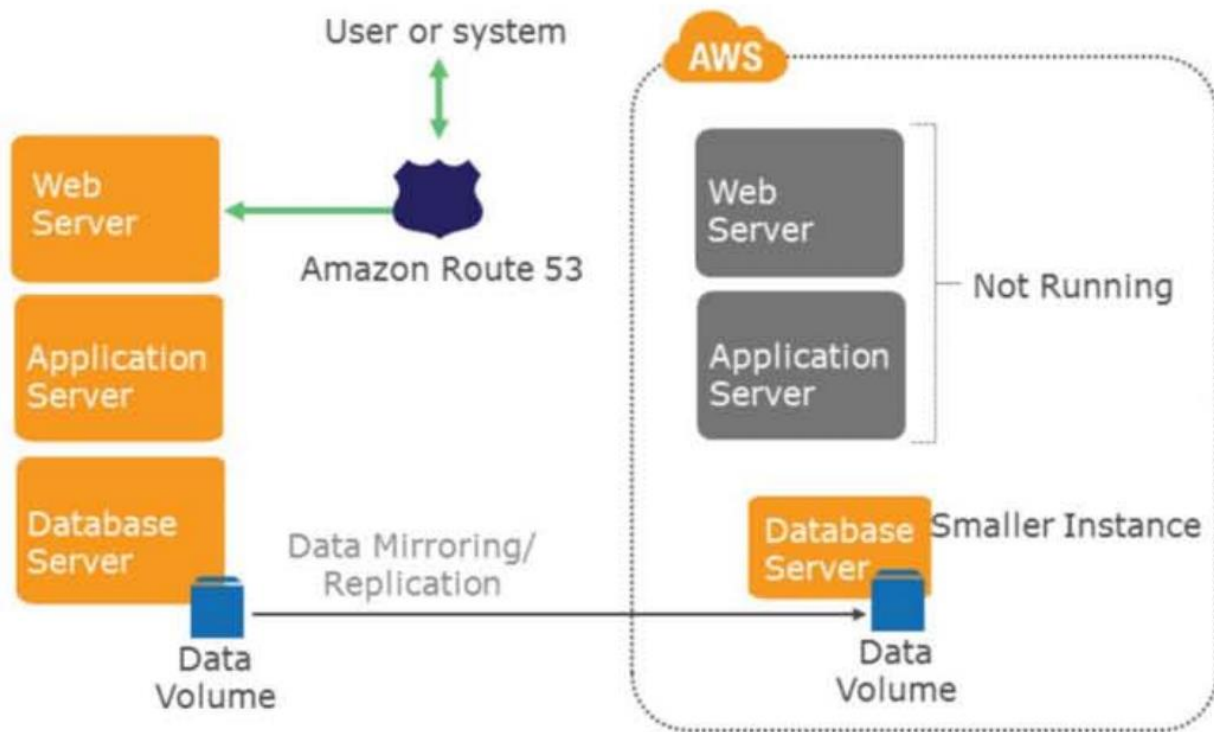
- In case of disaster

- Retrieve backups from S3
- Bring up required infrastructure
 - EC2 instance with prepared AMIs, Load Balancing, etc.
- Restore system from backup
- Switch over to the new system
 - Adjust DNS records to point to AWS


- Objectives

- RTO: as long as it takes to bring up infrastructure and restore system from backups
- RPO: time since last backup

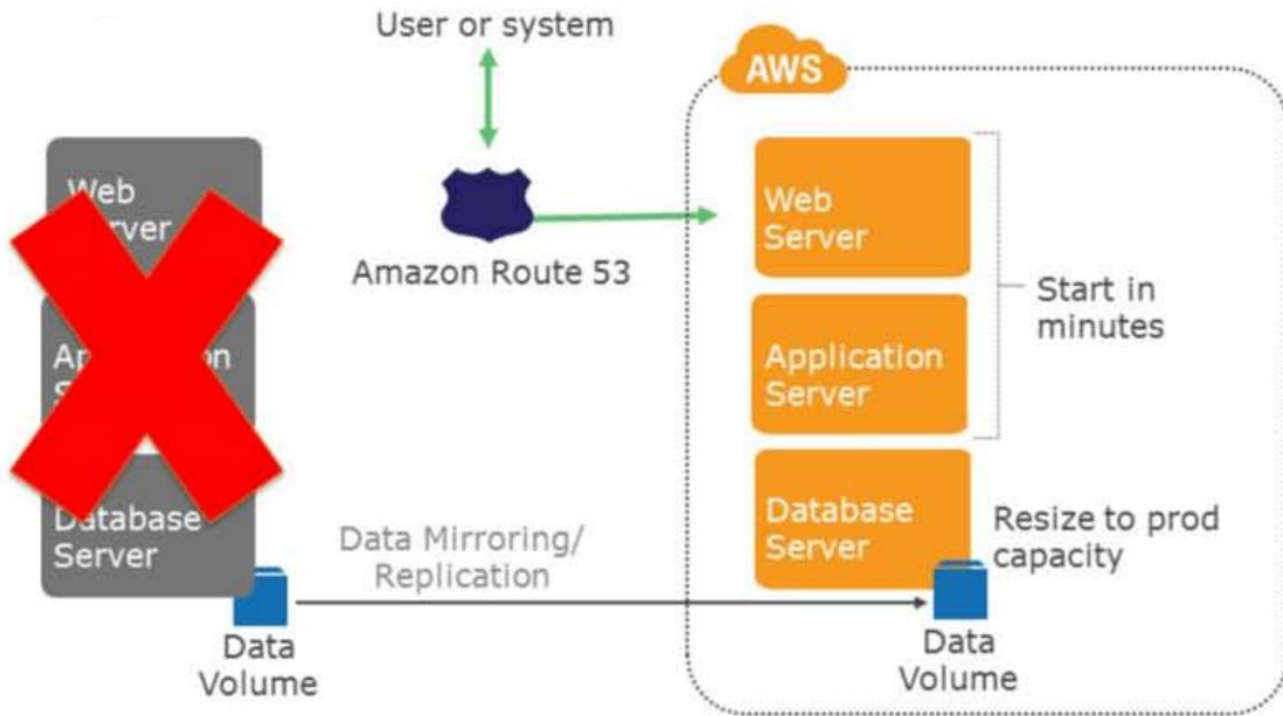
Pilot Light



Pilot Light



Pilot Light



Pilot Light

- Advantages

- Very cost effective (fewer 24/7 resources)

- Preparation Phase

- Enable replication of all critical data a AWS

- Prepare all required resources for automatic start

- AMIs, Network Settings, Load Balancing, ect..

- Reserved Instances

Pilot Light

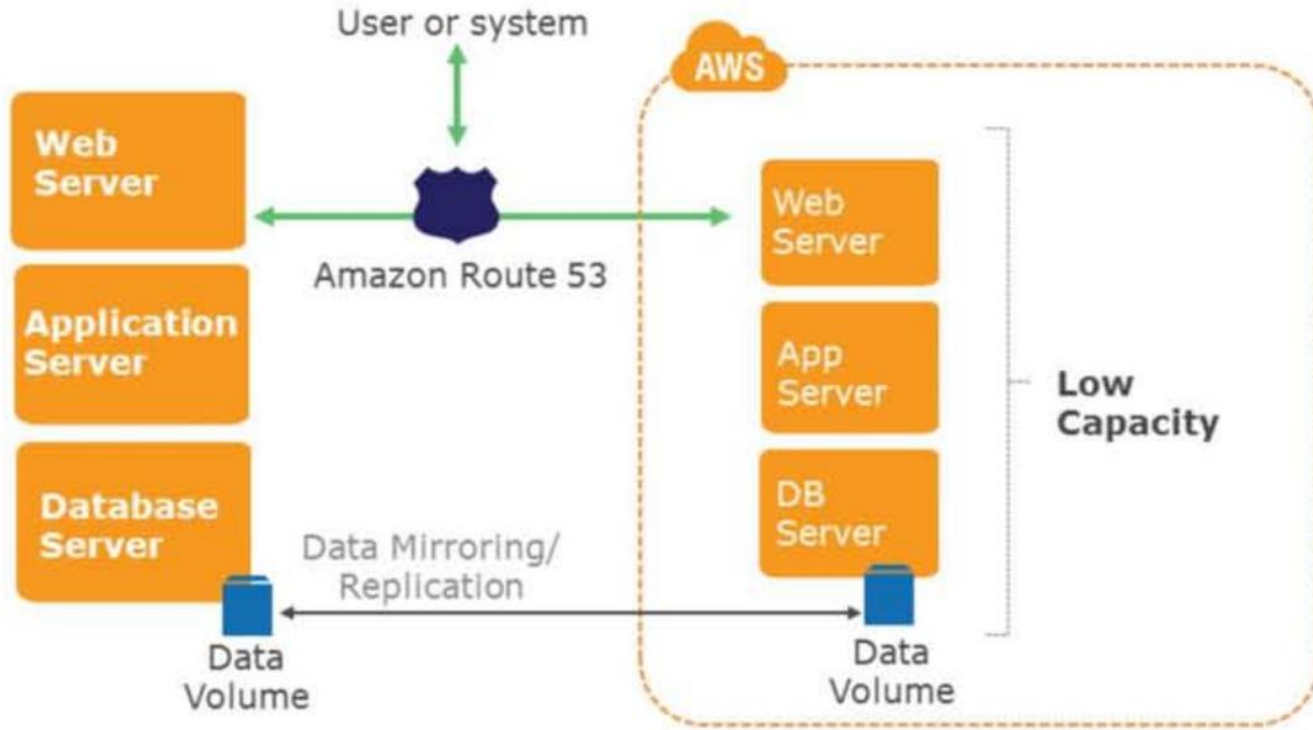
- In case of disaster

- Automatically bring up resources around the replicated core data set
- Scale the system as needed to handle current production traffic
- Switch over to the new system
 - Adjust DNS records to point to AWS

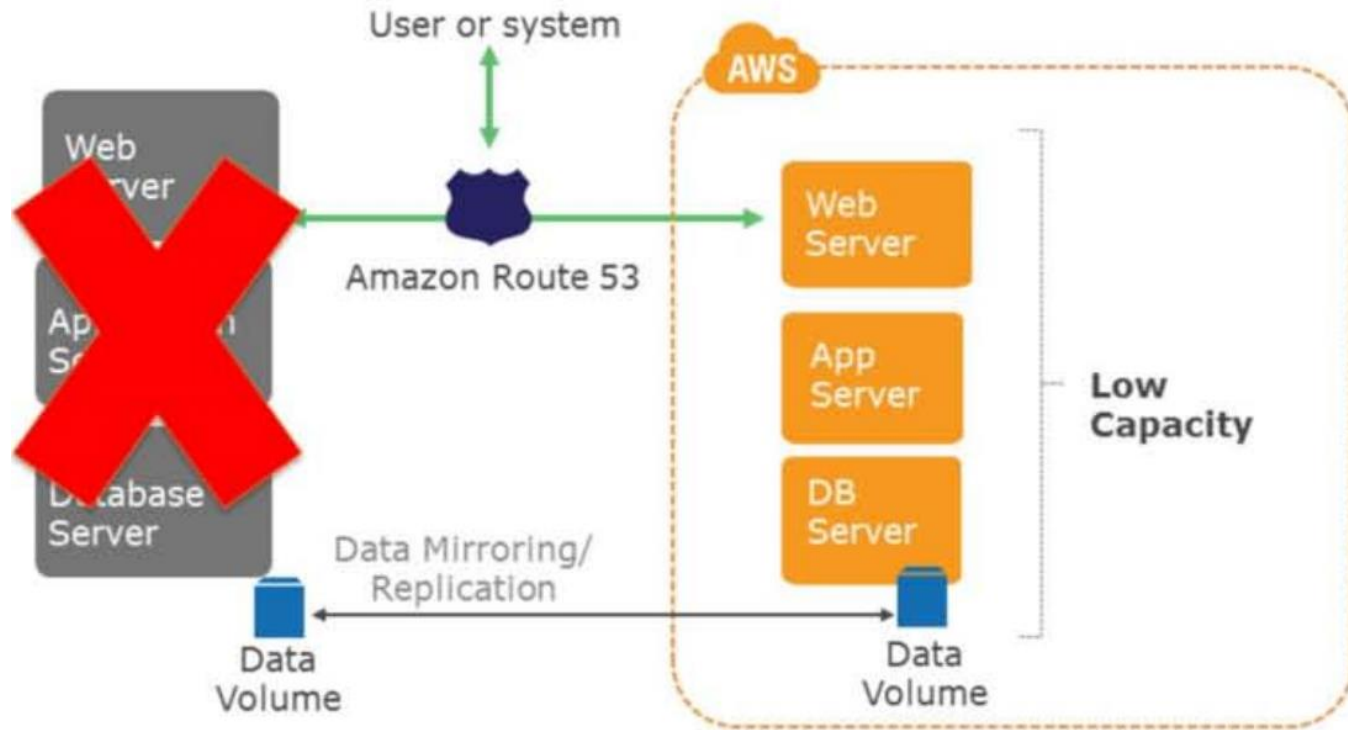
- Objectives

- RTO: as long as it takes to detect need for DR and automatically scale up replacement system
- RPO: depends on replication type

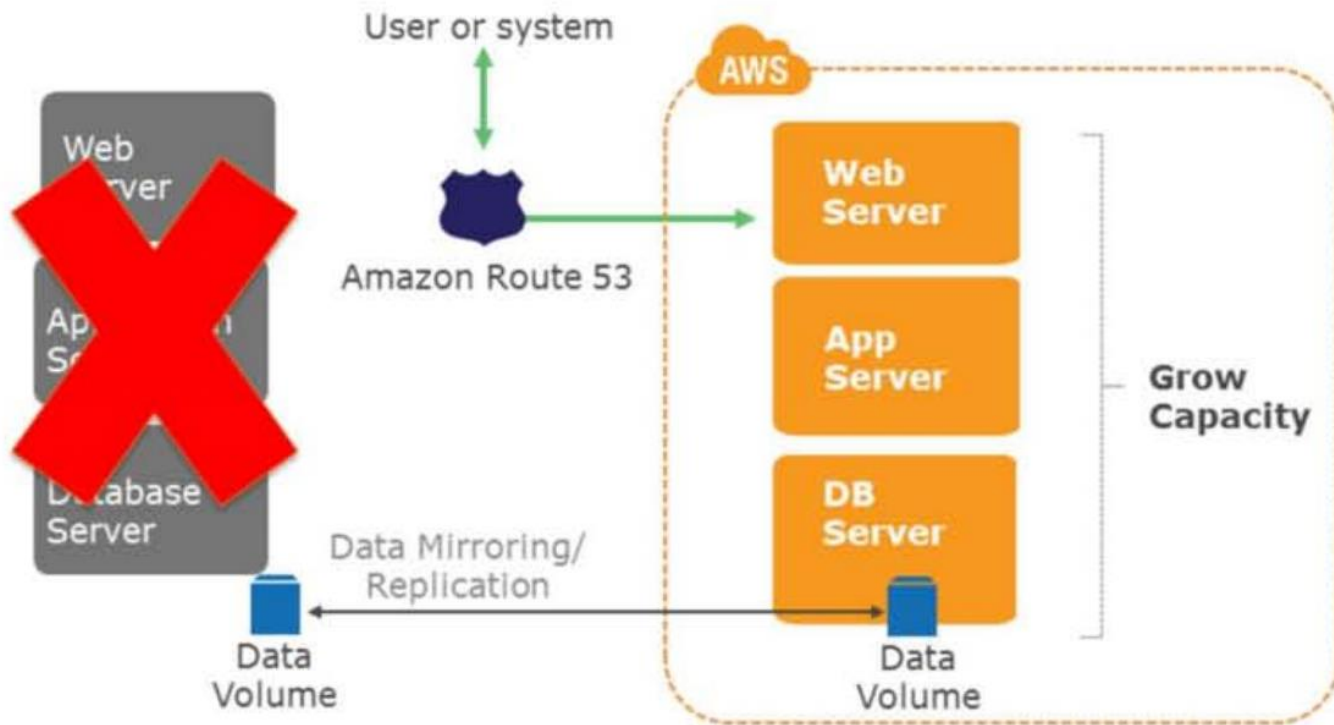
Fully-Working Low Capacity Standby



Fully-Working Low Capacity Standby



Fully-Working Low Capacity Standby



Fully-Working Low Capacity Standby

- Advantages

- Can take some production traffic at any time
- Cost savings (IT footprint smaller than full DR)

- Preparation

- Similar to Pilot Light
- All necessary components running 24/7, but not scaled for production traffic
- Best practice-continuous testing
 - “Trickle” a statistical subset of production traffic to DR site

Fully-Working Low Capacity Standby

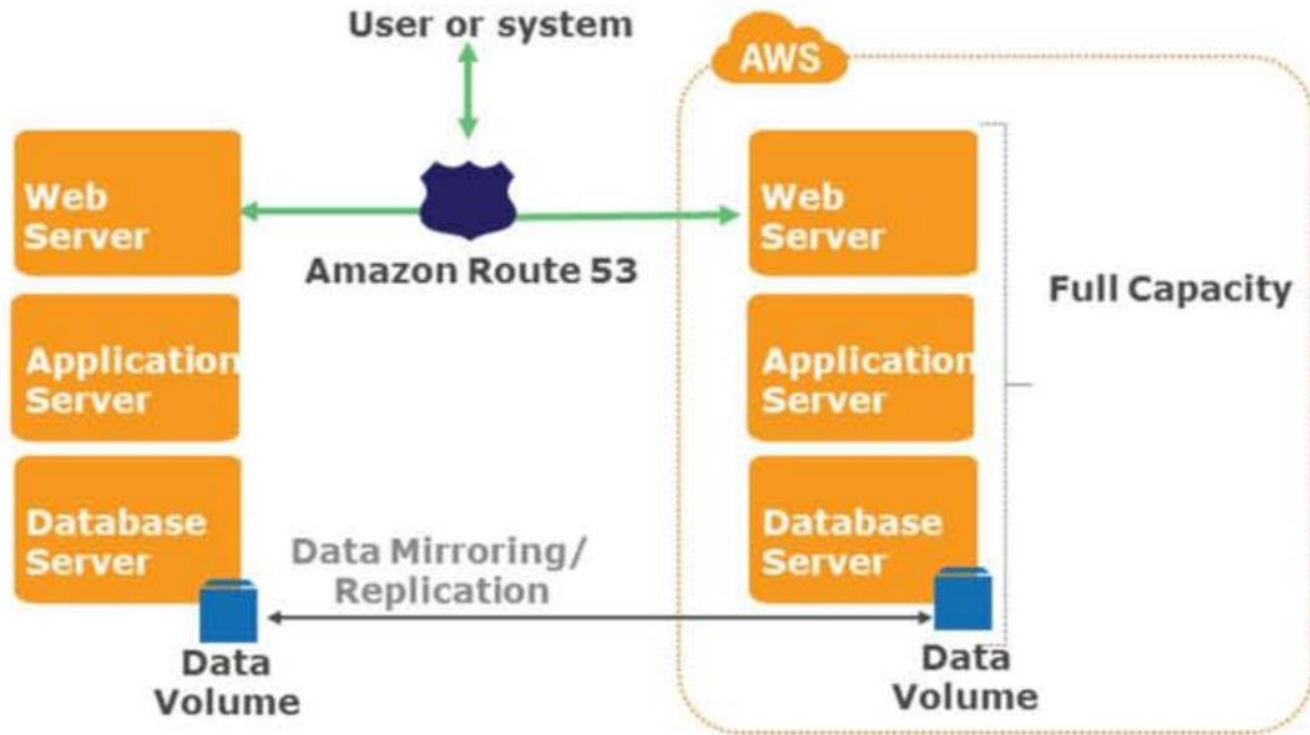
- In case of disaster

- Immediately fail over most critical production load
 - Adjust DNS records to point to AWS
- (Auto) Scale the system further to handle all production load

- Objectives

- RTO: for critical load: as long as it takes to fail over; for all other load, as long as it takes to scale further
- RPO: depends on replication type

Multi-site Active-Active



Multi-site Active-Active

- Advantages

- At any moment can take all production load

- Preparation

- Similar to Low-Capacity Standby

- Fully scaling in/out with production load

- In case of Disaster

- Immediately fail over all production load

- Adjust DNS records to point to AWS

- Objectives

- RTO: as long as it takes fail over

- RPO: depends on replication type

Hosted Desktop

- Advantages

- Replacement of workstations in case of disaster
- Pay only when used for DR

- Preparation

- Set up AMIs with appropriate working environment

- In Case of Disaster

- Launch desktop AMI and resume work

- Objectives

- RTO: as long as it takes to launch AMI and restore work environment on virtual desktop
- RPO: depends on state of AMI

Best Practices for Being Prepared

- Start simple and work your way up
 - Backups in AWS as first step
 - Incrementally improve RTO/RPO as continuous effort
- Check for any software licensing issues
- Exercise your DR Solution
 - Game Day
 - Ensure backups, snapshots, AMIs, etc.. Are working
 - Monitor your monitoring system

Conclusion – Advantages of disaster recovery with AWS

- Various building blocks available
- Fine control over cost vs. RTO/RPO tradeoffs
- Ability to scale up when needed
- Pay for what you use, and only when you use it (when an event happens)
- Ability to easily and effectively test your DR plan
- Availability of multiple locations world wide
- Hosted desktops available
- Variety of Solution Providers

Putting It Together: Incremental Improvement

- Start with existing on-premise app with traditional DR
- Gradually move DR (and thus the app) to the cloud:
 - Backup to S3 and Restore
 - “Pilot Light” on AWS for Quick Recovery
 - Fully Working Low Capacity Standby on AWS
- Migrate to primary on AWS, DR on-premise
- Add hot standby in second AWS AZ or Region
- Incrementally add HA features to primary app
- End State: Use HA techniques in Region, use DR