




# Architecting on AWS Student Guide

Version 3.1

100-ARC-31-EN-SG



## Module A: Reference Architecture

# DocStore

by aMedia Corp.

# DocStore

aMedia DocStore allows customers to upload and manage their documents

# DocStore

aMedia DocStore allows customers to upload and manage their documents

- **Upload and access content from anywhere** – browser and mobile apps
- **PDF, Office docs, images, videos** – text content extracted, indexed, and searchable

# DocStore

## Two Account Types

- Free
  - 5GB storage
  - Search on name, created/modified, and tags
    - No full-text indexing
  - All content may be securely downloaded by owner

# DocStore

## Two Account Types

- Premium
  - 20GB storage
  - Full-text indexing and search for all documents (OCR on images)
  - All content may be securely downloaded by owner

# DocStore

## Two Major Components



# DocStore

## #1: Web Application

- Allow users to upload, manage, and view all documents
- All content uploaded directly to S3

# DocStore

## #1: Web Application

- Both powered by public DocStore API
- Both must scale
- Both must be highly available

# DocStore

## #1: Web Application

- Multi-region deployment for global support
- User can log in to any region
  - Always redirected to “home” region (i.e., region where account was created)

# DocStore Requirements

## Account Types

- **Free** – 5GB cap, basic document indexing, download original files
- **Premium** – 20GB cap, full-text Indexing/search + OCR

## Web App/API

- **User authentication**
- **Session state stored off-instance**
- **Powered by API tier**
- **Highly Available**

## Content

- **Stored in S3**
- **Efficient upload/download to/from S3**
- **Static assets in one S3 bucket; distributed globally**

## Deployment

- **Oregon (us-west-2) and Sydney (ap-southeast-2)**
- **Replicate minimal information (i.e., don't replicate uploaded content)**

# DocStore

## Sample Architecture

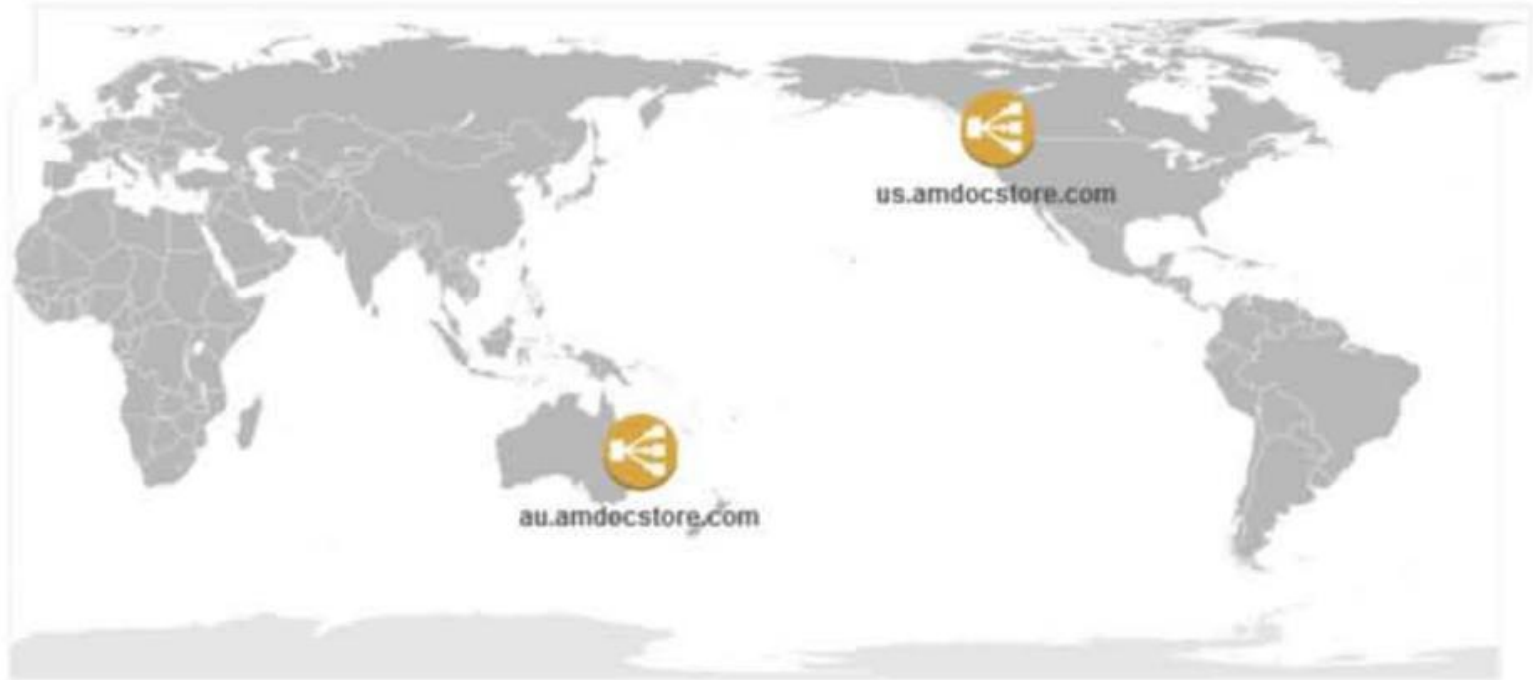
Application deployed in two regions: us-west-2 (Oregon) and ap-southeast-2 (Sydney)

**DocStore**



Each ELB has a region-specific CNAME in Route 53

**DocStore**



**Latency Based Routing** feature of **Route 53** directs customers to closest endpoint

**DocStore**





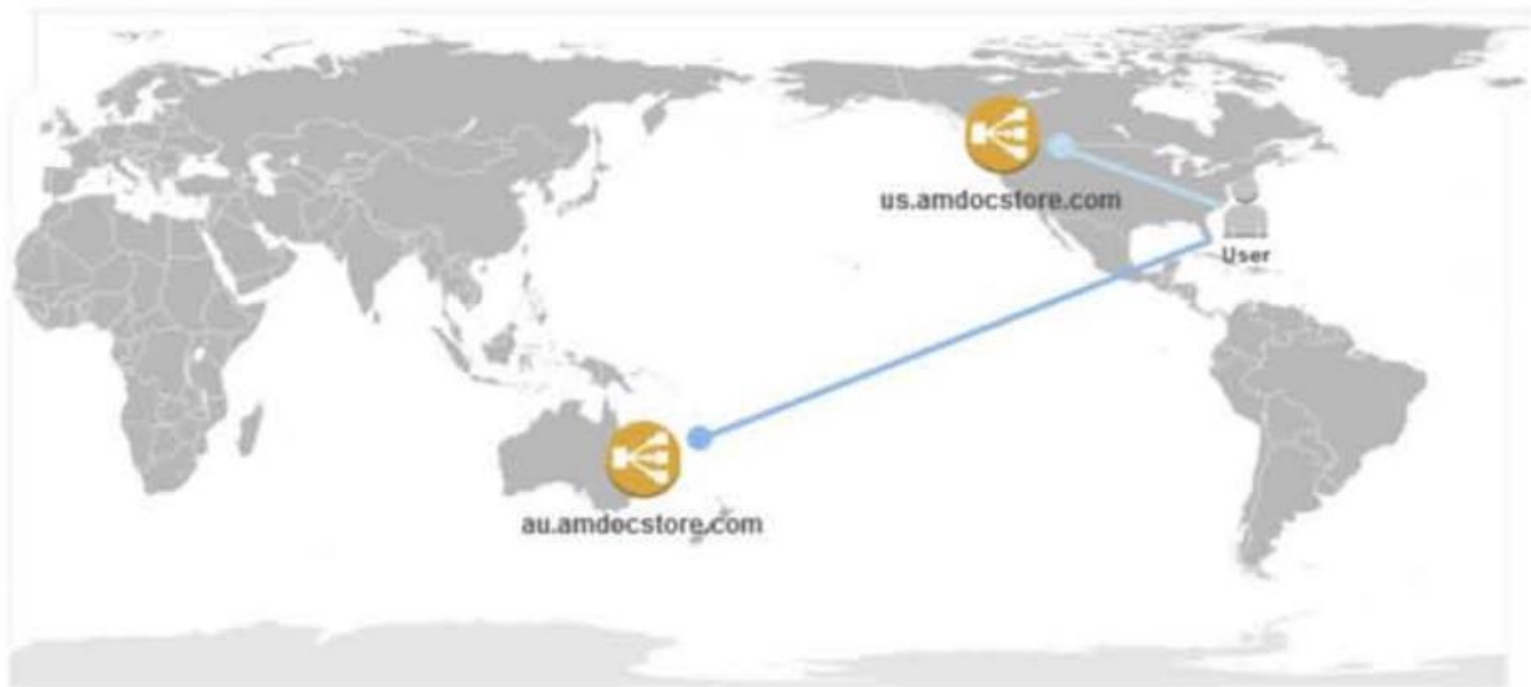
**Closer to Oregon?** That's where you go. In Perth?  
Head to the ELB in Sydney.

**DocStore**



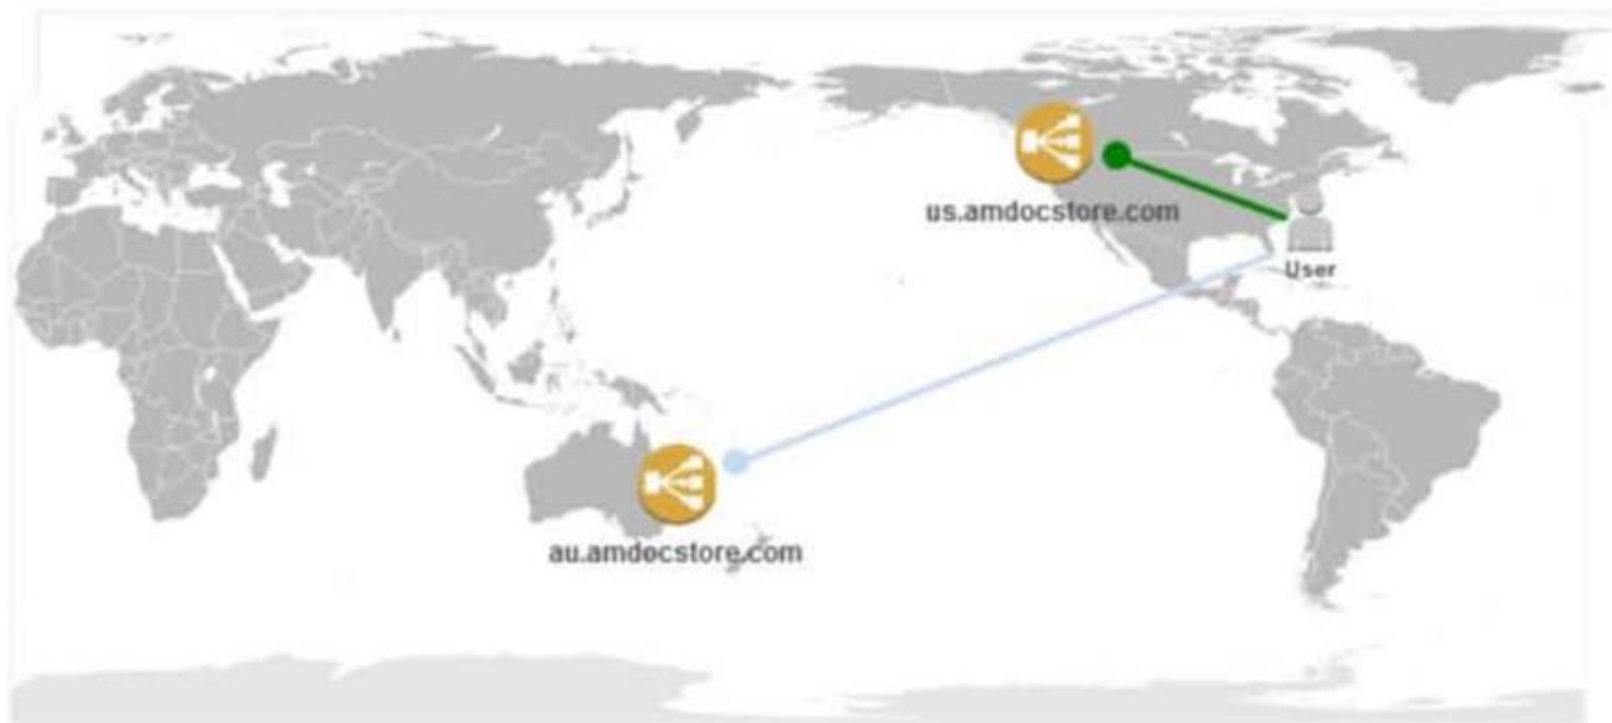
Closer to Oregon? That's where you go.

# DocStore



Closer to Oregon? That's where you go.

# DocStore



Let's focus on one region

# DocStore



us.amdocstore.com


**Route 53 LBR** directs a user to the **nearest Elastic Load Balancer** (in this case, the ELB in us-west-2)

**DocStore**




DocStore stores all **static web assets** (CSS, images, JavaScript, etc.) in **S3**

## DocStore




This bucket exists in the us-west-2 origin, but is the **canonical source for static asset requests globally**

## DocStore



CloudFront distributes static web assets in S3 to end users with low latency using a **global network of edge locations**

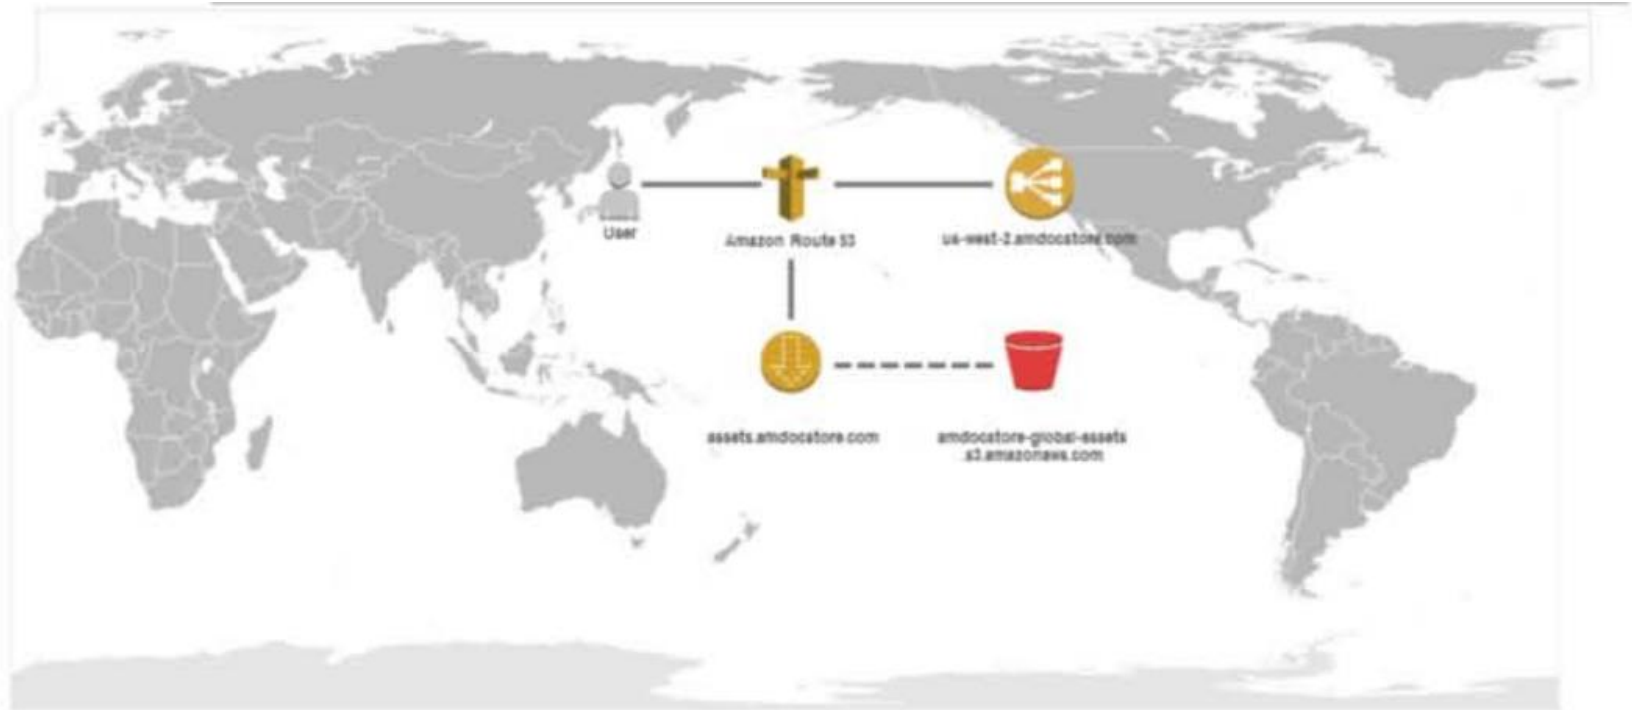
## DocStore





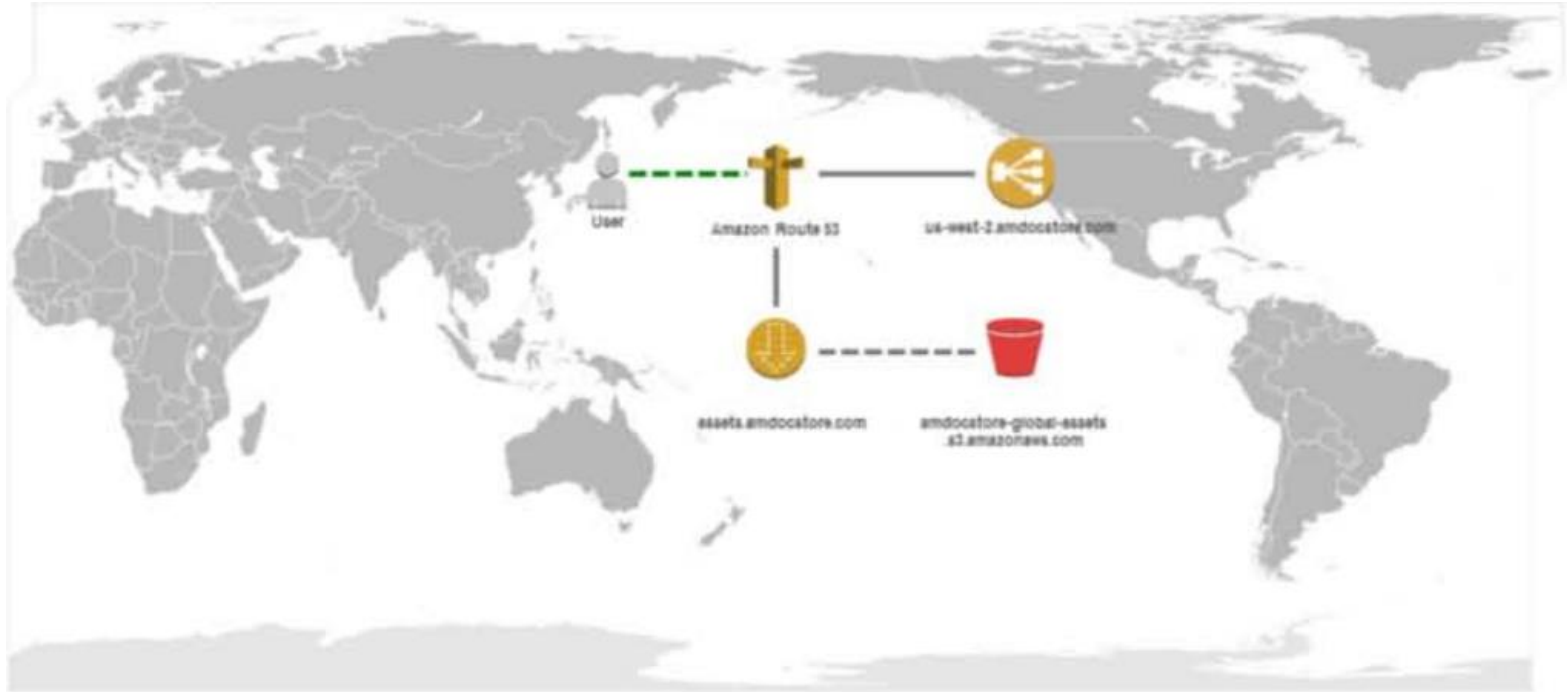
Let's see how a user in **Tokyo** would access  
<http://assets.amdocstore.com/css/style.css>

## DocStore



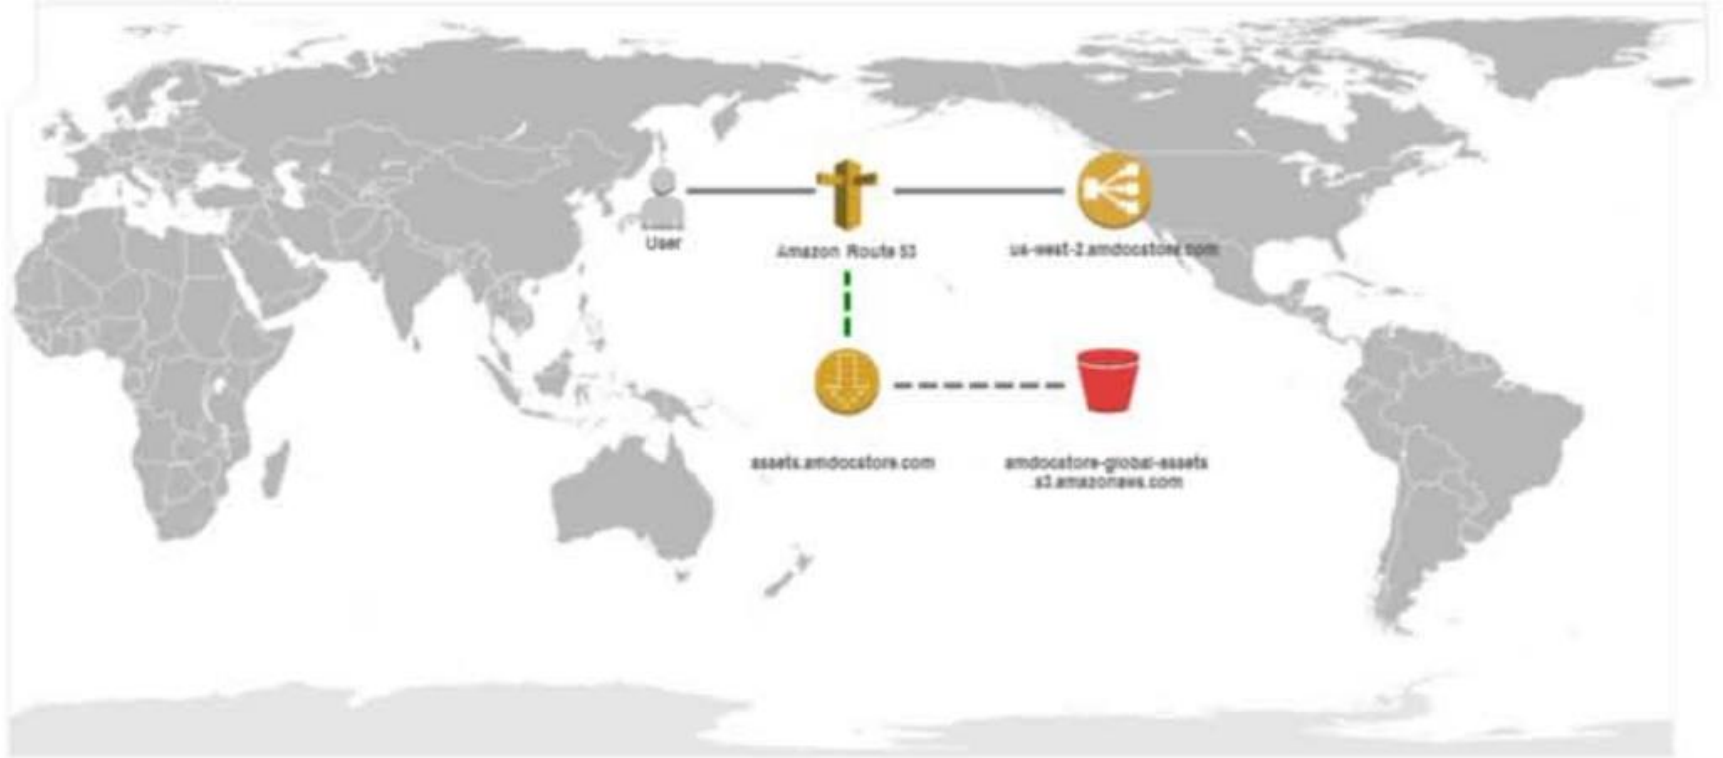
Any user request for assets.amdocstore.com will  
**automatically route to the nearest edge location**

**DocStore**



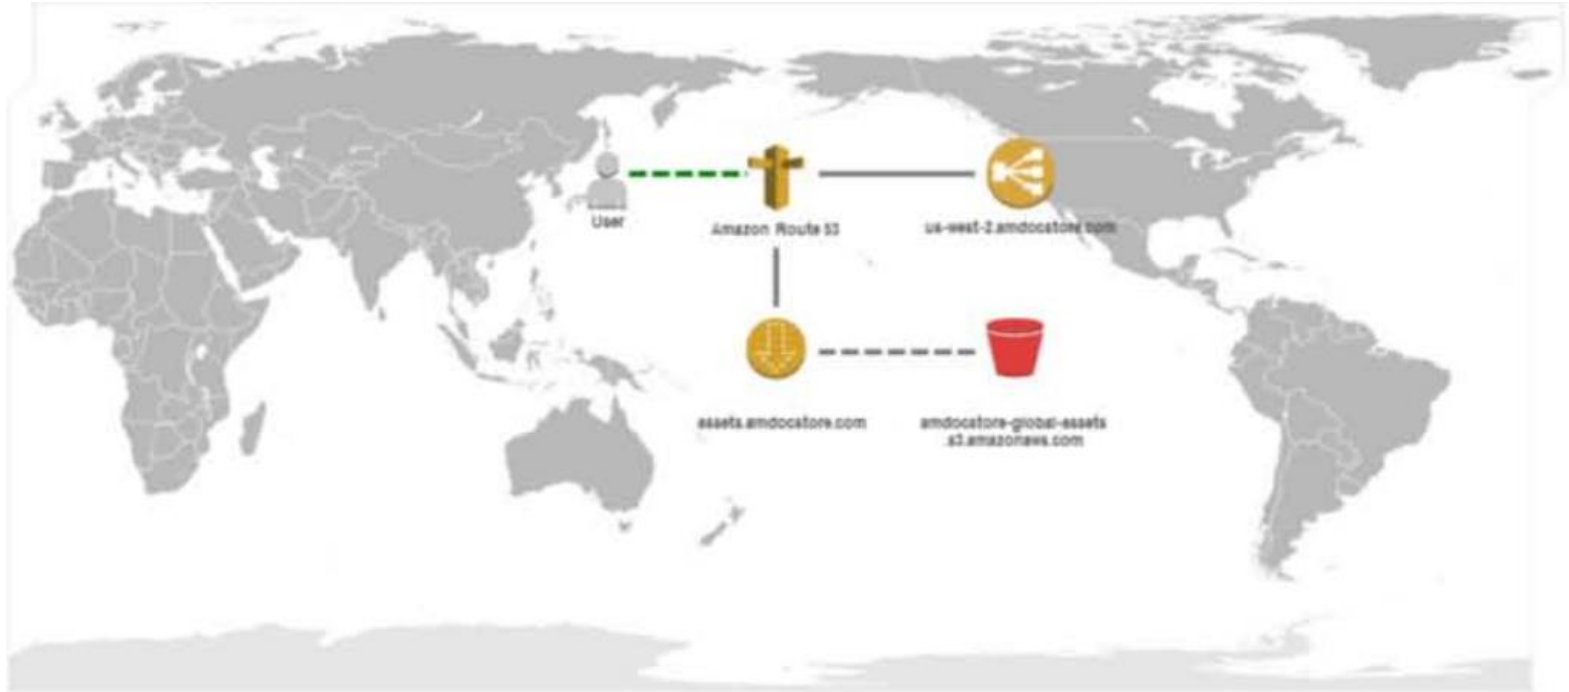
Any user request for assets.amdocstore.com will  
**automatically route to the nearest edge location**

**DocStore**




Any user request for assets.amdocstore.com will  
**automatically route to the nearest edge location**

**DocStore**



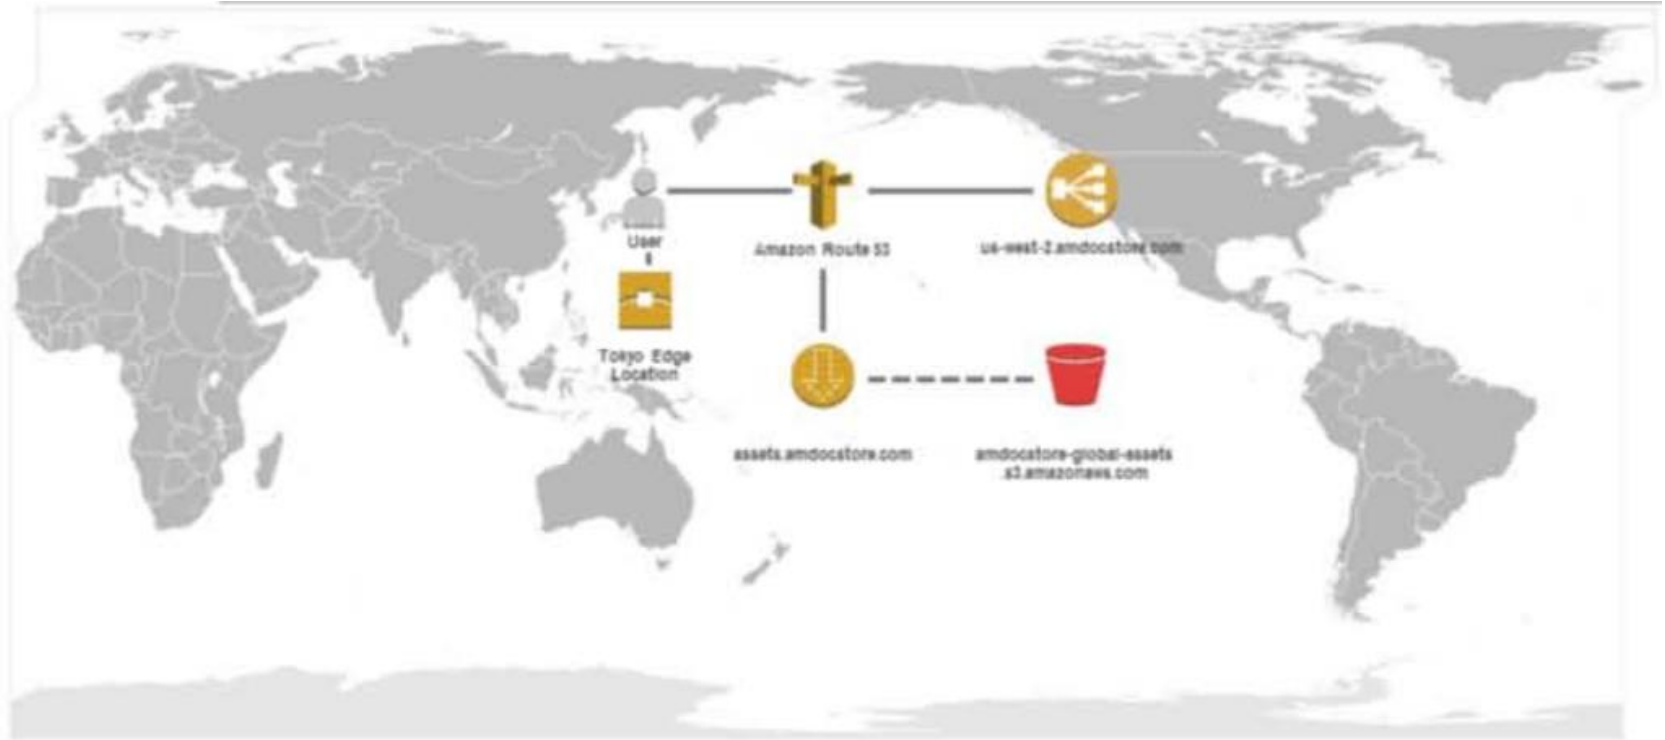
Any user request for assets.amdocstore.com will  
**automatically route to the nearest edge location**

**DocStore**



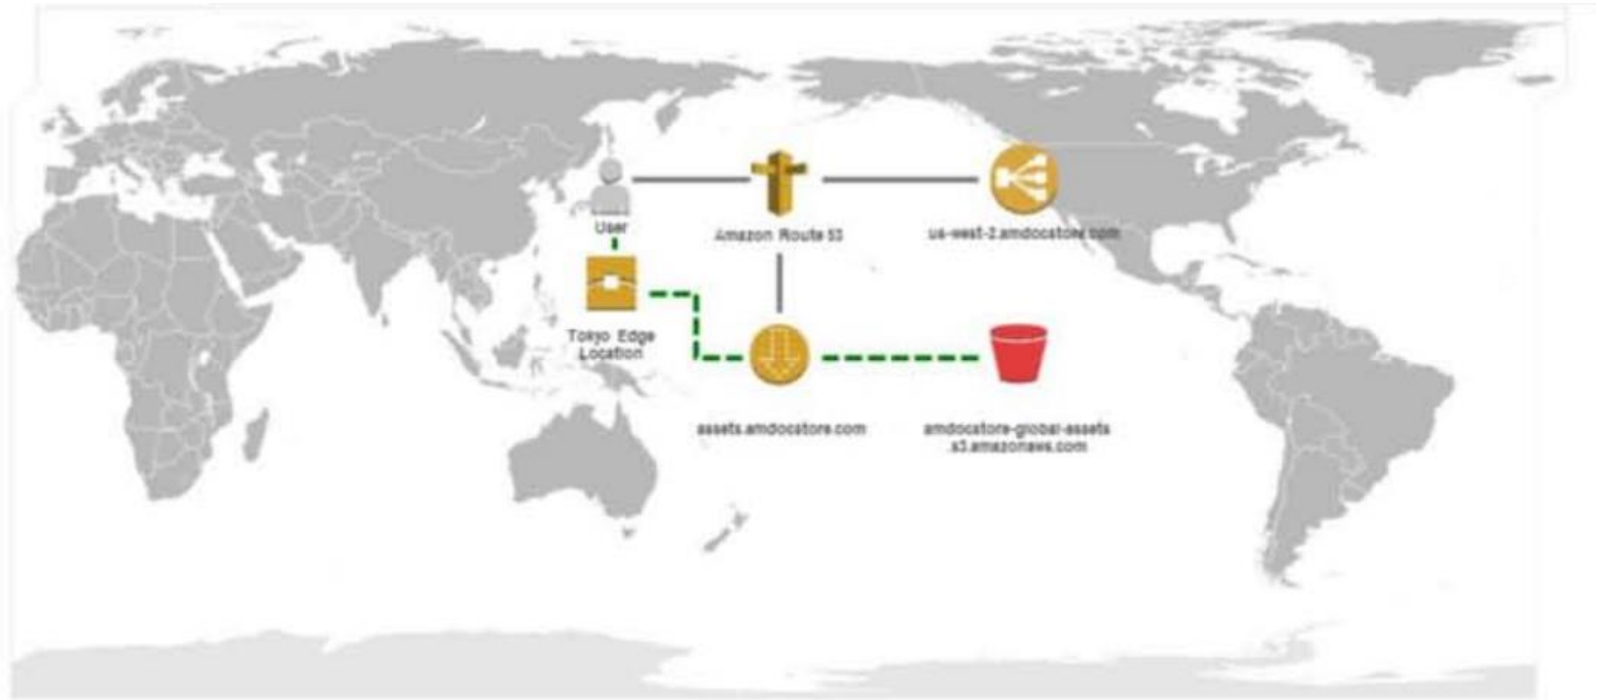
The requested **file will be served from the Tokyo edge** if it is in the cache

# DocStore



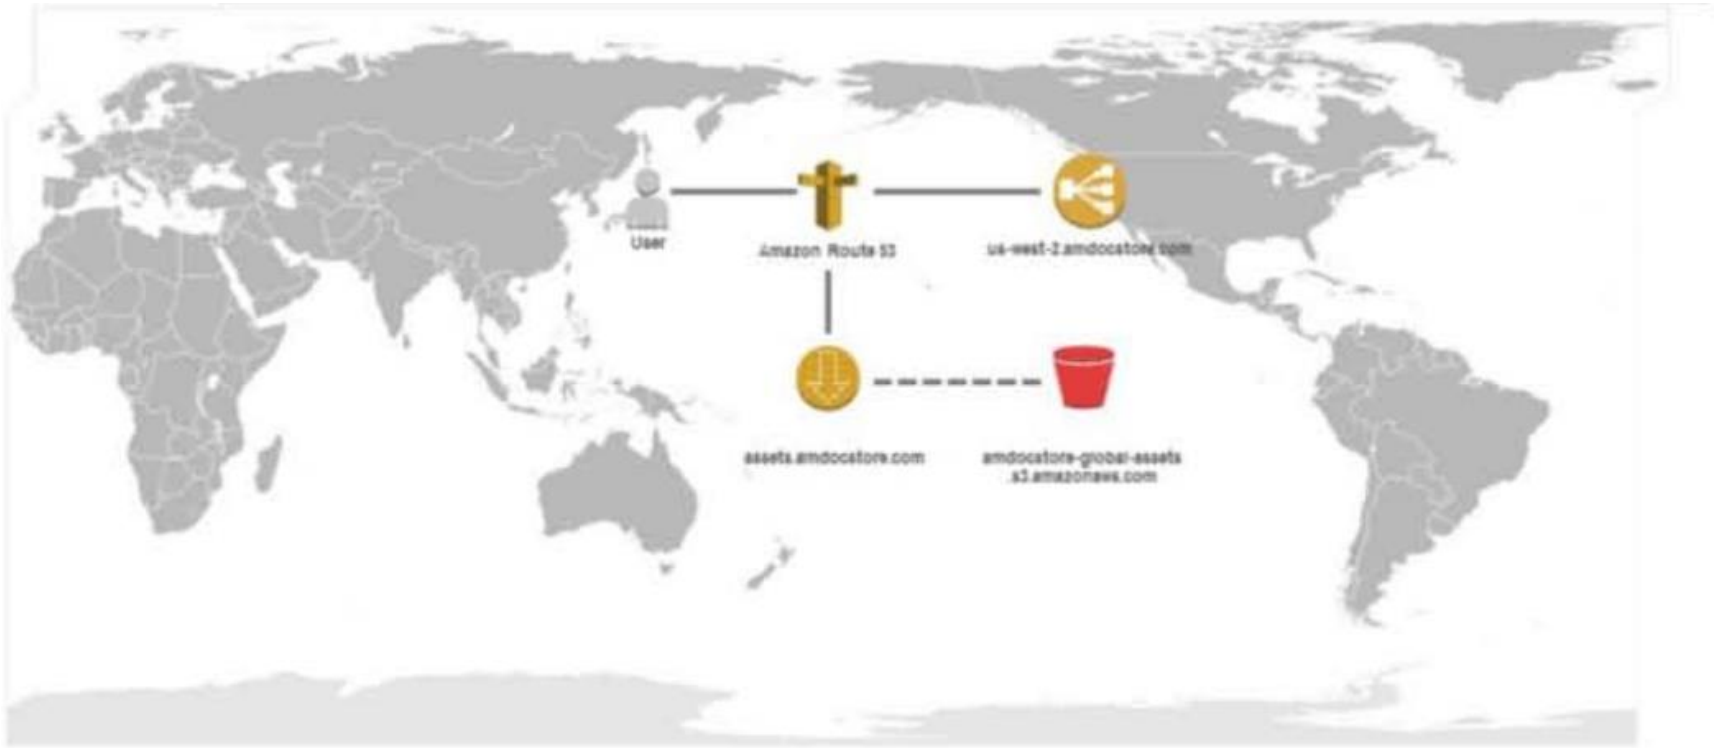
The requested file will be served from the Tokyo edge if it is in the cache, or **retrieved from the origin and cached at the edge** for future requests

## DocStore



Let's go back to our application deployment in **us-west-2**

## DocStore





Let's go back to our application deployment in **us-west-2**

## DocStore



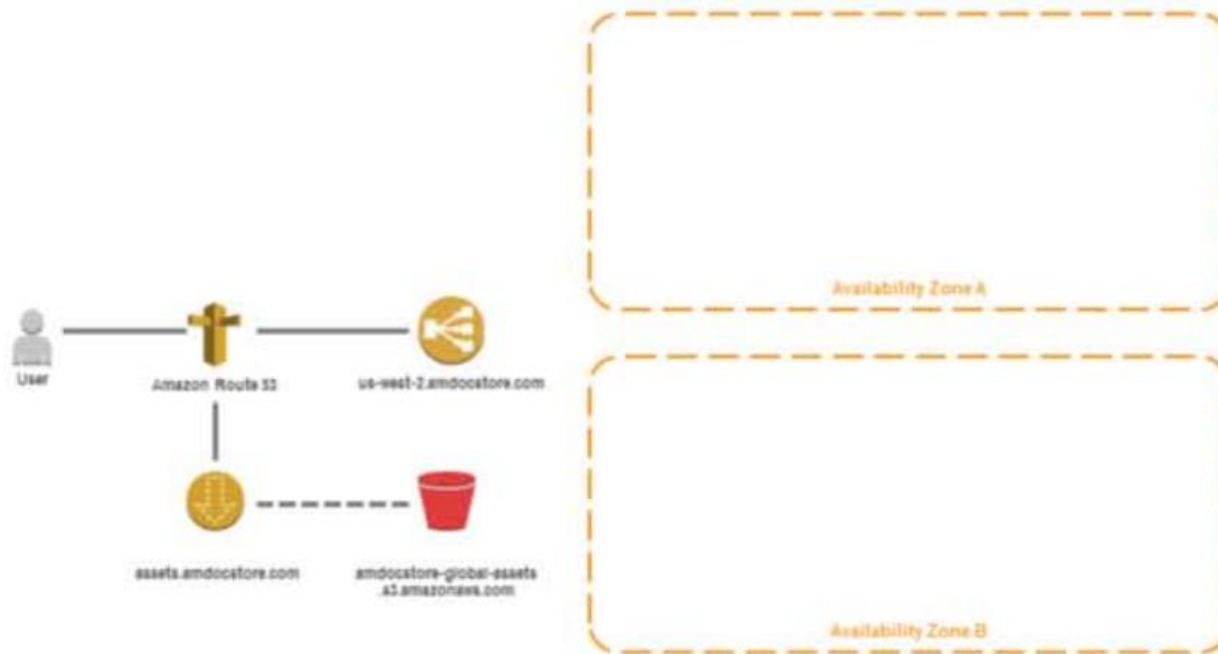
Let's go back to our application deployment in **us-west-2** and **focus on the browser component of DocStore**

## DocStore



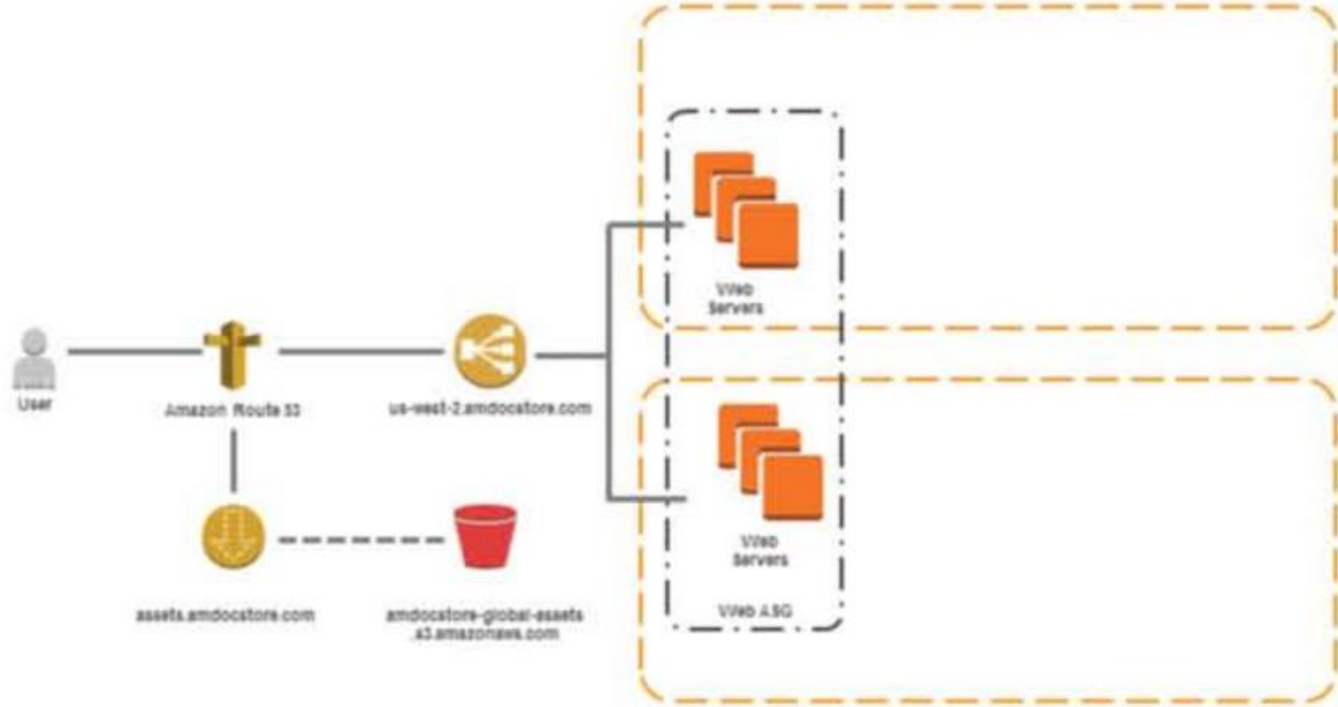
**Web servers** will be deployed across multiple Availability Zones using Auto Scaling

**DocStore**



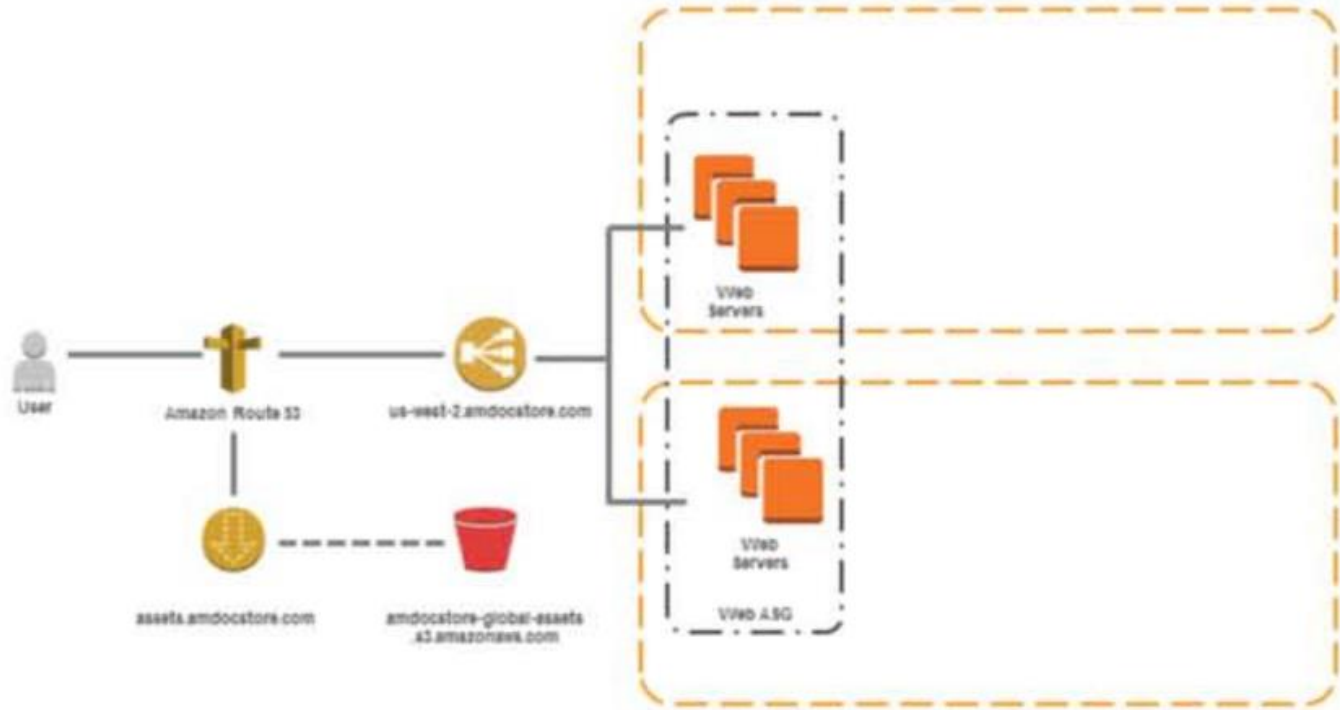
**Web servers** will be deployed across multiple Availability Zones using Auto Scaling

# DocStore



**Web servers** host the DocStore Interface that users will access with their web browsers

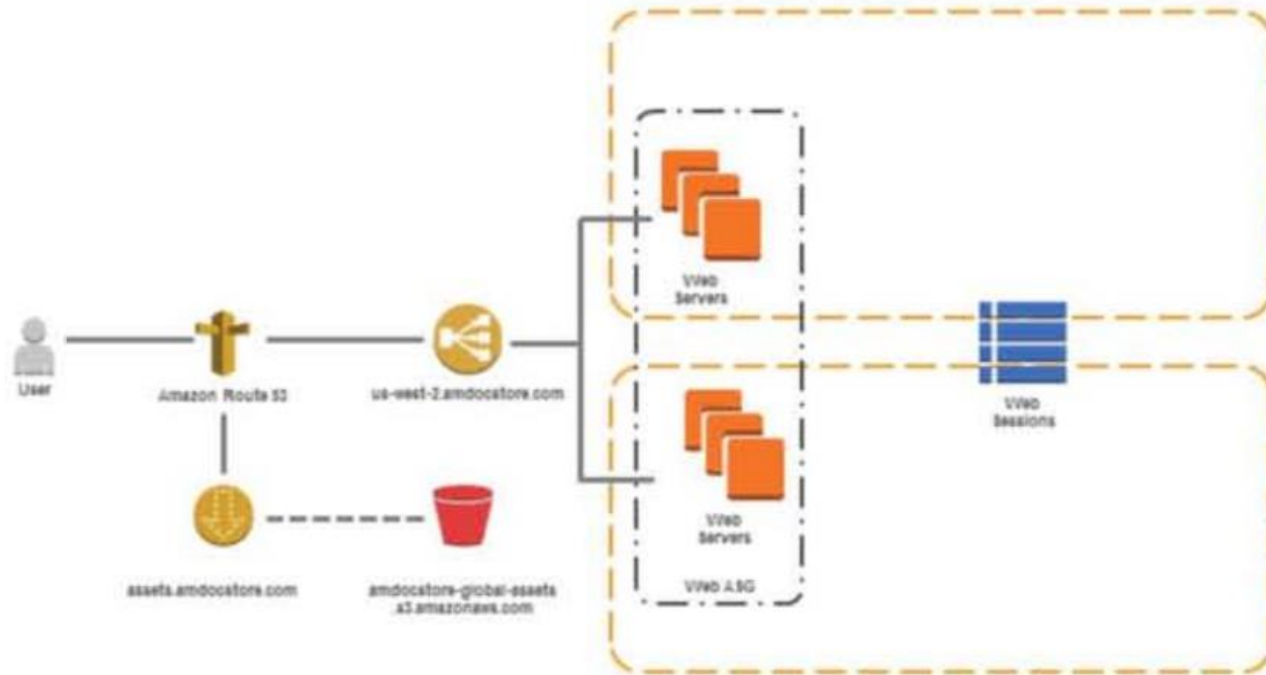
# DocStore



**Session state** is stored out of the web tier in a DynamoDB table. DynamoDB is a regional service, meaning the data is automatically distributed across AZs



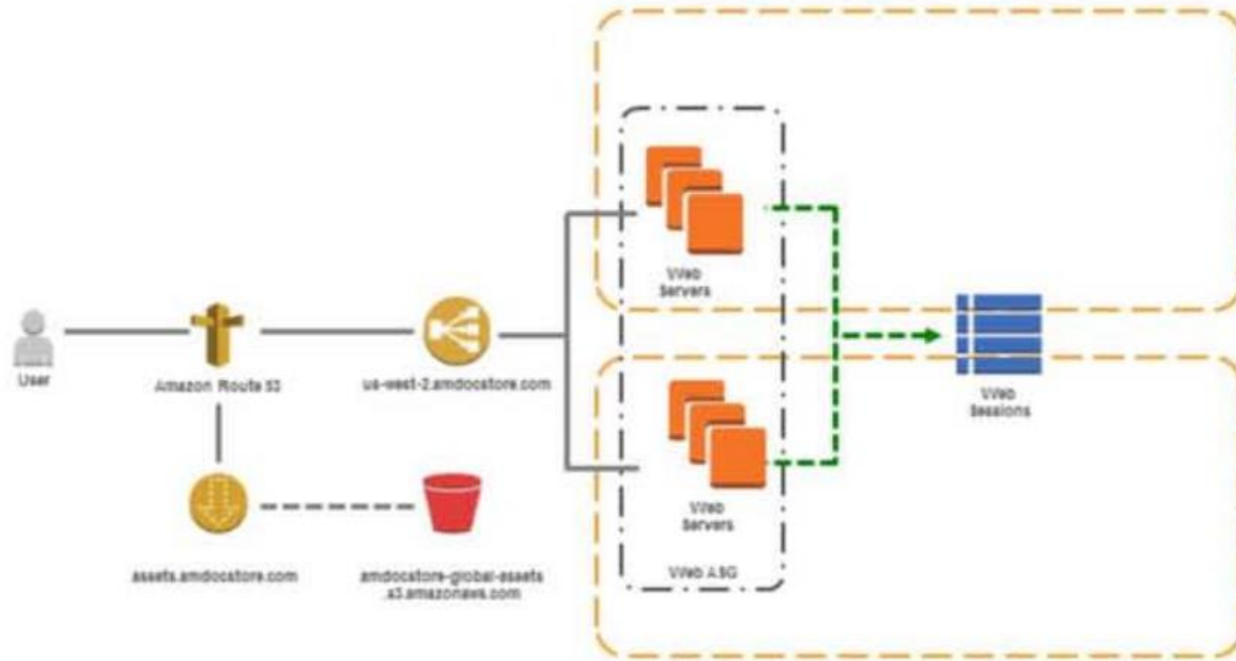
# DocStore



**Session state** is stored out of the web tier in a DynamoDB table. DynamoDB is a regional service, meaning the data is automatically distributed across AZs




# DocStore



Core functionality is provided by the **DocStore API**.  
Web servers consume this API

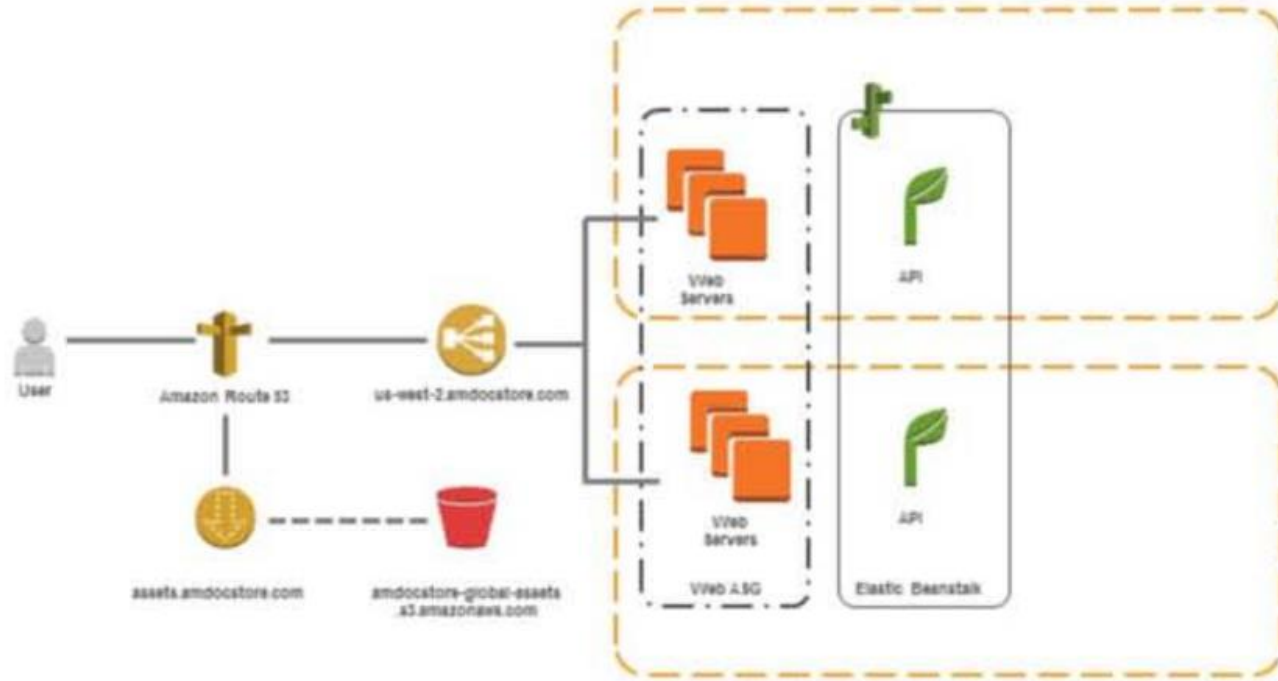
## DocStore





The DocStore API will be deployed as an application in Elastic Beanstalk container

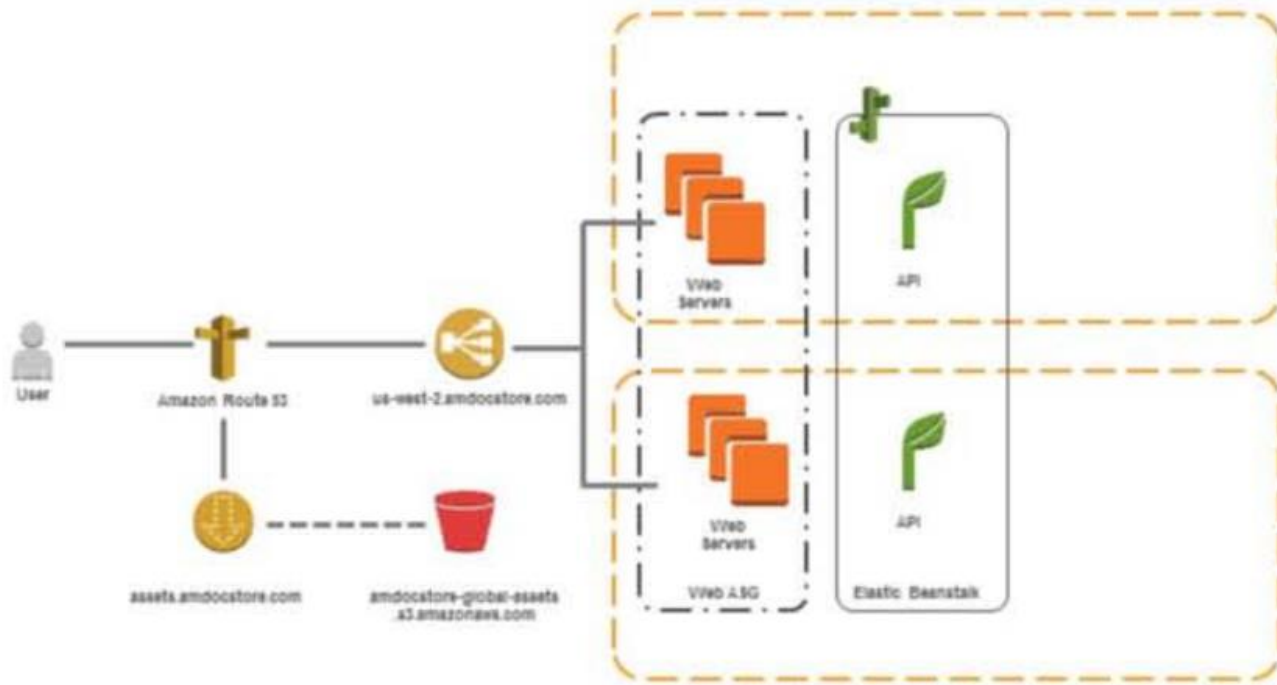
## DocStore





**Elastic Beanstalk** provides Multi-AZ deployment,  
Auto Scaling, and load balancing

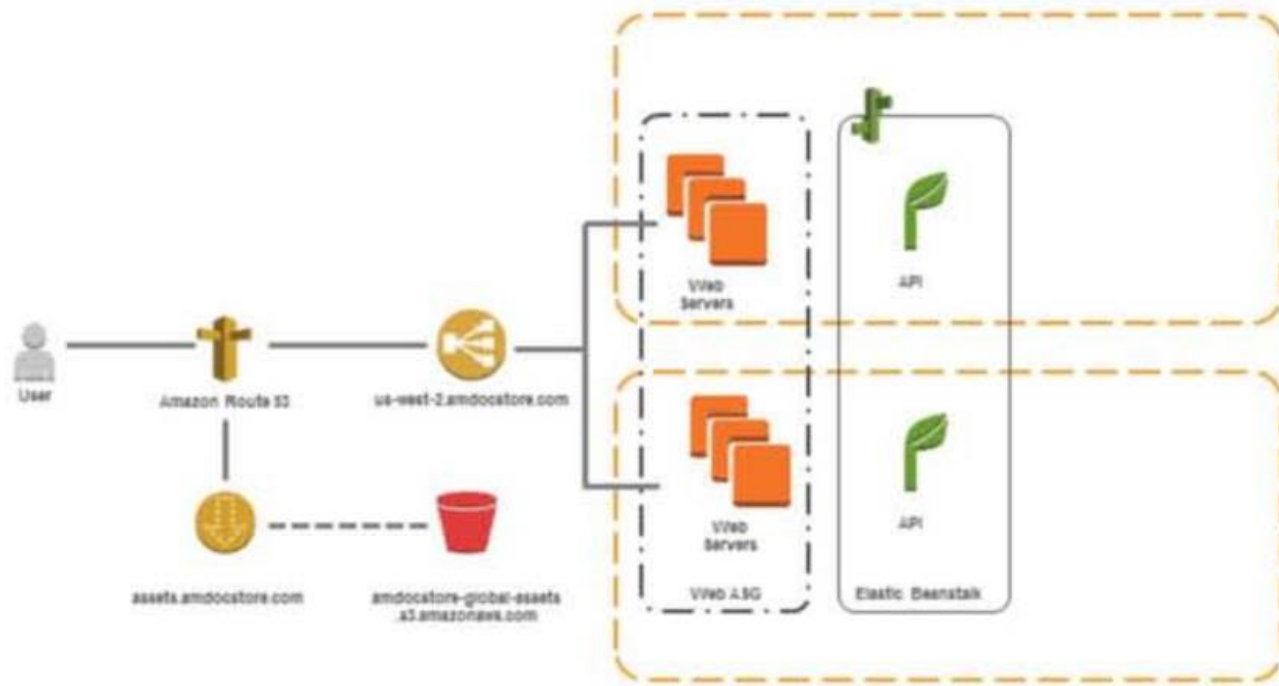
# DocStore





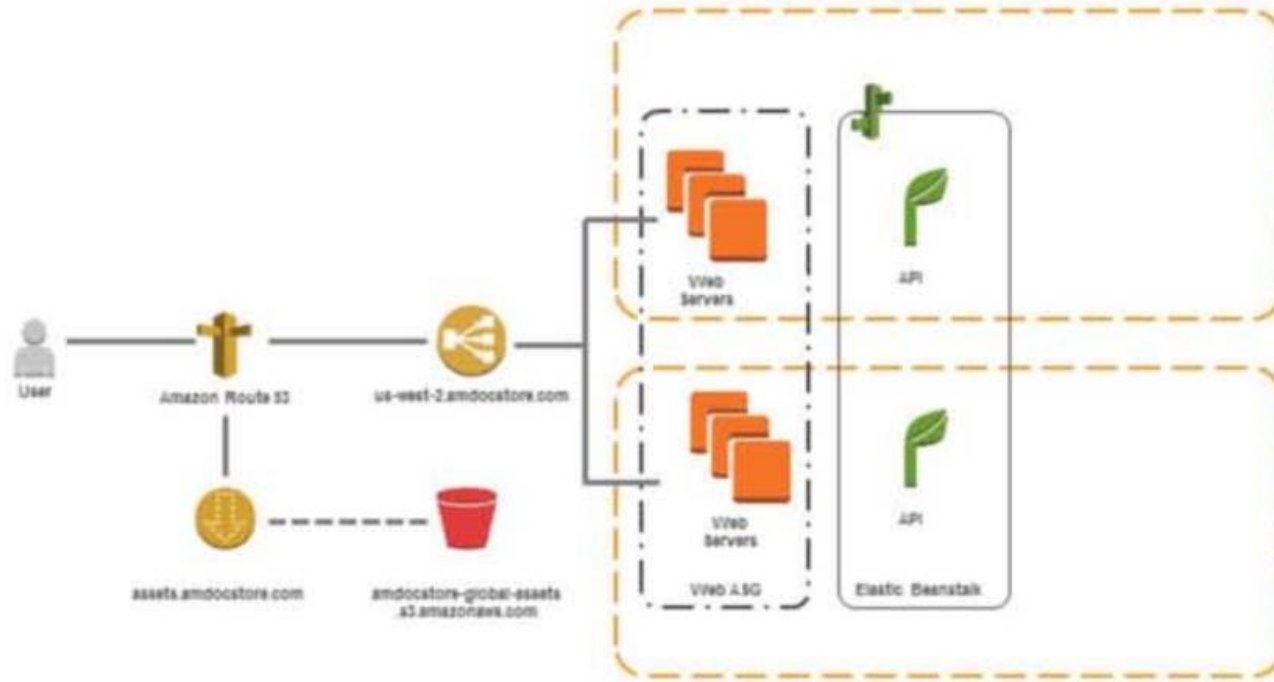
**Elastic Beanstalk** supports Java, .NET, Python, PHP, Ruby, and node.js applications

# DocStore



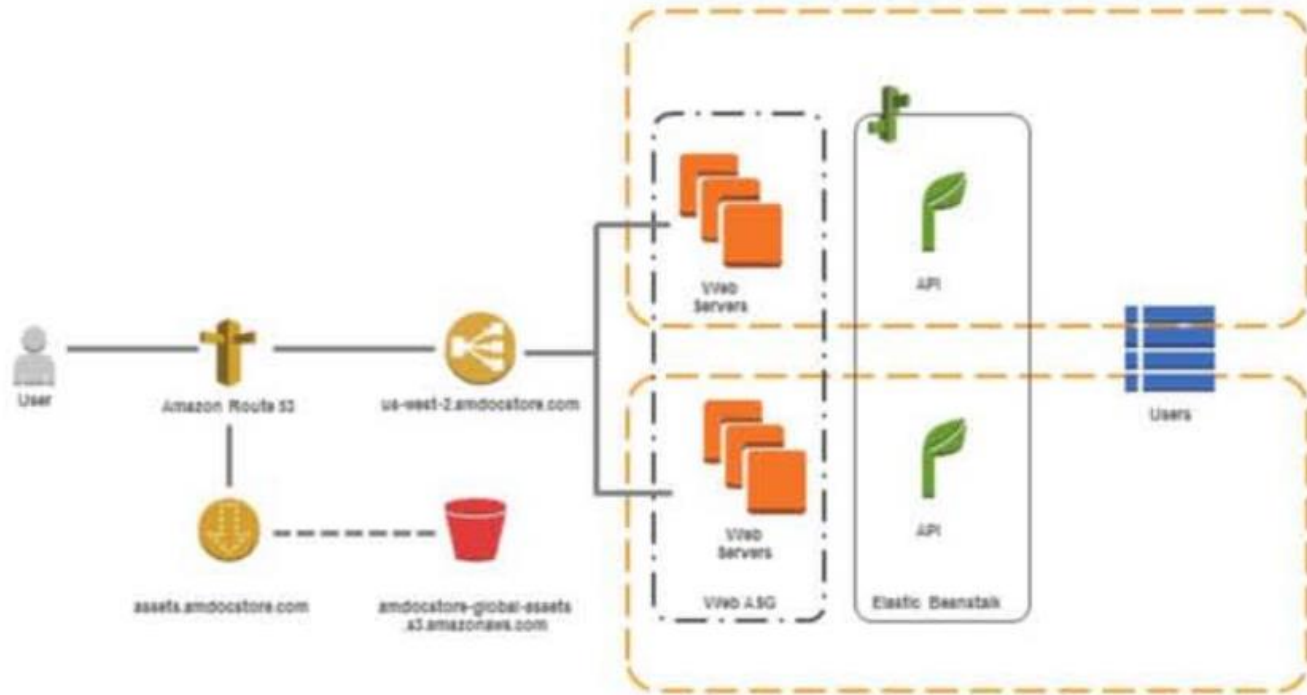
The API will use several different data stores for different types of data

## DocStore

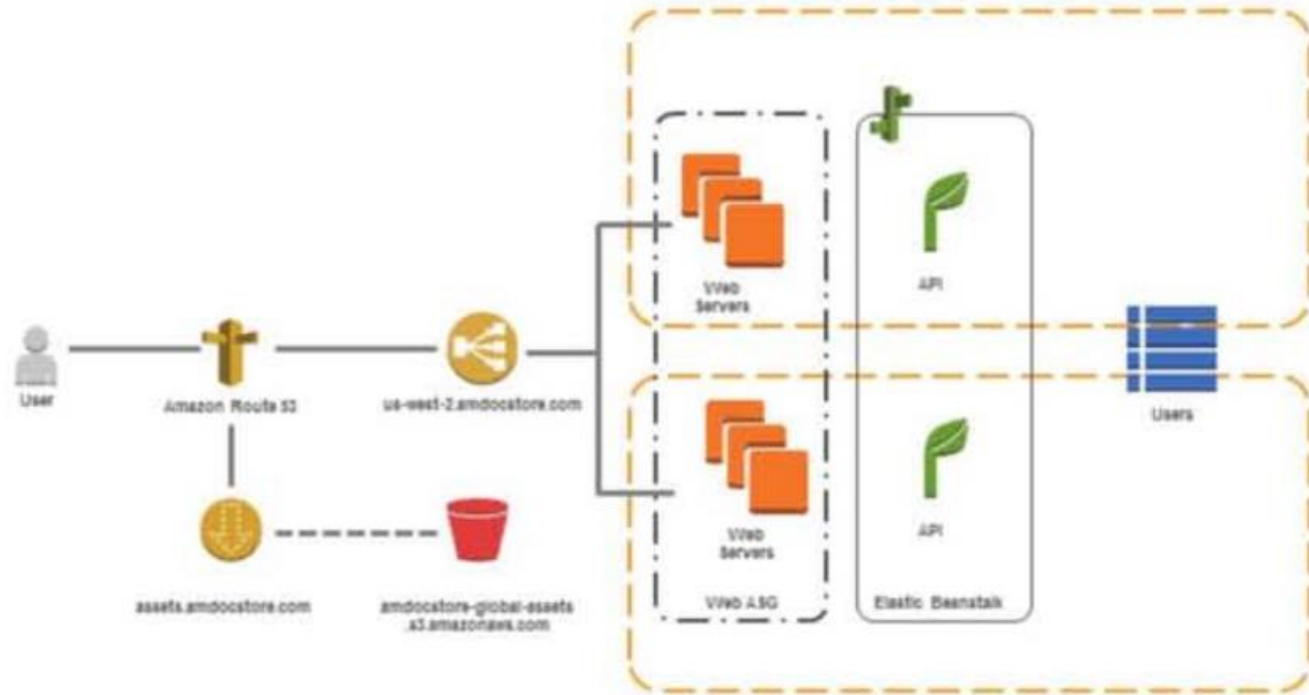


**User account information** will be stored in a **DynamoDB** table

**DocStore**

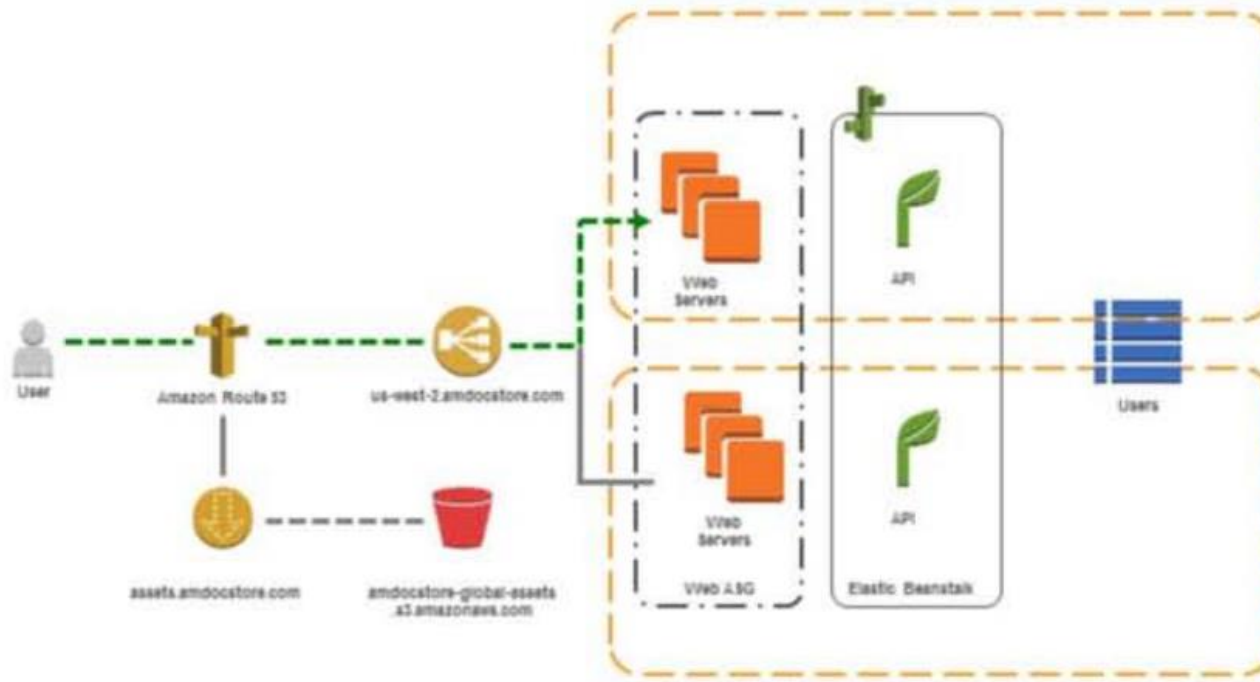


Let's see what happens when a new user signs up for **DocStore**



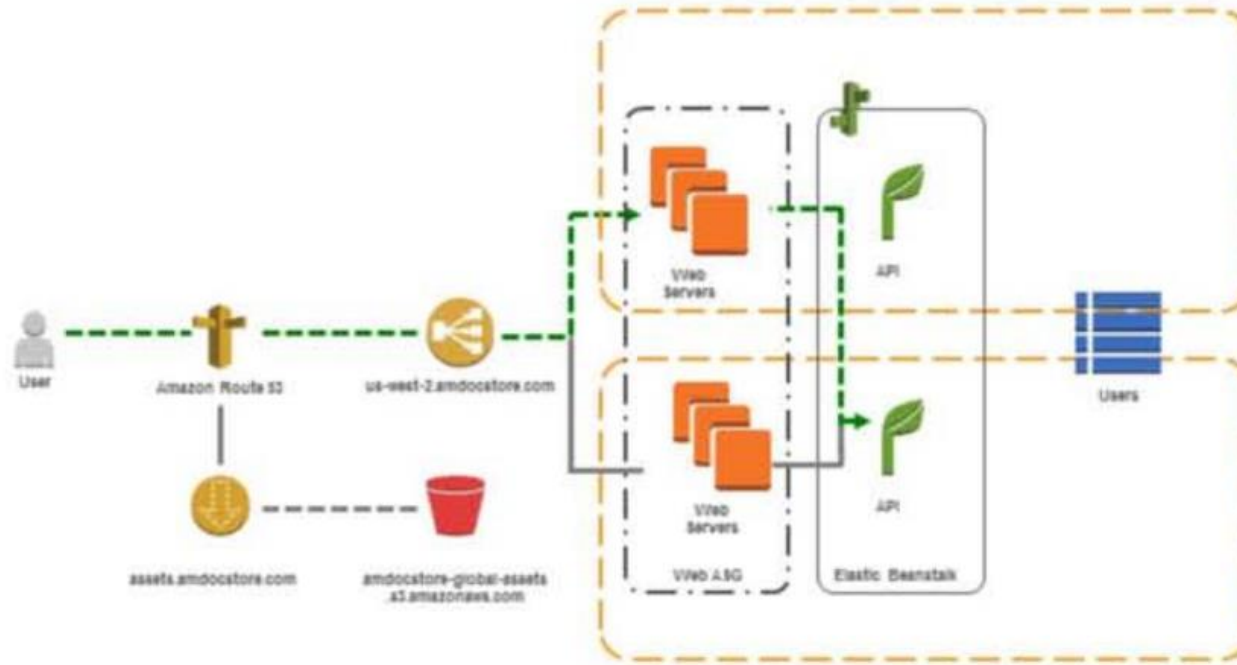
**Create Account:** User submits sign-up form to web servers with username, password, etc.

## DocStore



**Create Account:** Web servers invoke  
CreateAccount API

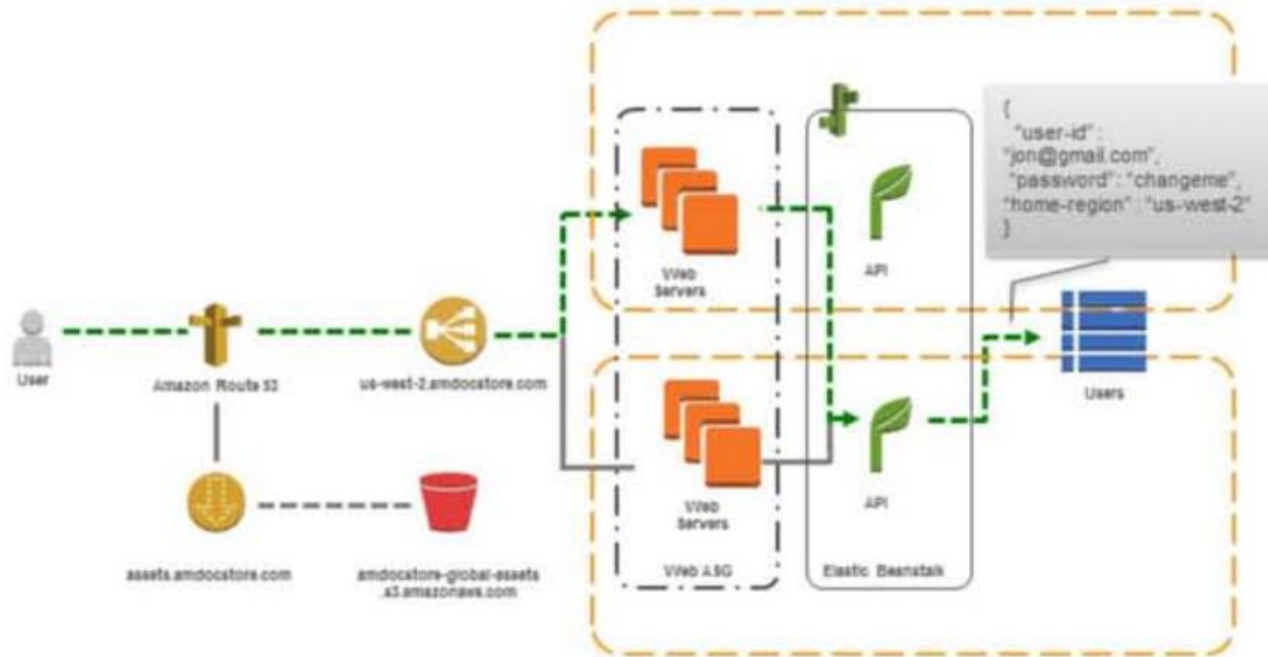
**DocStore**





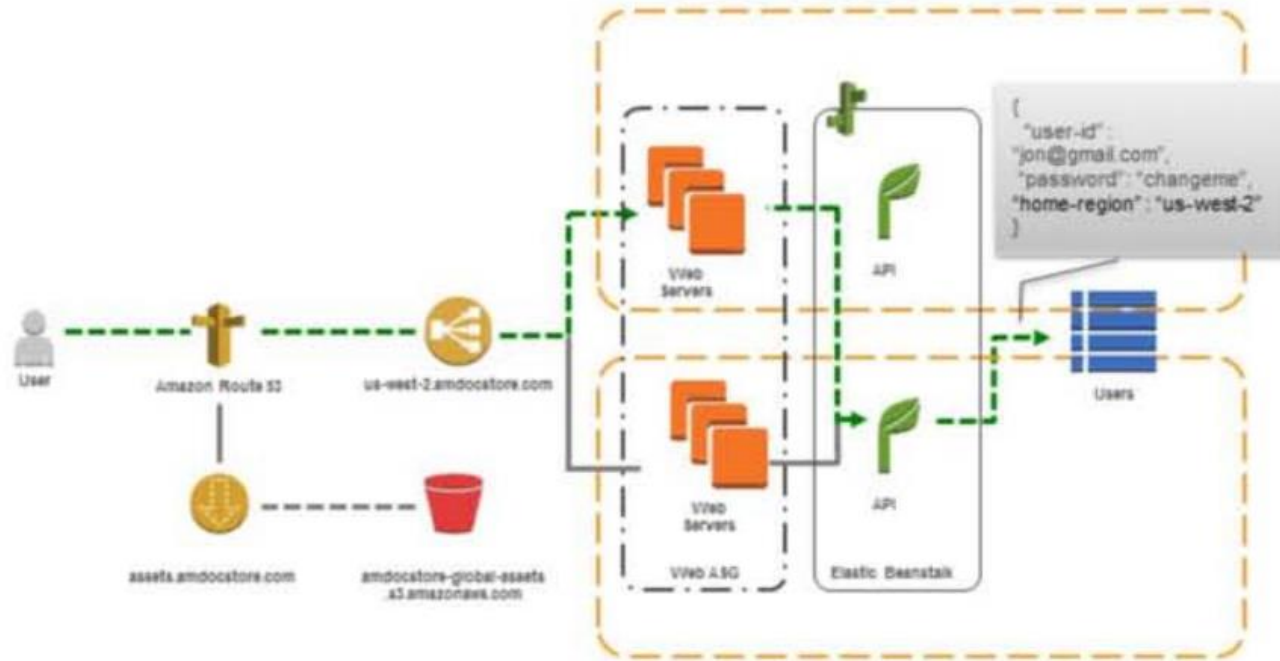
**Create Account:** API creates a **new item** in the datastore

## DocStore



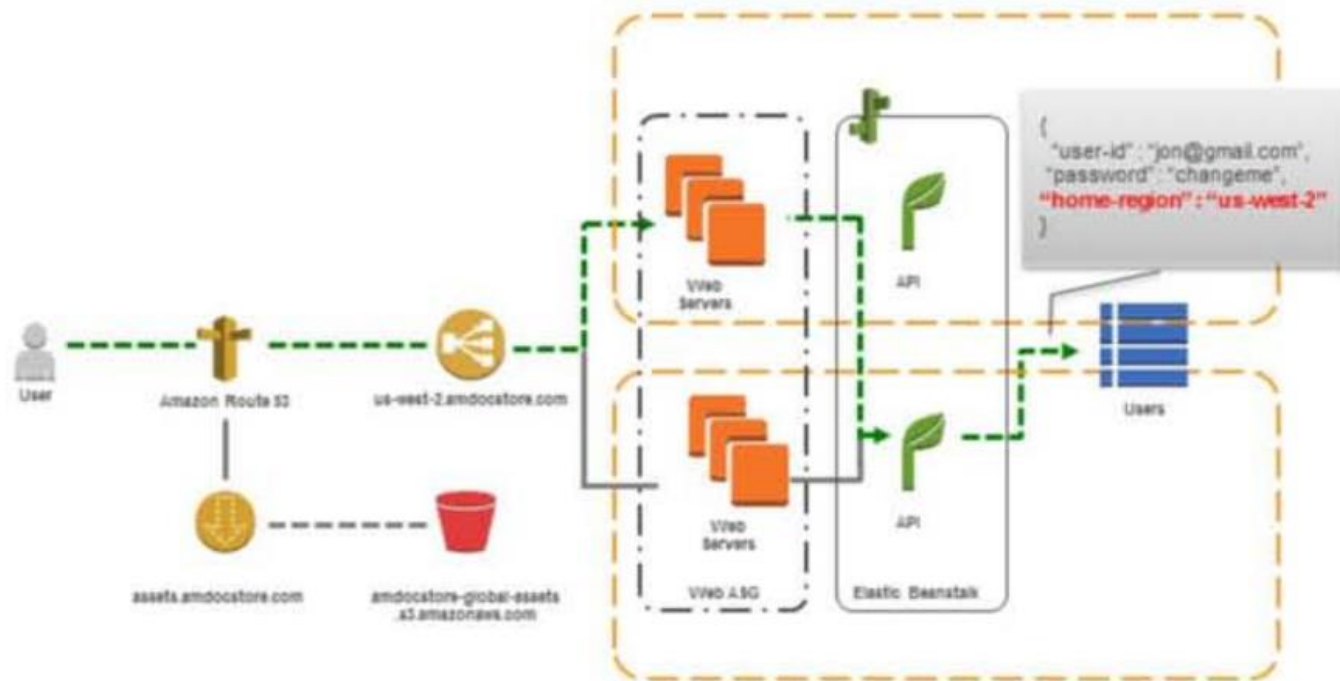
**Create Account:** The API app deployed in Elastic Beanstalk has a configuration setting making it aware of the region it is running in

# DocStore




**Create Account:** In addition to standard information, the API stores the user's "home region", i.e., the region the user signed up in

## DocStore



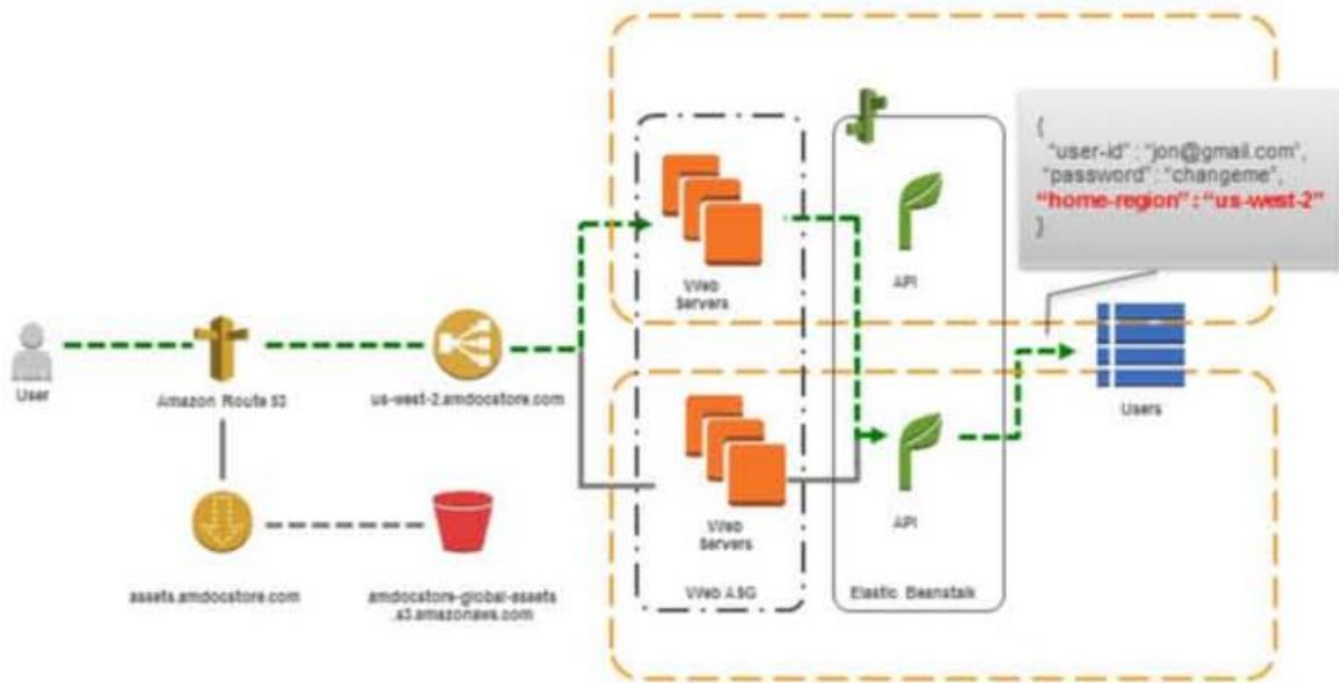
**Create Account:** Replicating account data to our Sydney deployment will allow the API in that region to redirect the user to his home region if he ever accesses from there

# DocStore



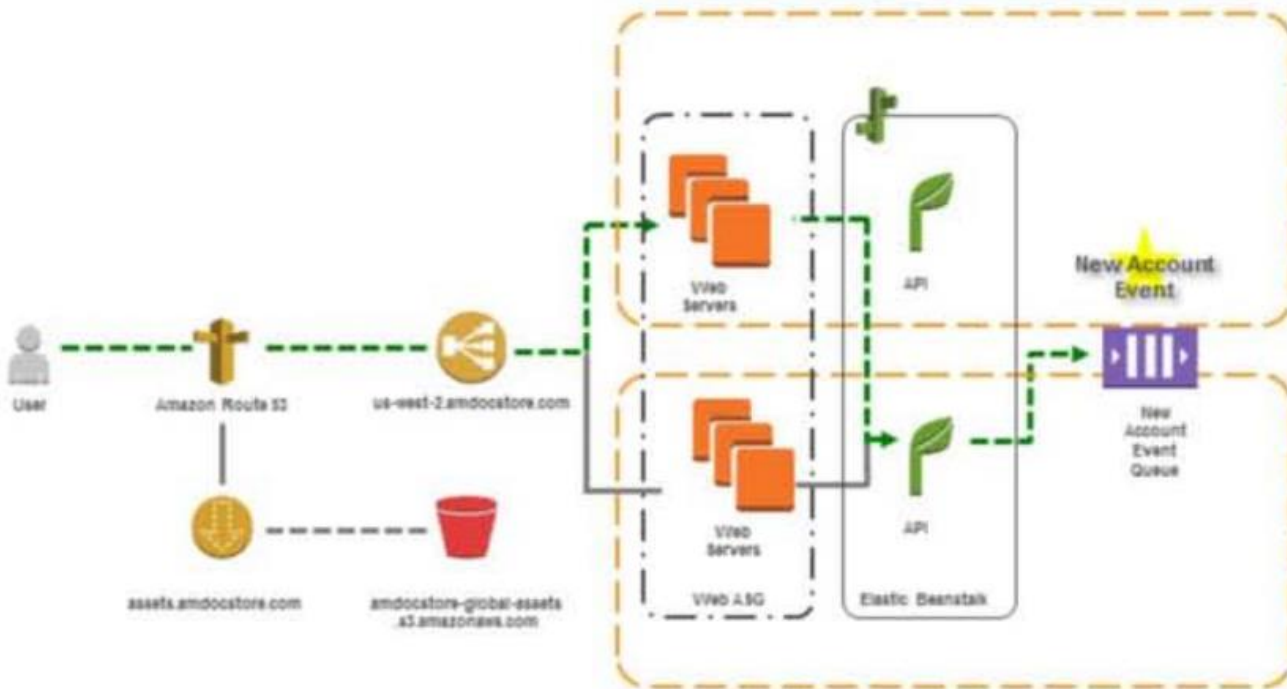
**Create Account:** We'll consider account signup an event that other components of our application might be interested in.

# DocStore



**Create Account:** We'll let listeners interested in the event (i.e., a cross-region account replication service) receive the event by putting a message in a queue

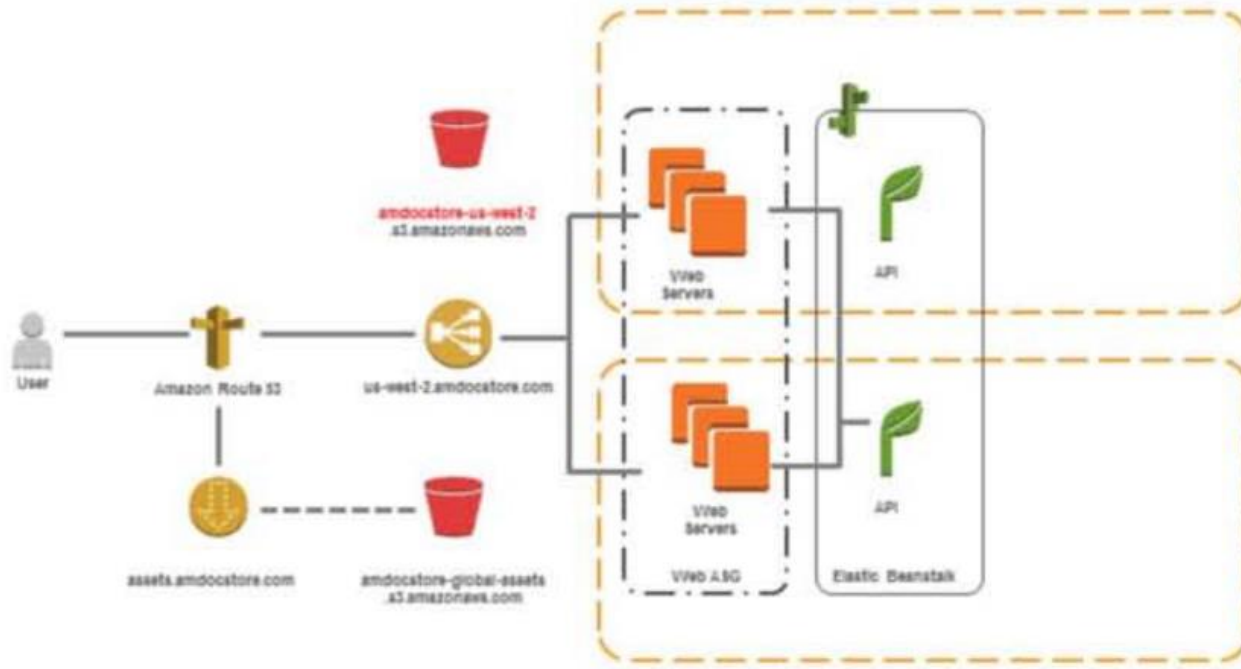
## DocStore





The **S3 bucket name** includes the name of the region it was created in.

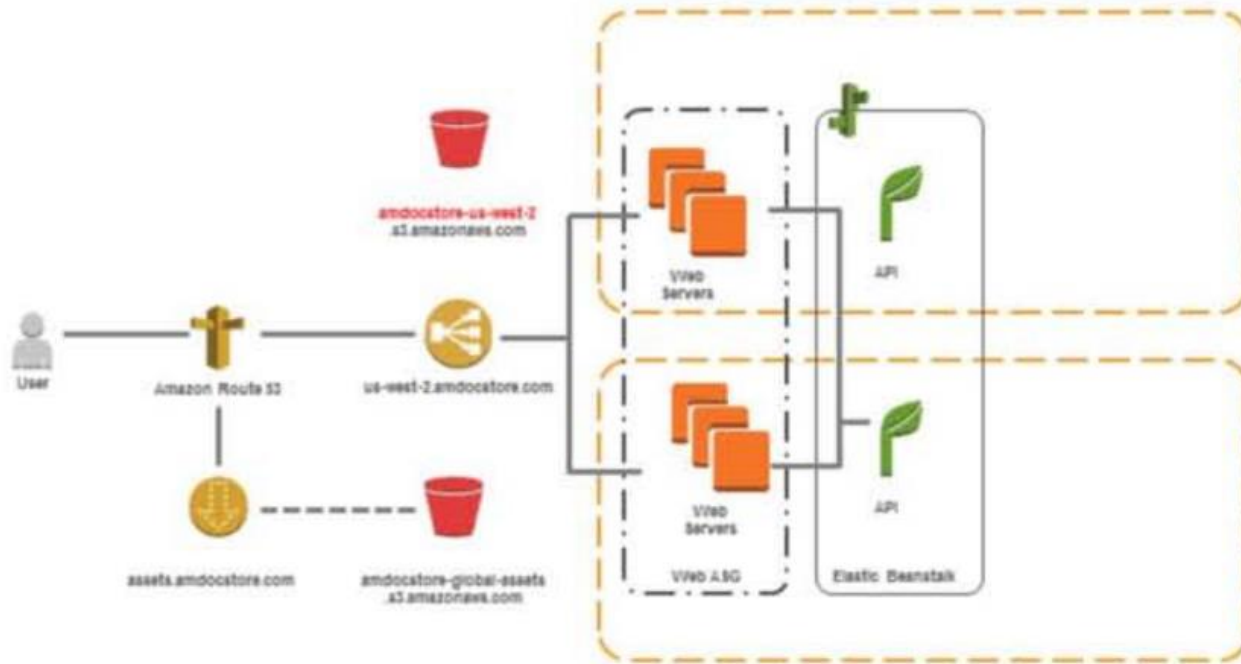
## DocStore





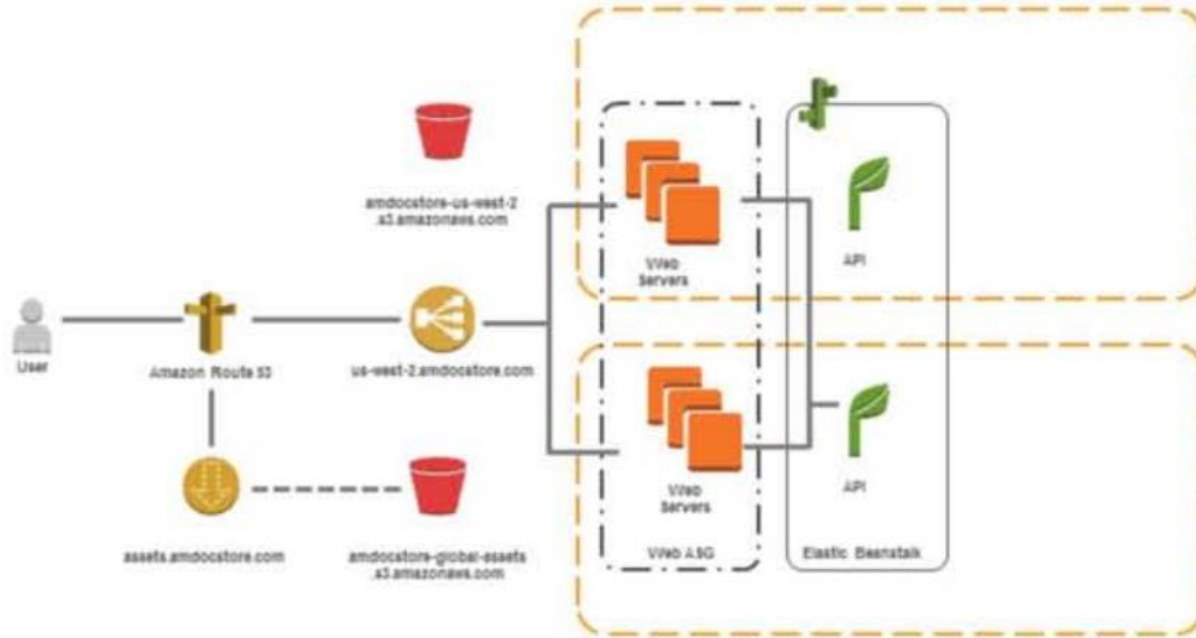
The S3 bucket name includes the name of the region it was created in. Users **upload documents to the bucket in their home region**

## DocStore



Let's see what happens when a user wants to  
upload a PDF document

## DocStore

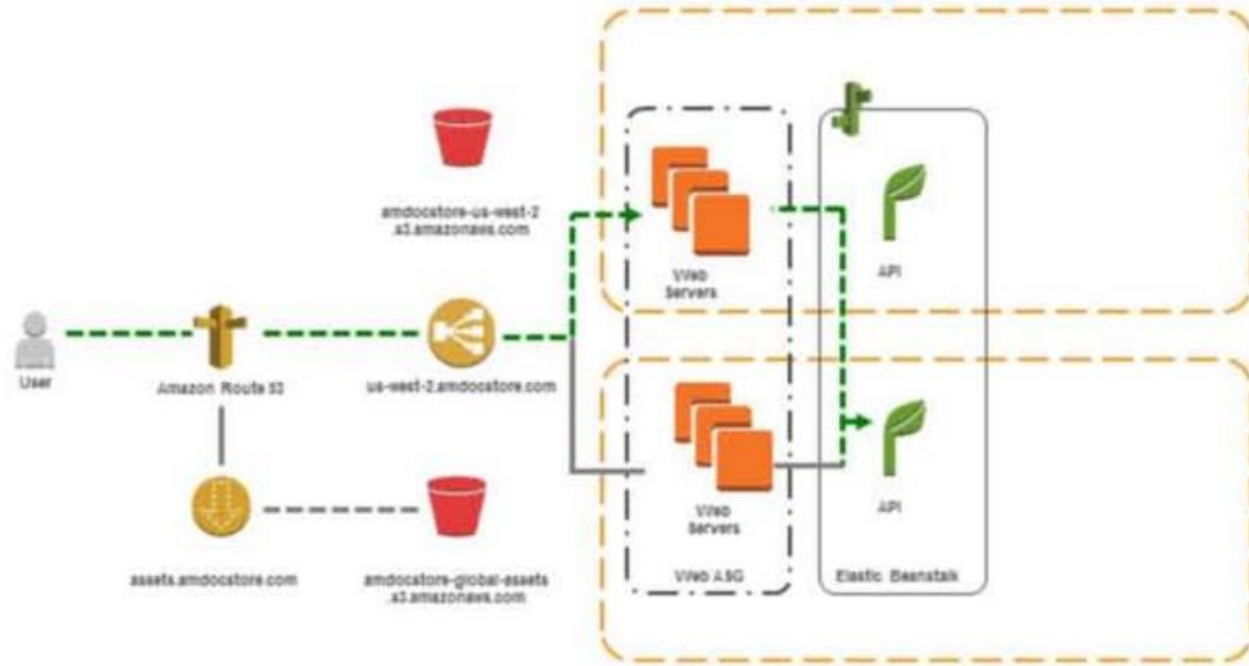


# DocStore



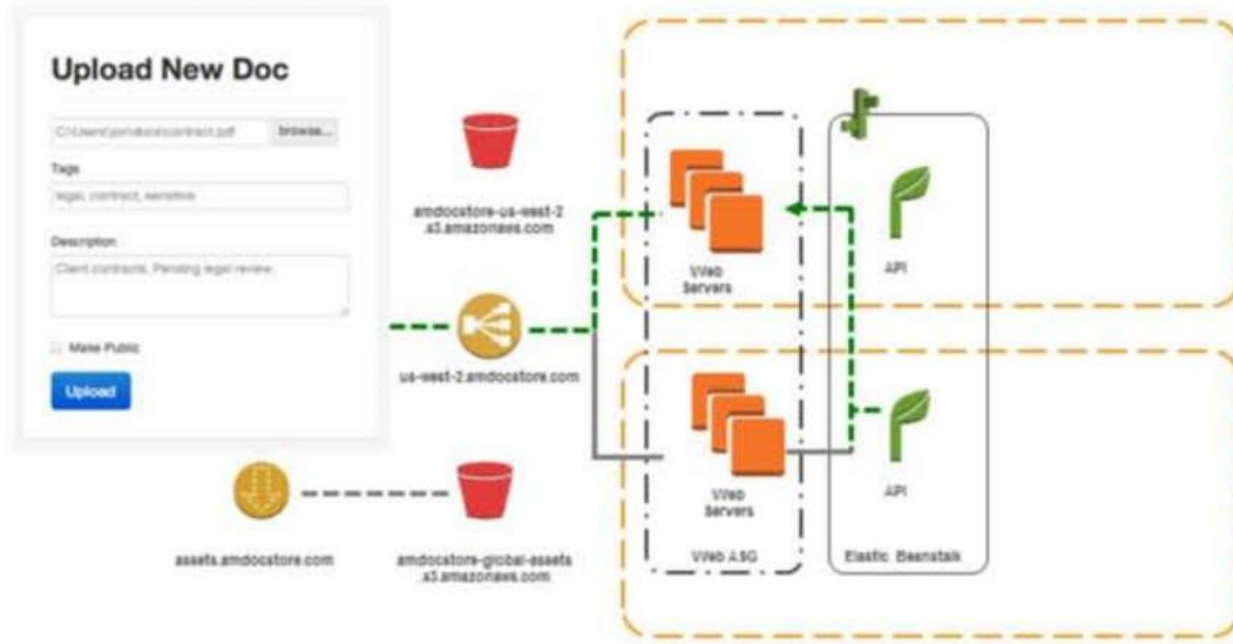
**Upload Document:** The web servers request a signed HTML form from the API and return it to the user's browser

## DocStore



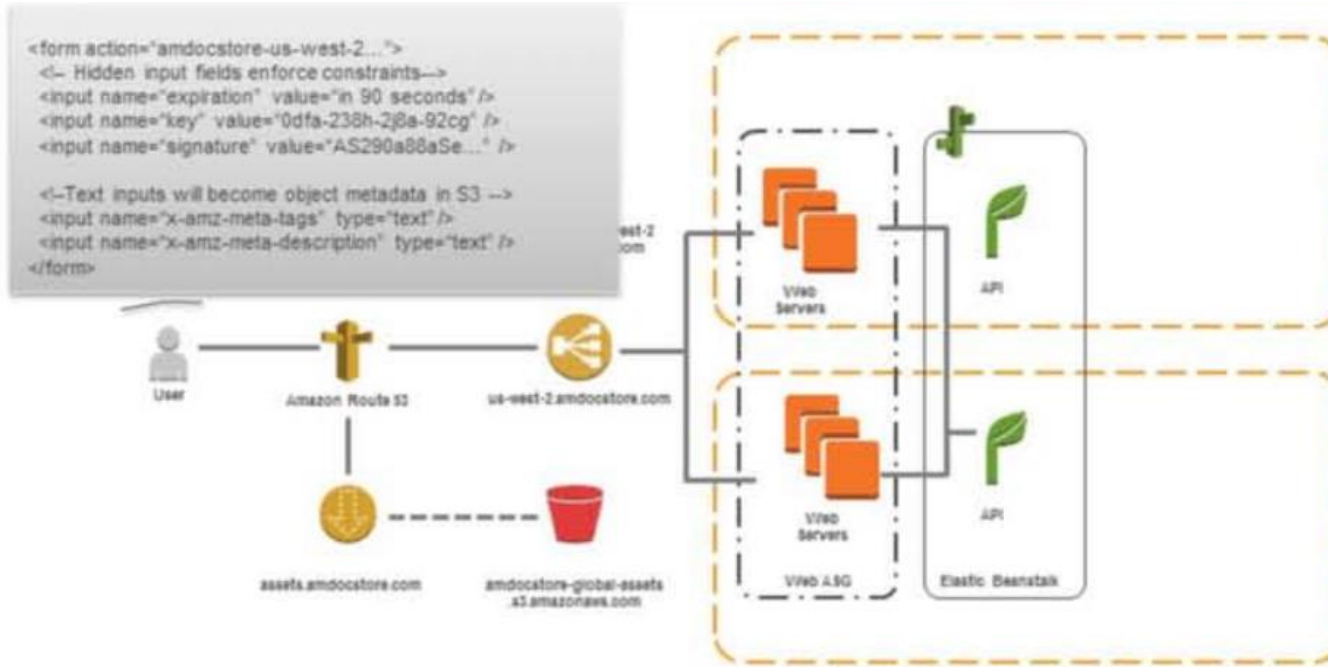
**Upload Document:** The signed form is delivered to the user's browser

**DocStore**



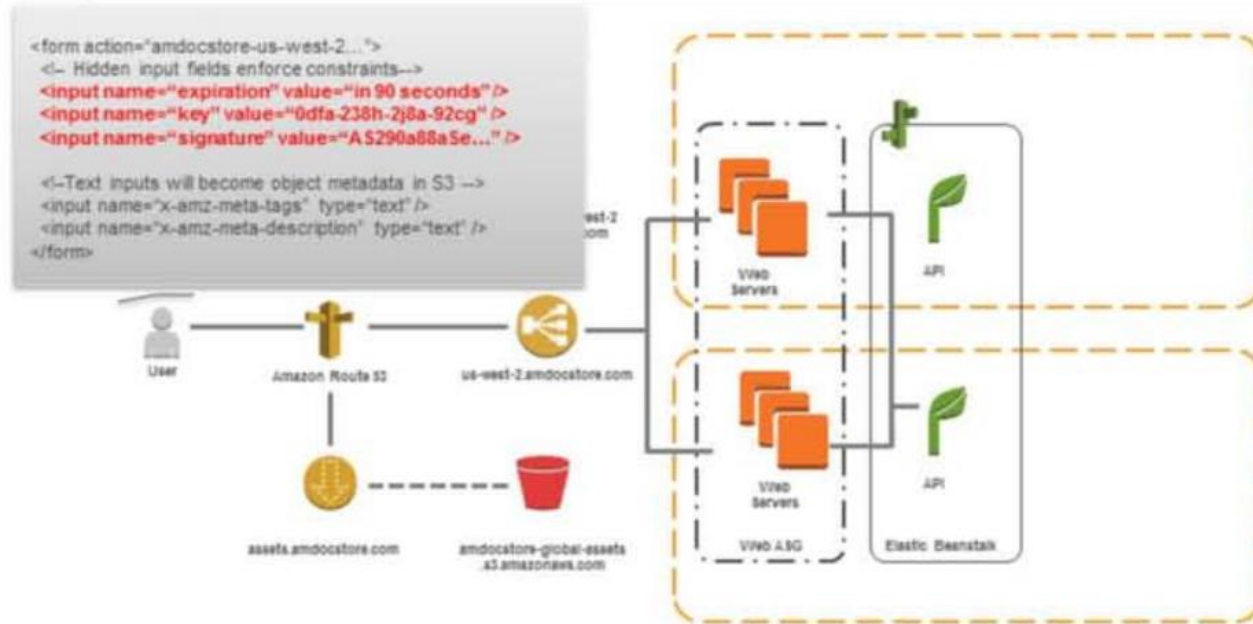
**Upload Document:** The signed form allows the user to upload his document **directly to S3**

## DocStore



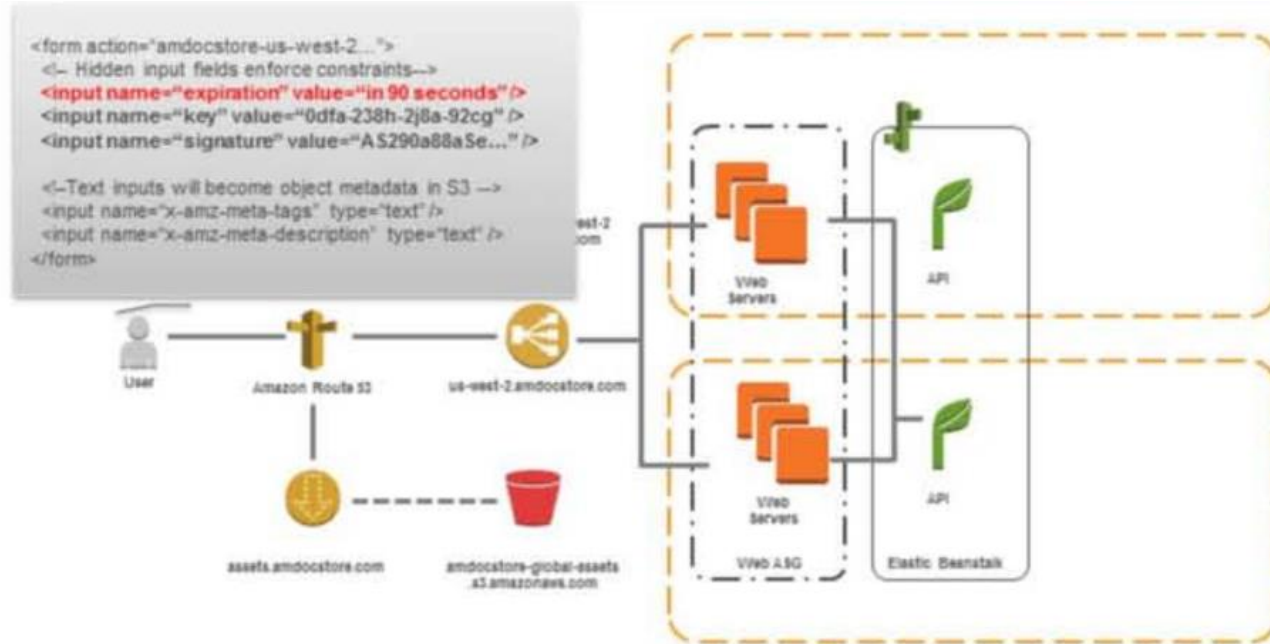
**Upload Document:** The form includes some **hidden inputs that constrain** what the user can upload.

## DocStore



**Upload Document:** The form includes some **hidden inputs that constrain** what the user can upload. The **expiration** indicates for how long the form is valid

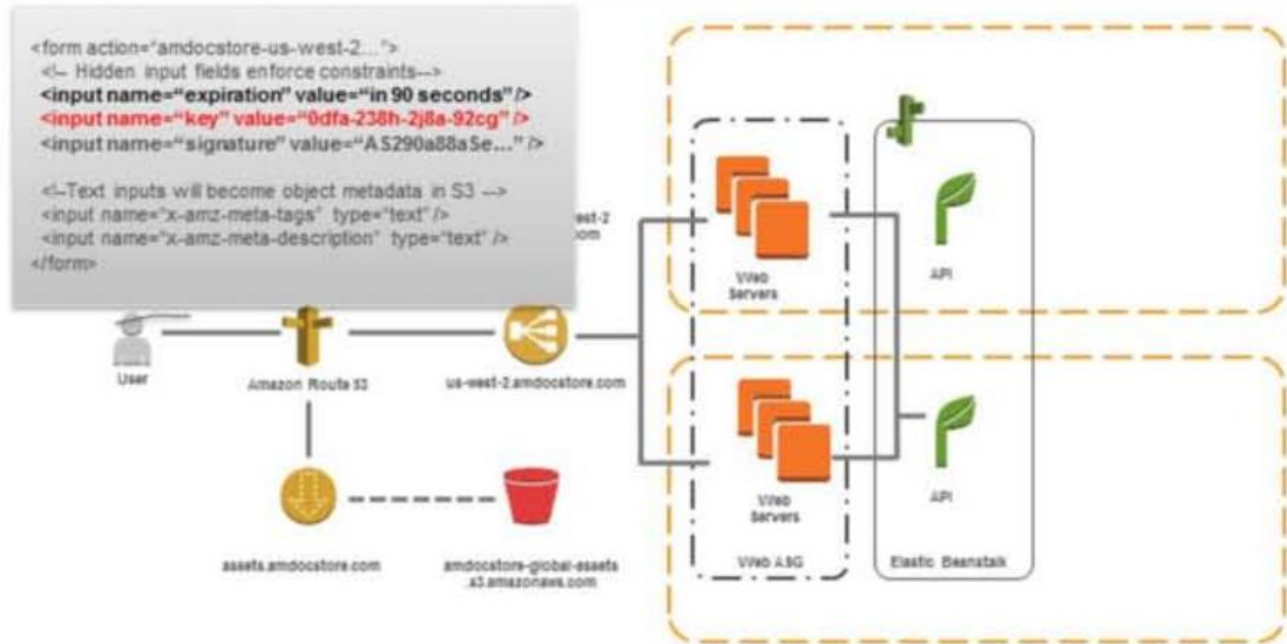
## DocStore





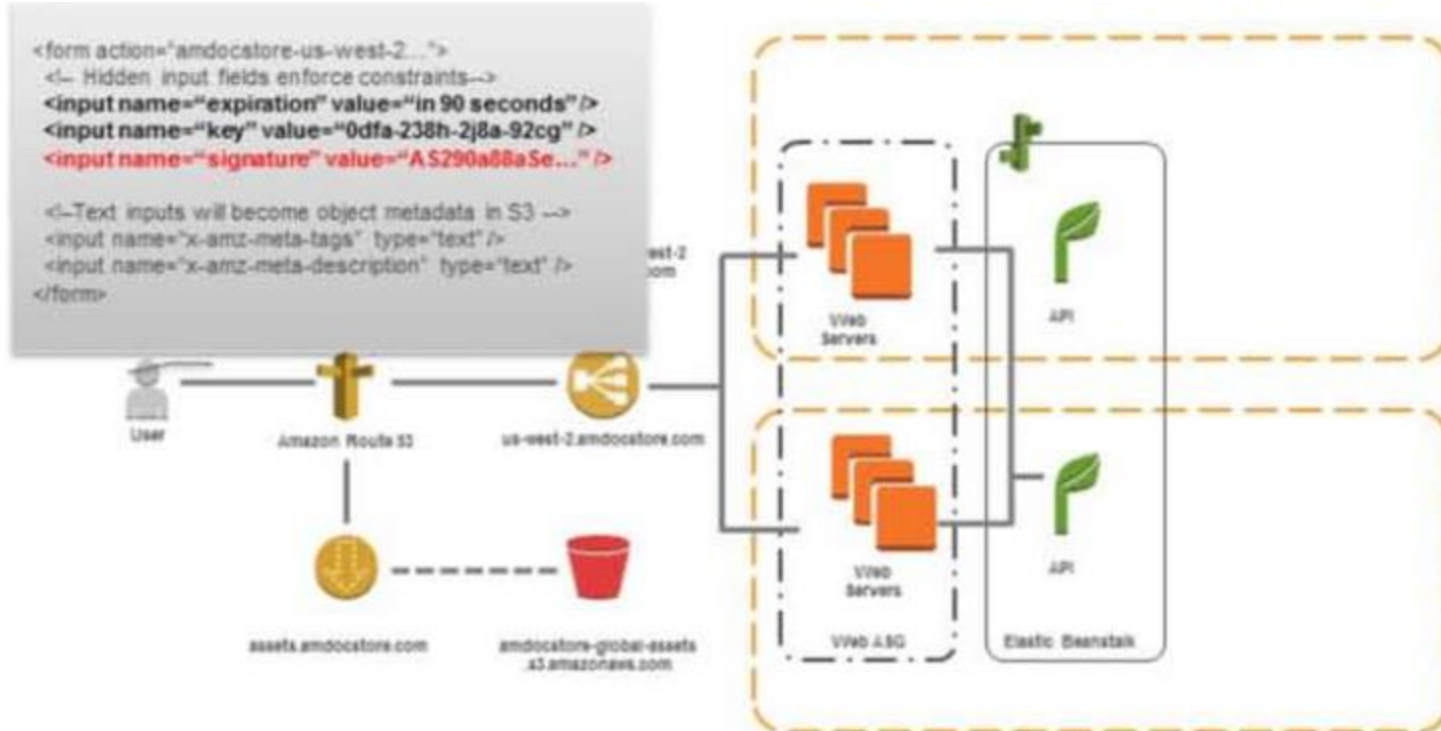
**Upload Document:** The form includes some **hidden inputs** that **constrain** what the user can upload. **The key** allows us to specify what the object will be called in S3 (thus avoiding naming conflicts)

# DocStore



Upload Document: The form includes some hidden inputs that constrain what the user can upload. The signature ensures the upload will be rejected if the form is tampered with

# DOCSTORE



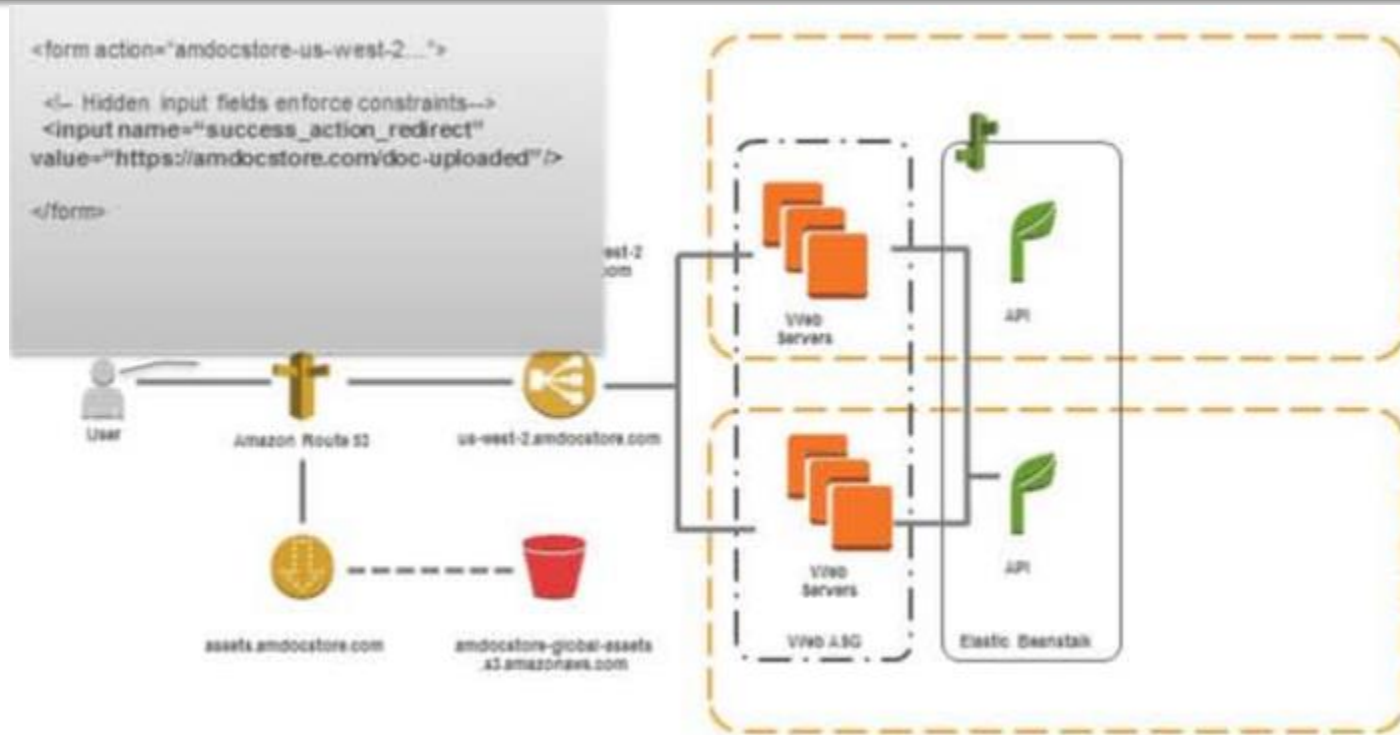
Upload Document: User-editable text input fields prefixed with x-amz-  
meta-will be associated with the uploaded object as S3 object metadata

# DOCSTORE



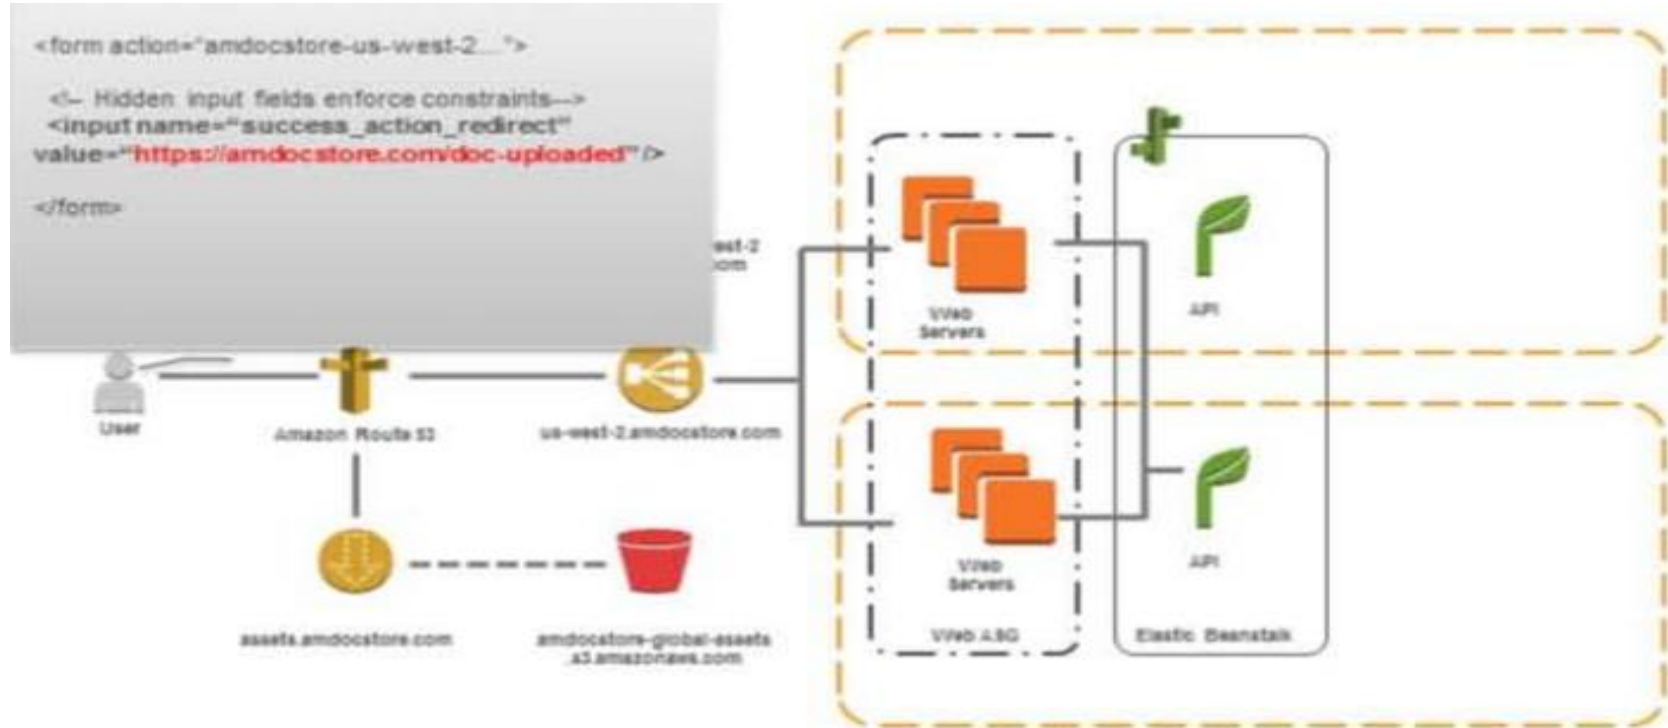
Upload Document: Finally, the success\_action\_redirect specified when the API generated the form tells S3 where to redirect the user's browser when the upload completes

# DOCSTORE



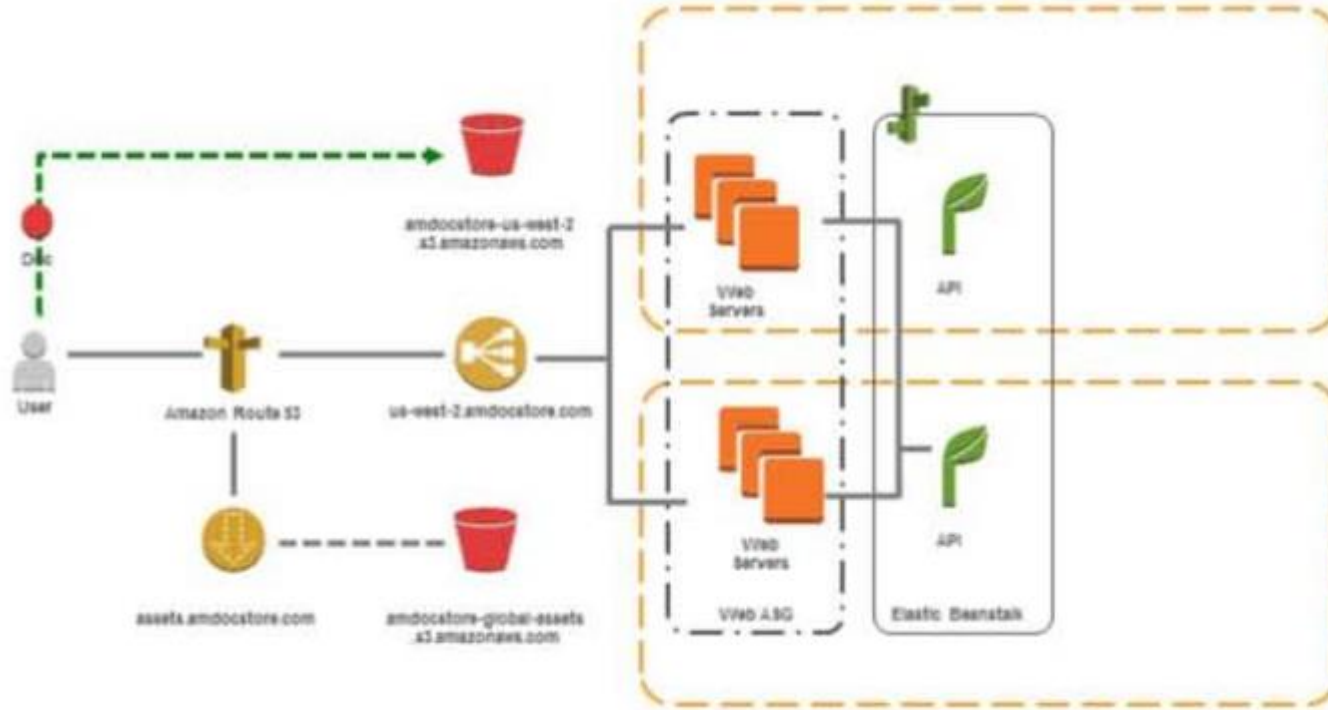
Upload Document: We send the user to a special page on the web servers that will capture this event

# DOCSTORE



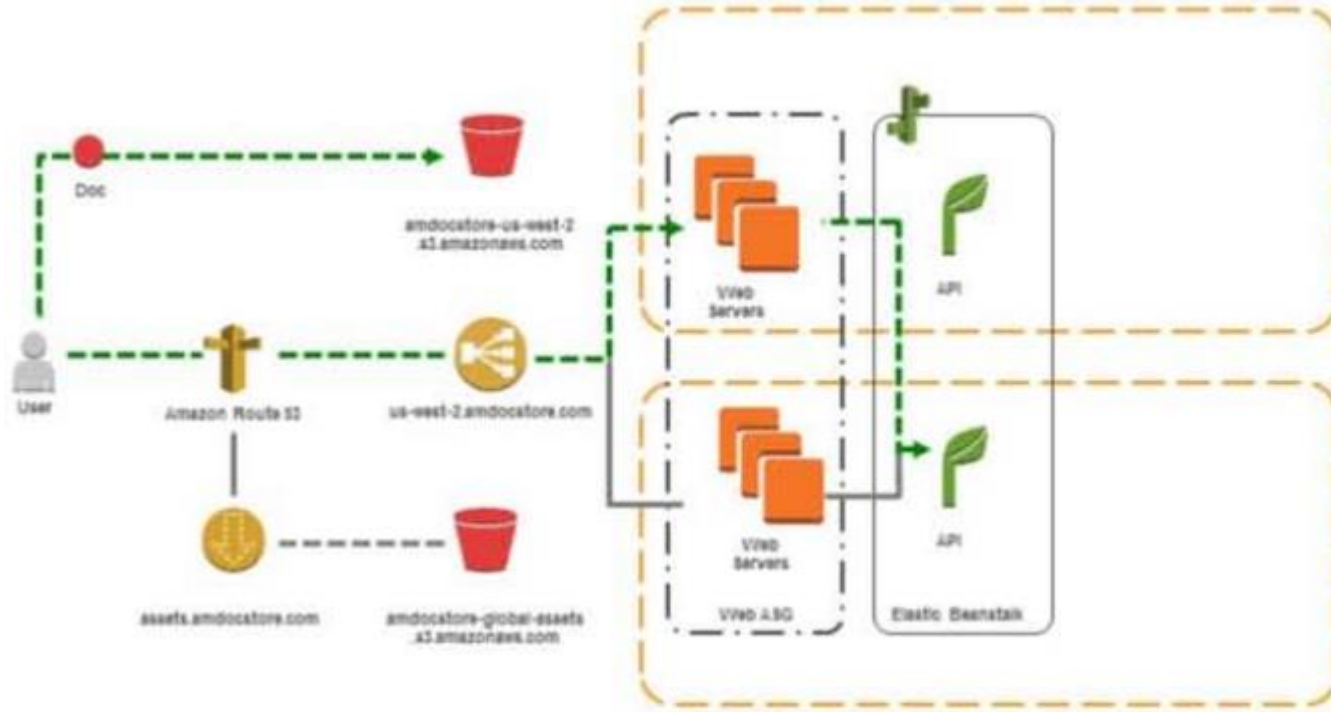
Upload Document: While uploading the document to S3

# DOCSTORE



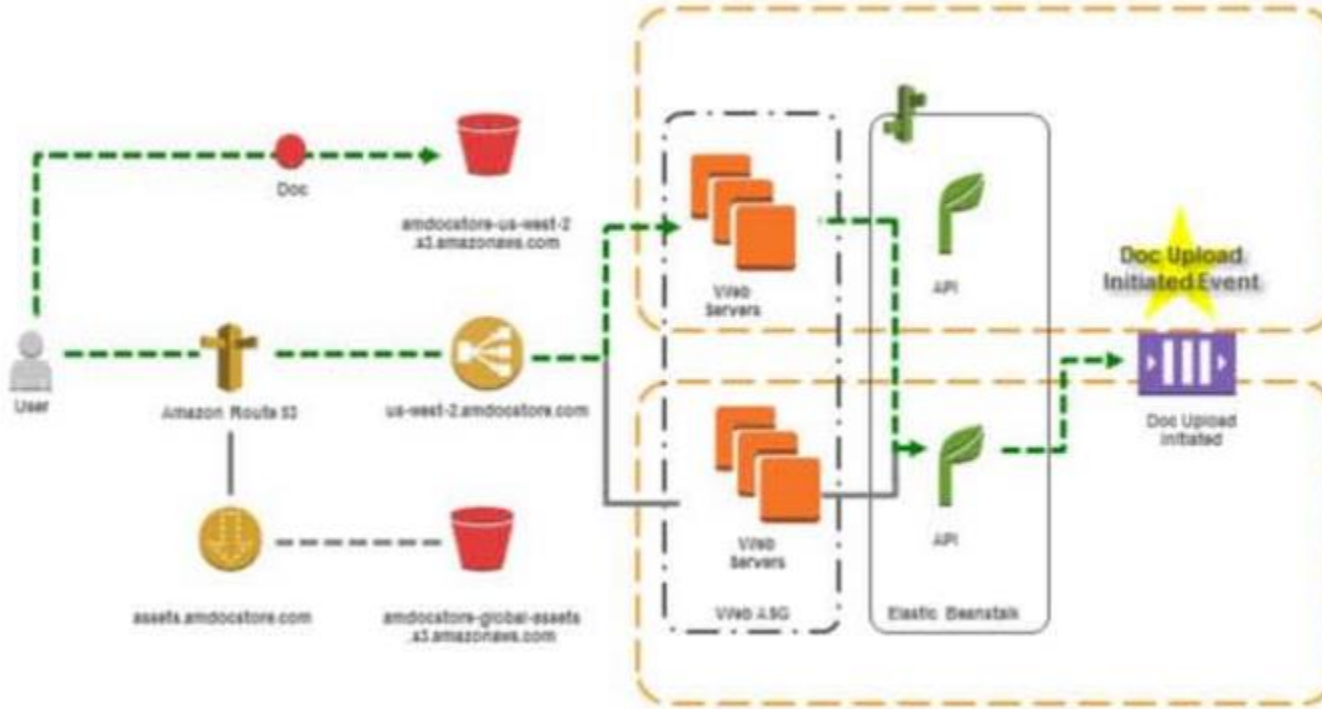
Upload Document: While uploading the document to S3, we asynchronously notify the web servers that the upload has begun

# DOCSTORE



Upload Document: The API places a message in a queue indicating as much. Other components of DocStore may be interested in this event later

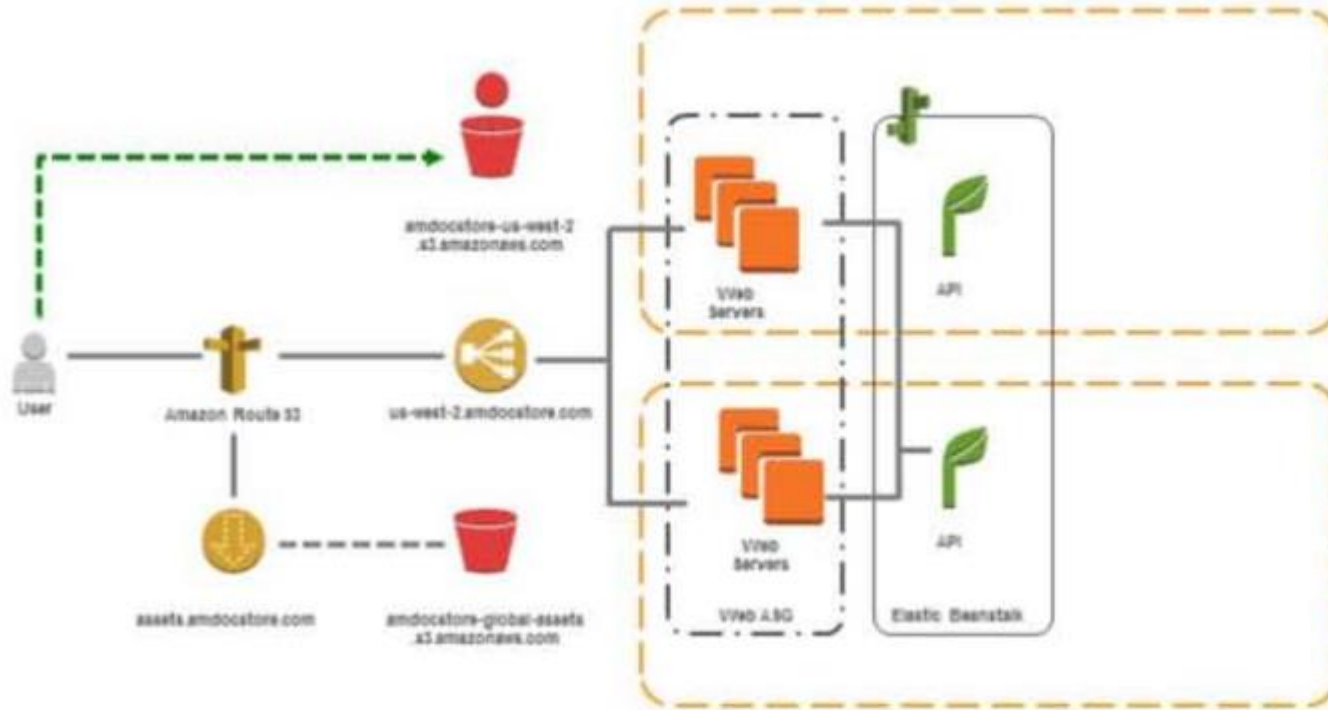
# DOCSTORE





Upload Document: When the upload completes

# DOCSTORE

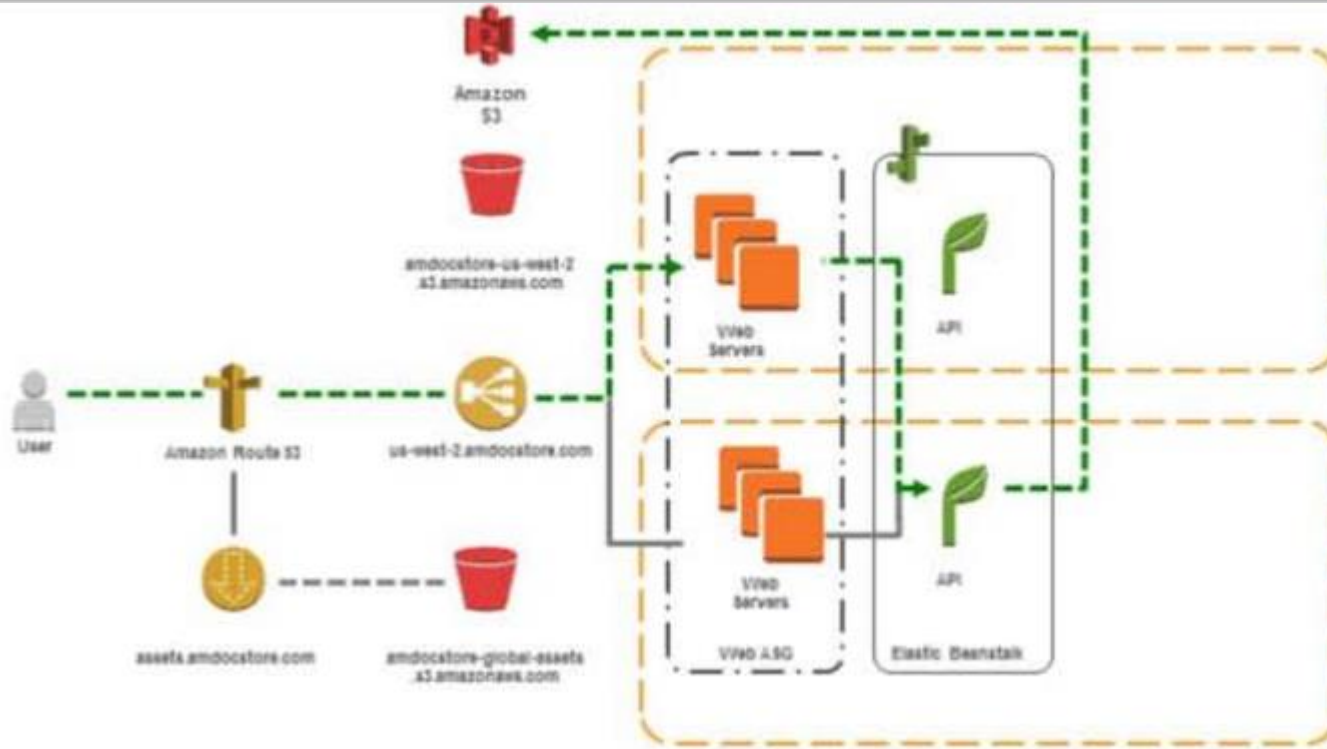


# DOCSTORE



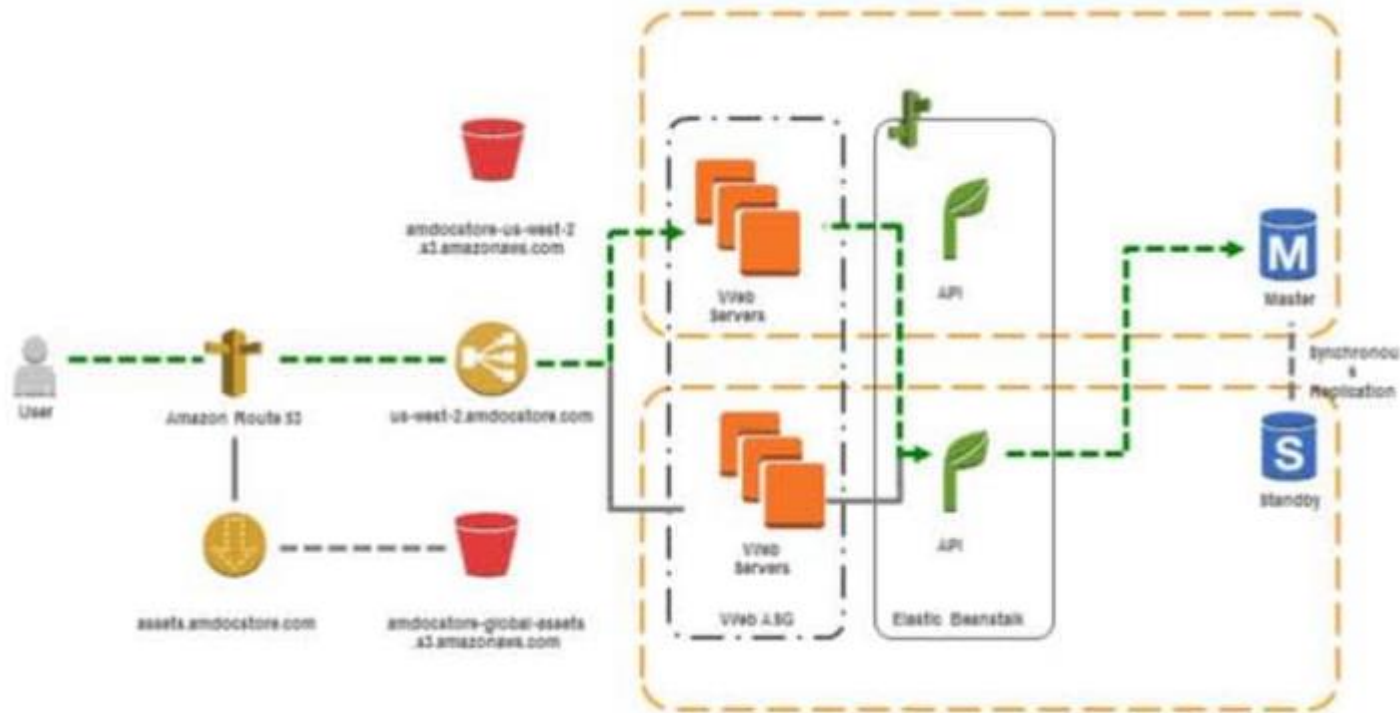
Upload Document: The API retrieves metadata about the object from the S3 API

# DOCSTORE



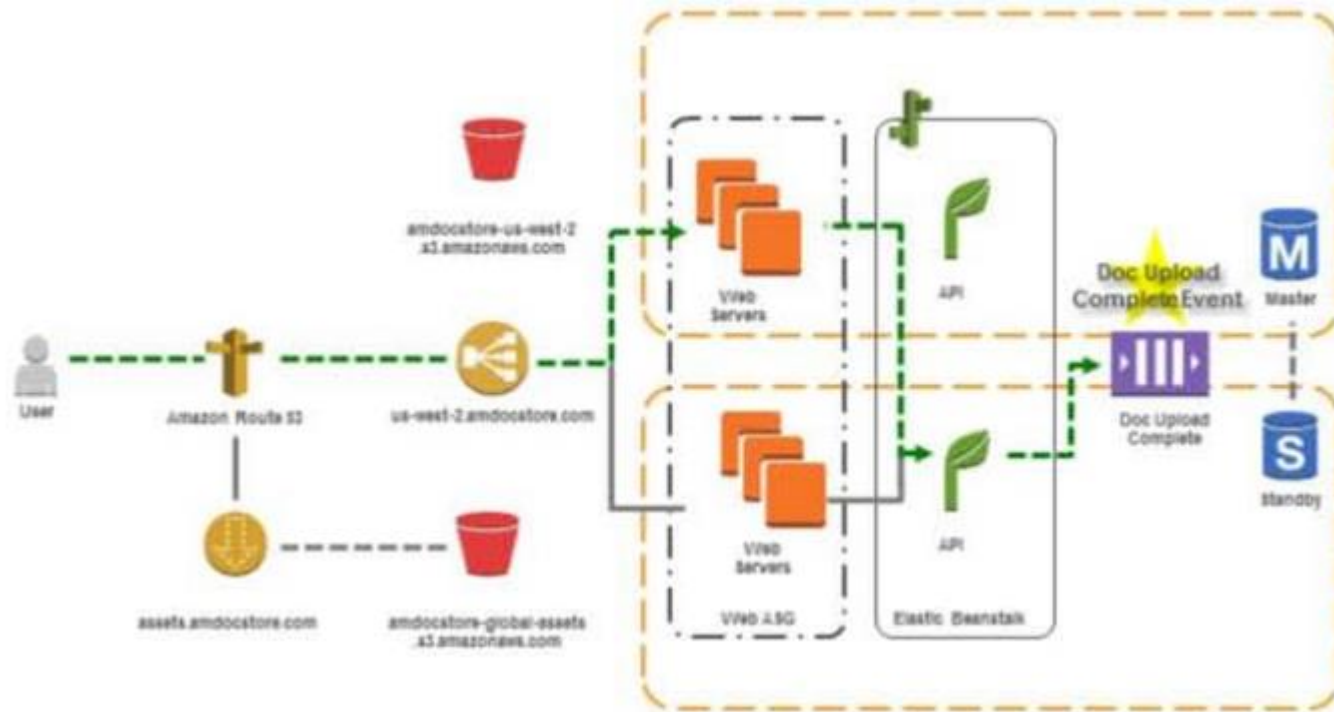
Upload Document: The Object metadata is stored in an RDS instance running in Multi-AZ mode for high availability

# DOCSTORE



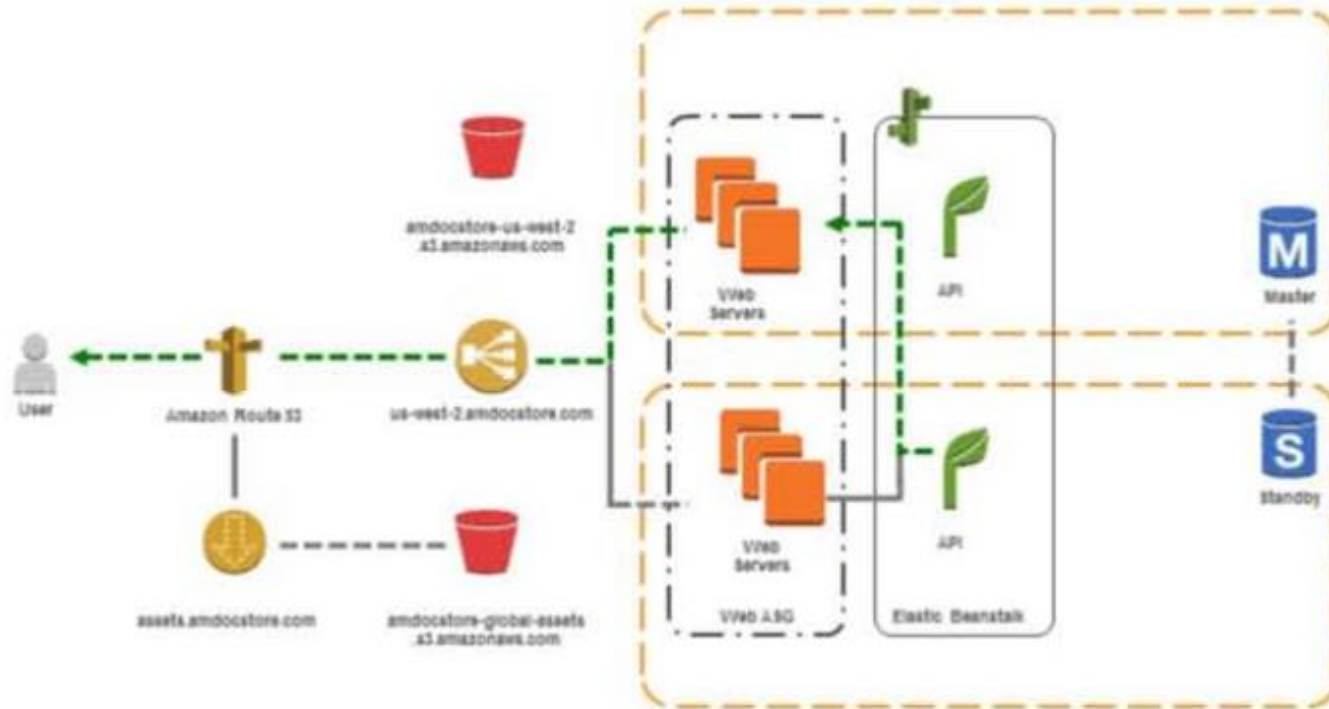
Upload Document: Finally, the API places a message in a queue indicating that the document upload was complete, and includes the ID of the newly uploaded object in S3

# DOCSTORE



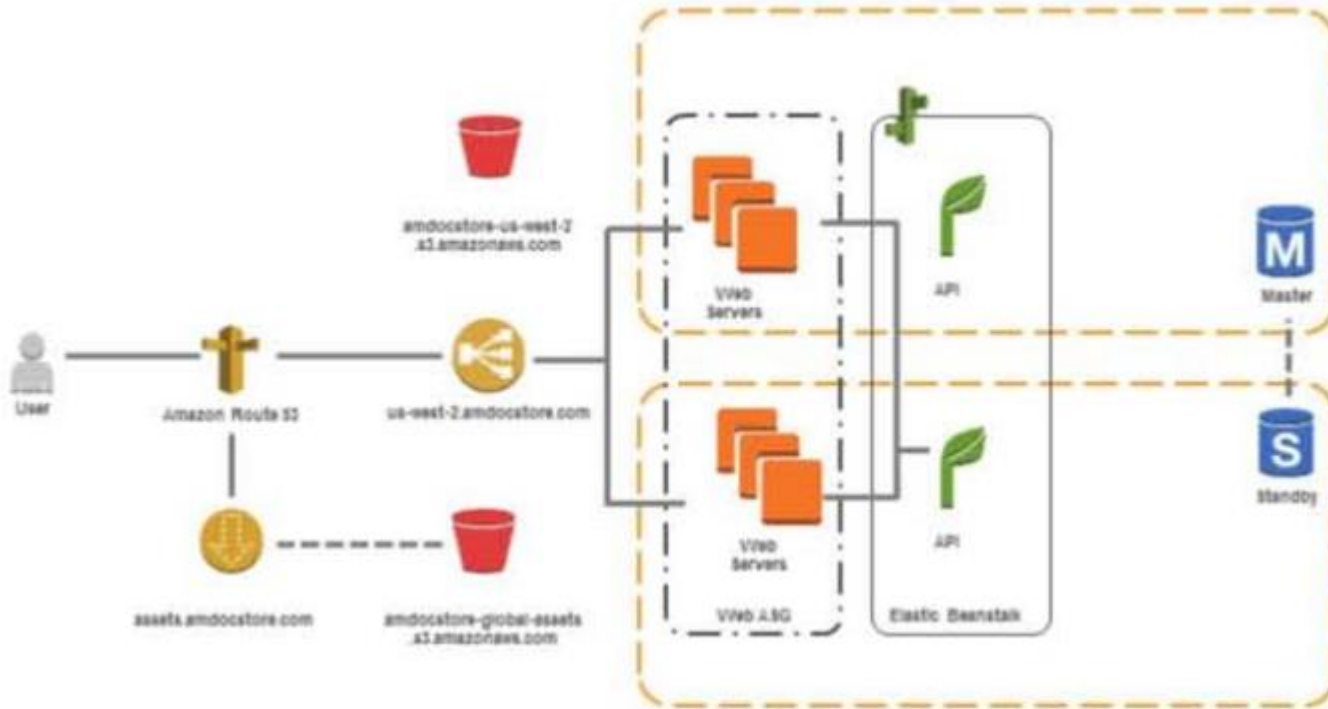
Upload Document: The user is redirected to a list of his documents

# DOCSTORE



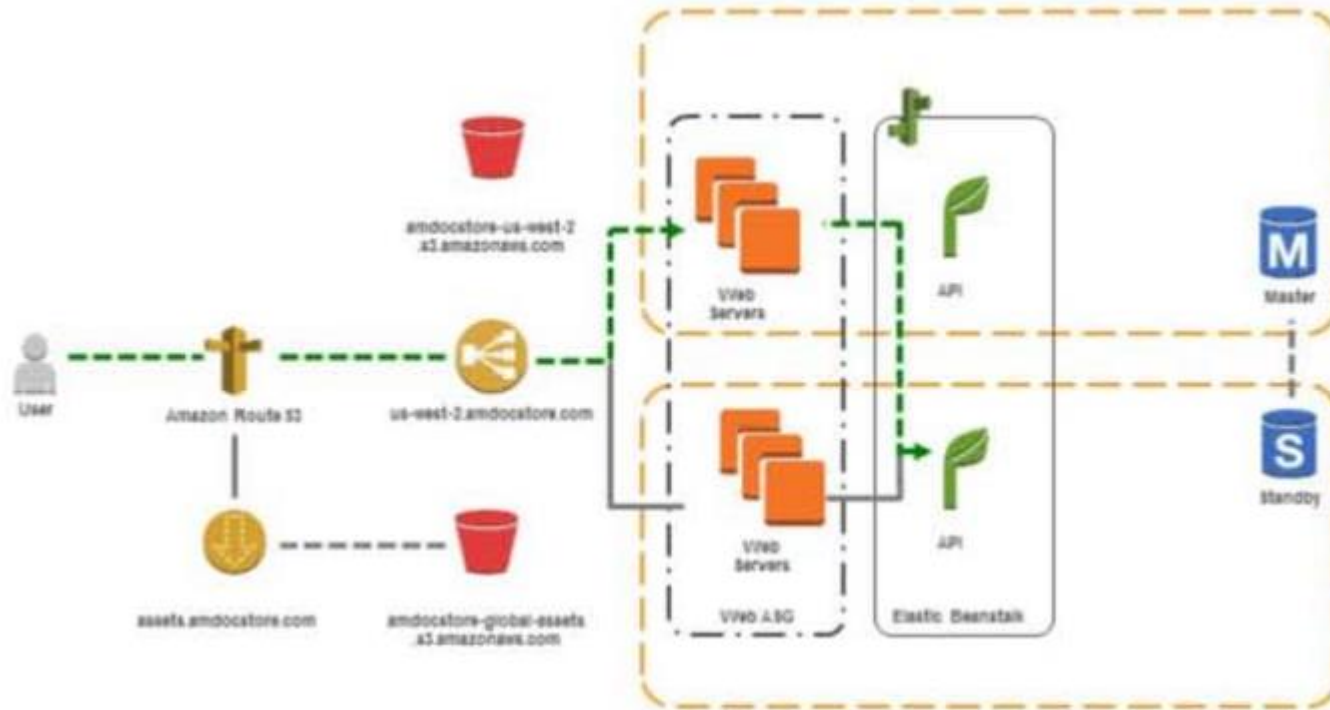
Let's see what happens when a user wants to view and download his documents

# DOCSTORE



View and Download Docs: An authenticated user requests the web site URL to view his docs

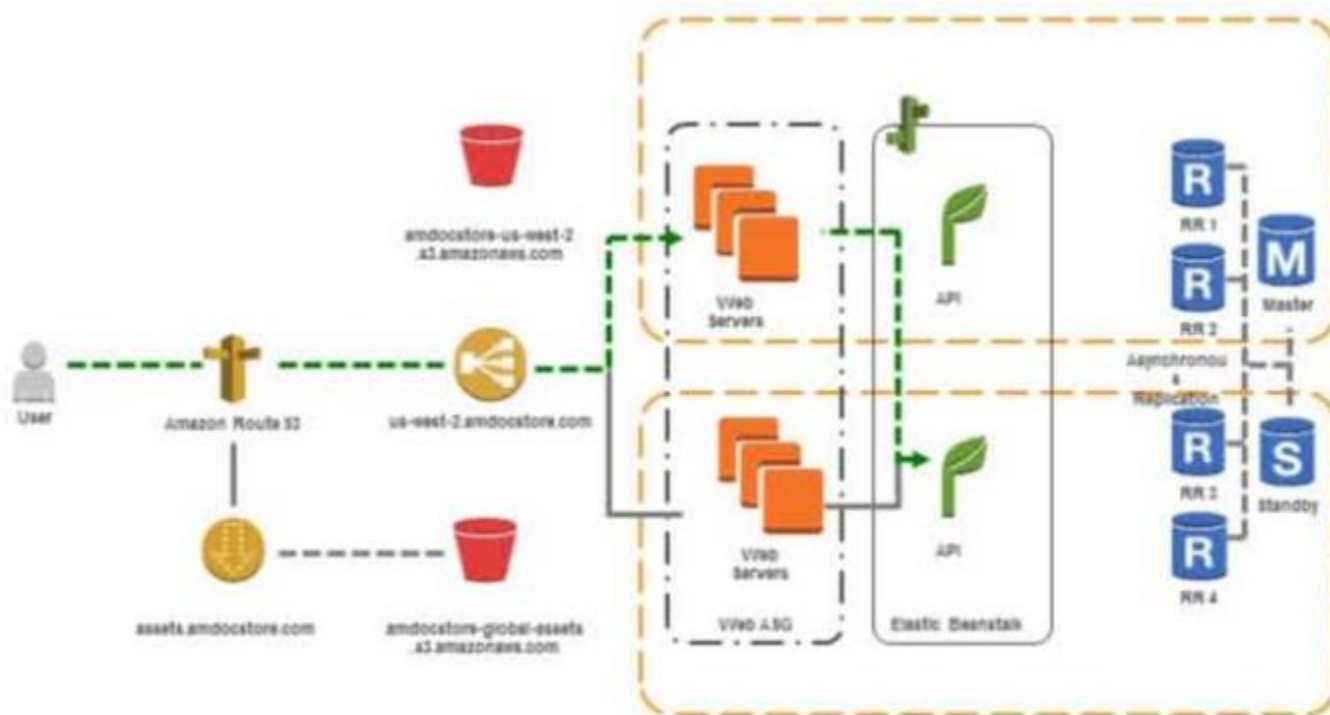
# DOCSTORE





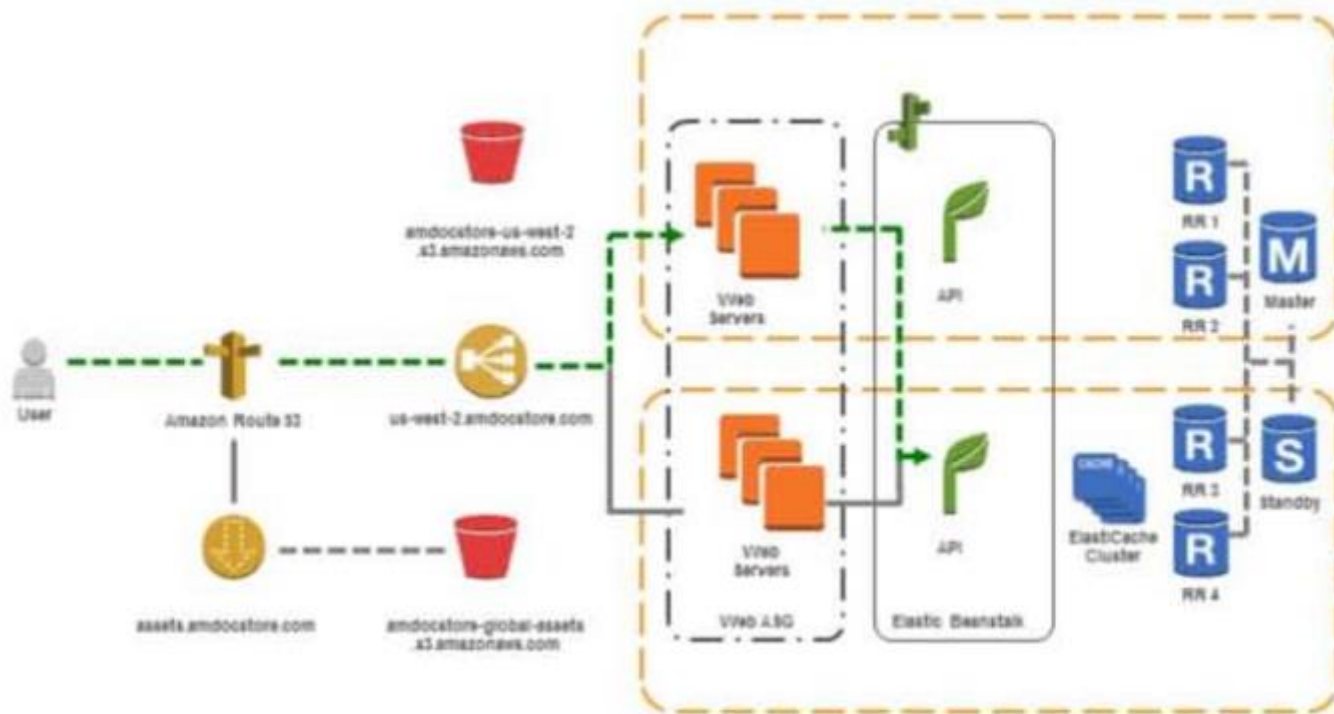
View and Download Docs: To increase read performance of the relational database, multiple RDS Read Replicas are provisioned. Read Replicas are suitable for read workloads, like the api selecting a list of documents owned by a user

# DOCSTORE



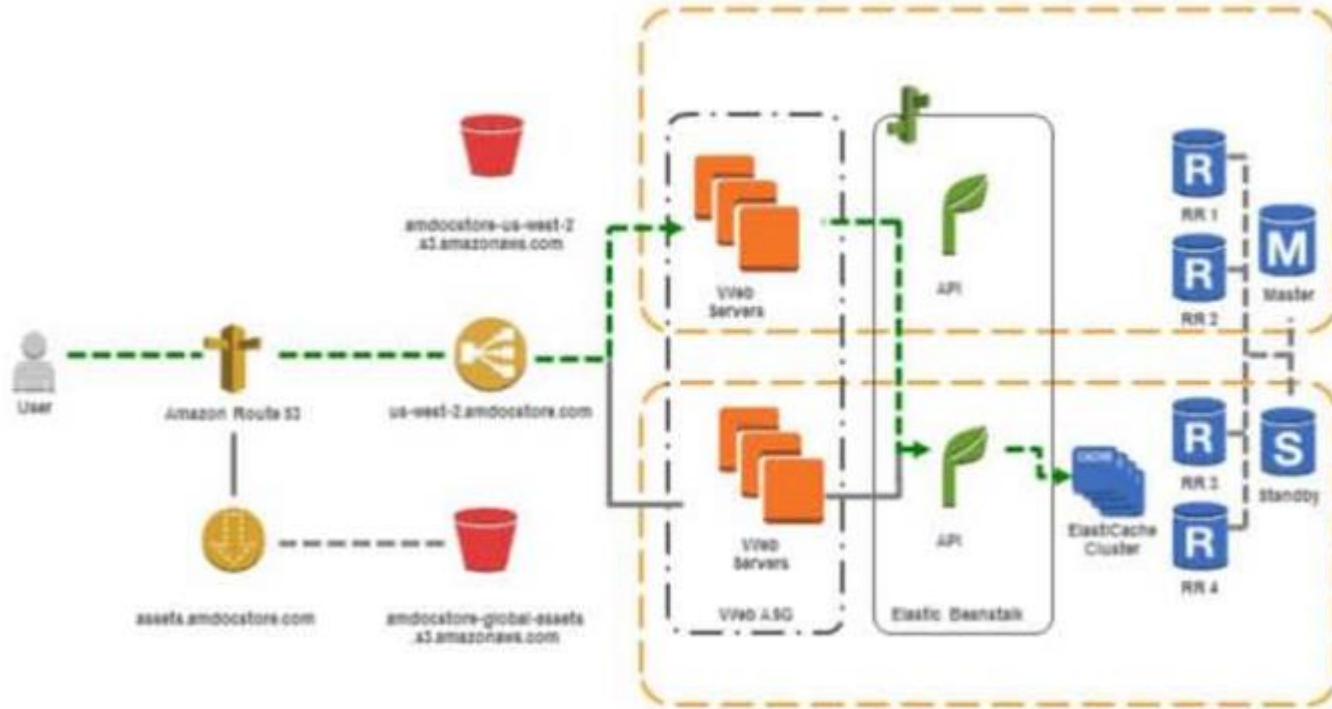
View and Download Docs: ElasticCache Cluster in front of the Read Replicas would provide additional performance when listing the user's documents

# DOCSTORE



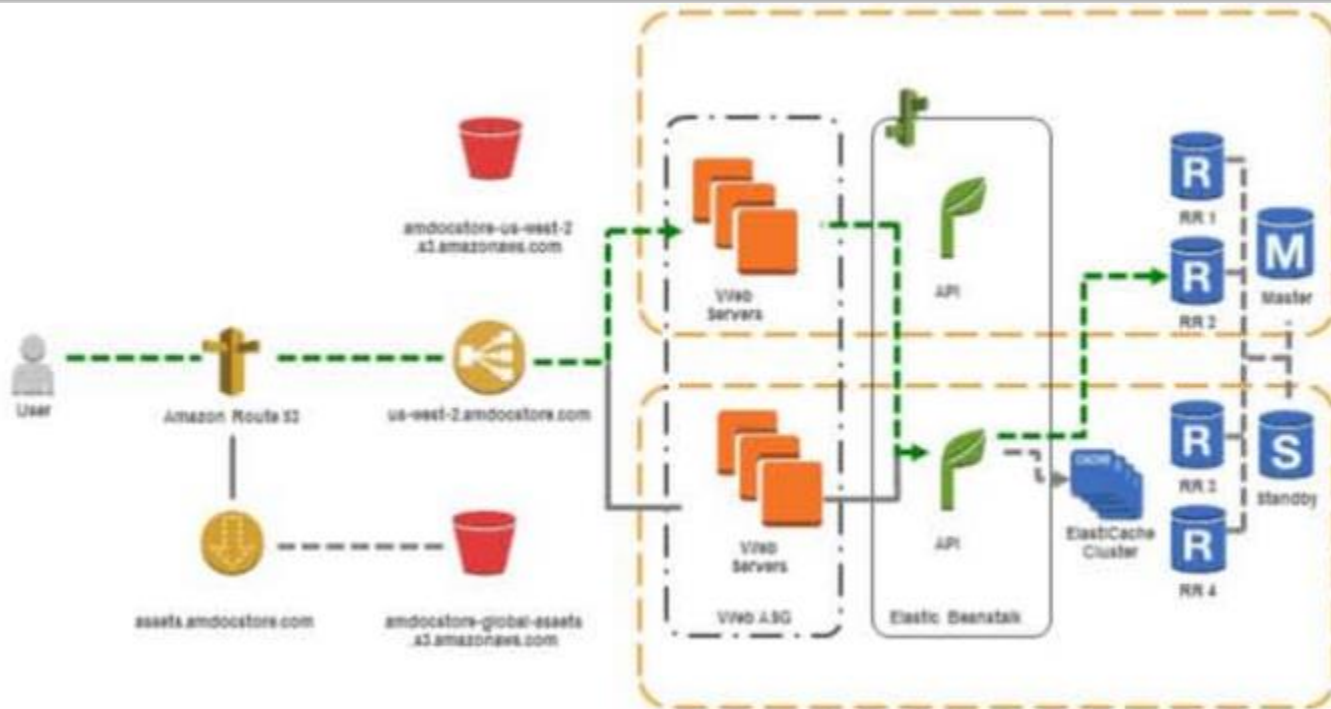
View and Download Docs: The API first checks the cache for a recent list of the user's documents

# DOCSTORE



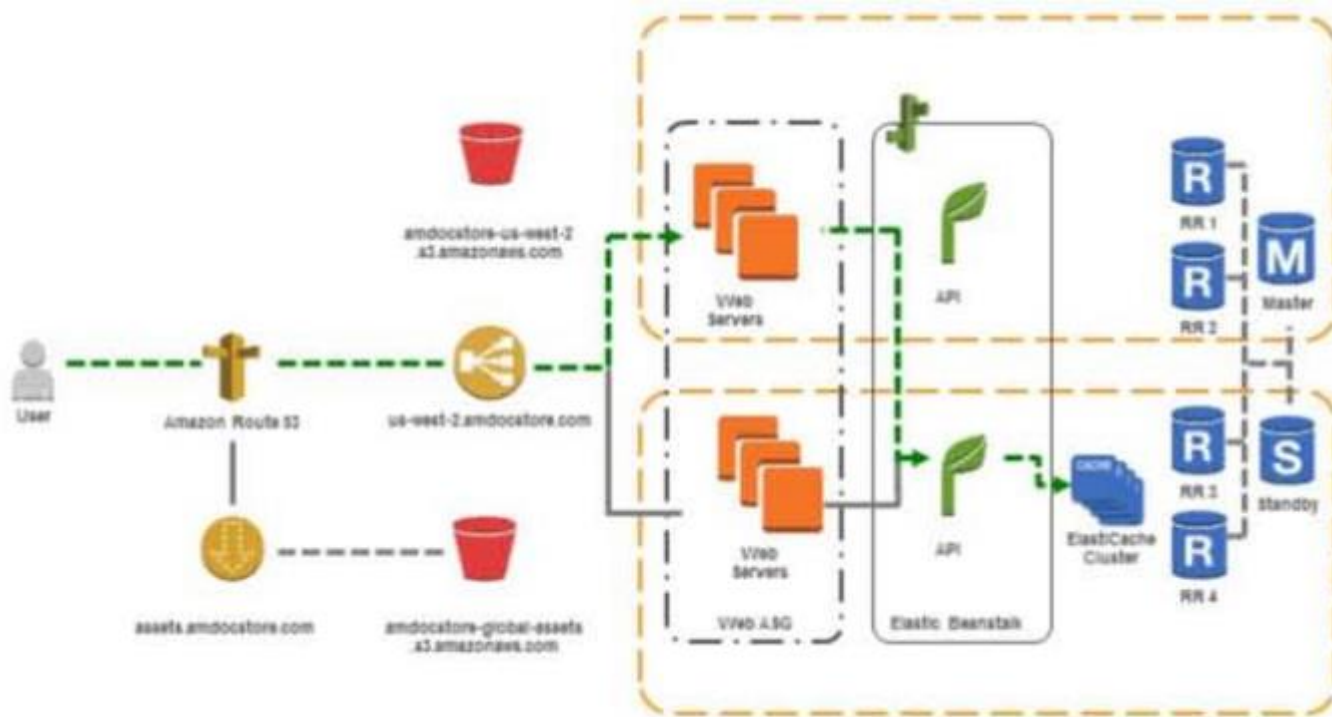
View and Download Docs: If no cached data is available, the API chooses one of the Read Replicas and queries for the list of documents

# DOCSTORE



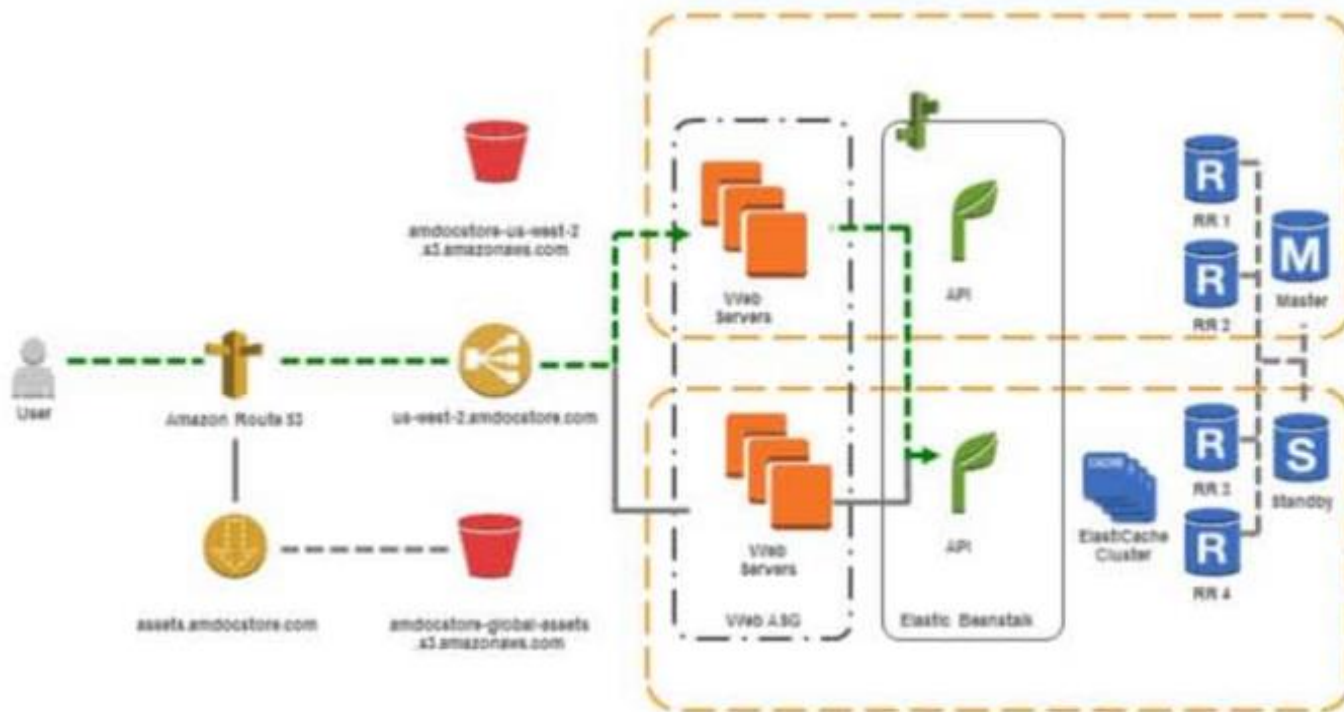
View and Download Docs: The result of the query is stored in the cache for future access

# DOCSTORE



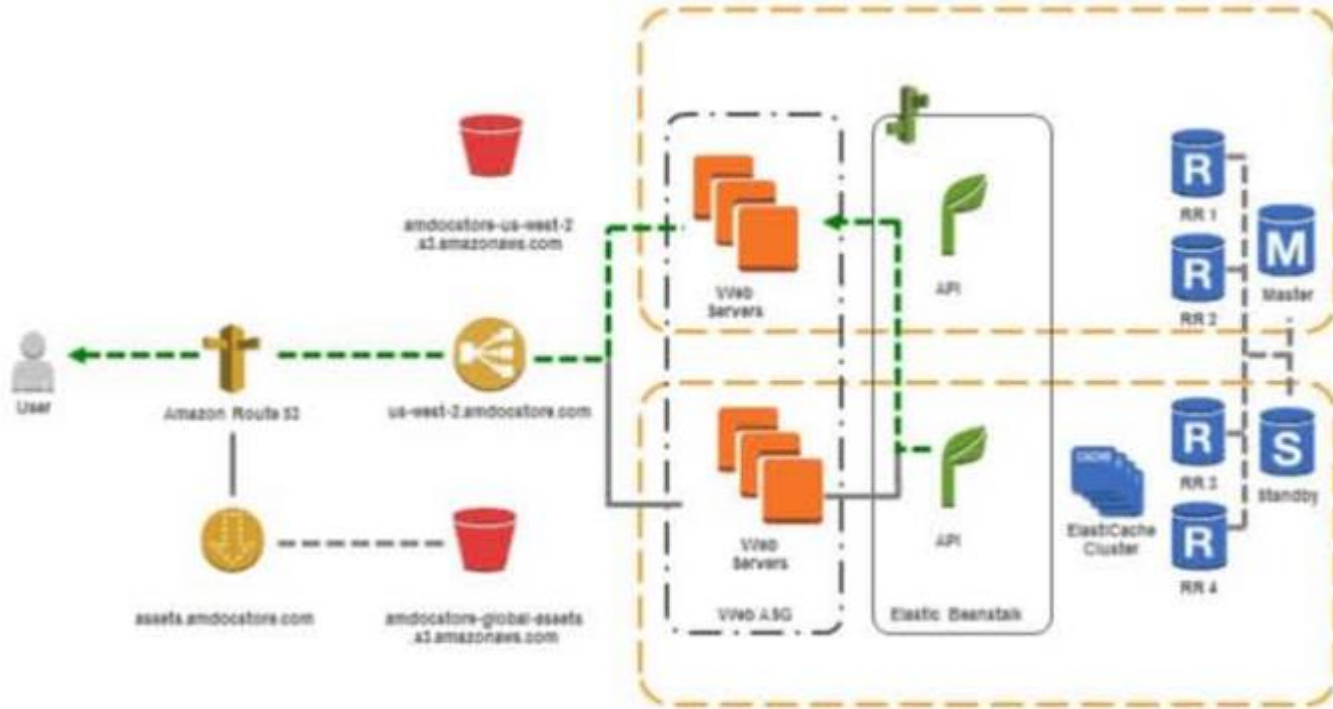
View and Download Docs: For each document, the API constructs a secure, auto-expiring, signed URL that will allow the user to download the document directly from S3

# DOCSTORE



View and Download Docs: The list of documents- and links for direct, secure download-are delivered to the user's browser

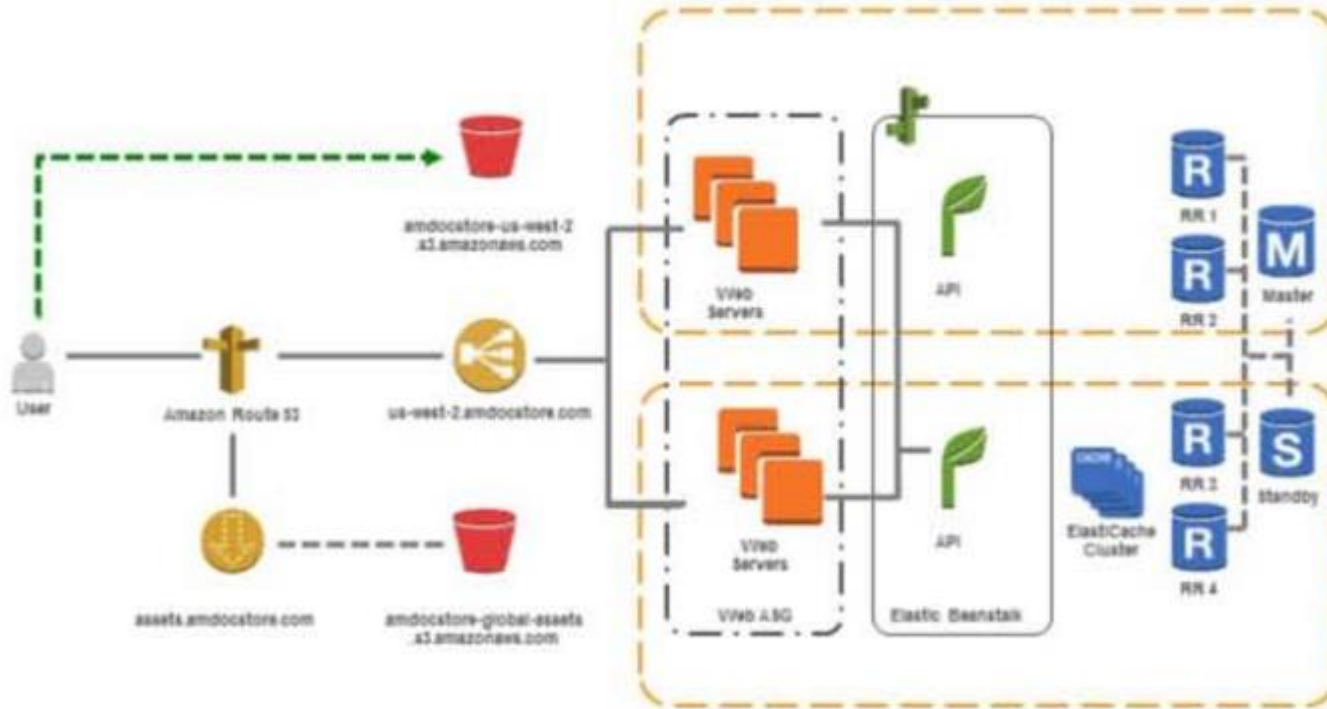
# DOCSTORE





View and Download Docs: The user may download documents directly from S3

# DOCSTORE



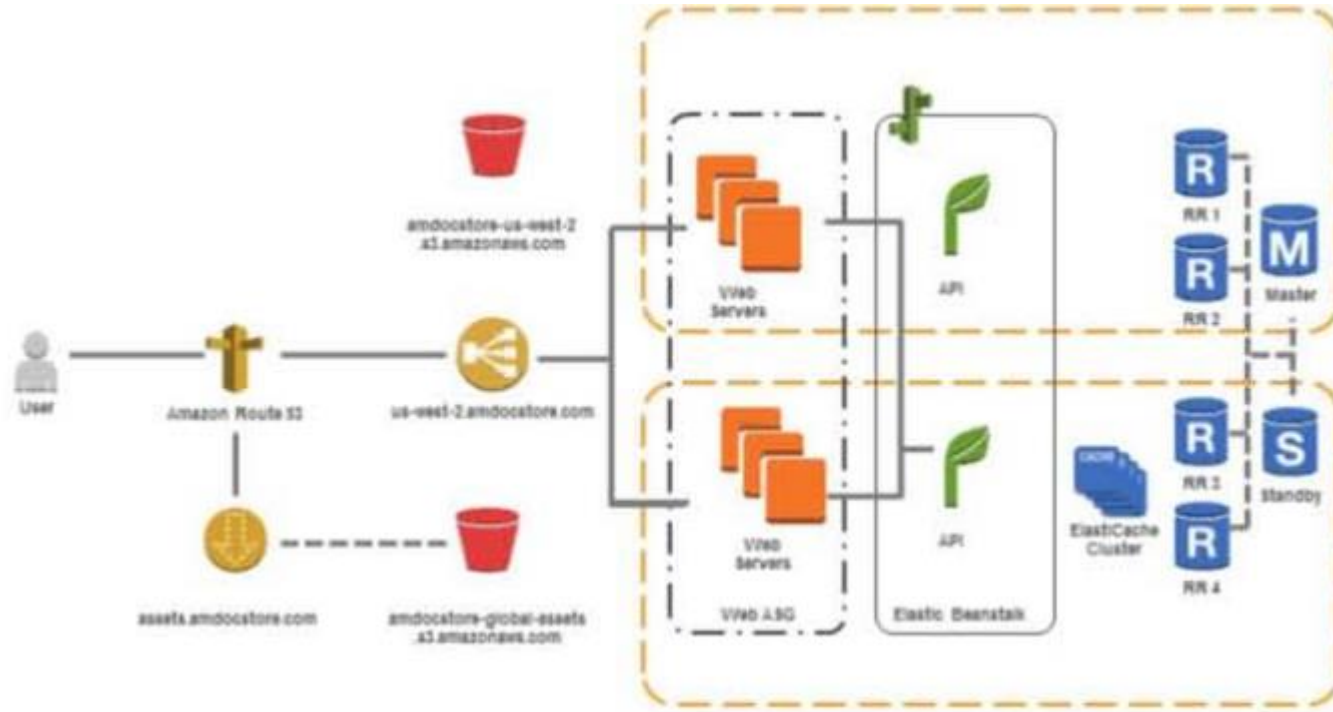


View and Download Docs: Let's look at options for security

DOCSTORE

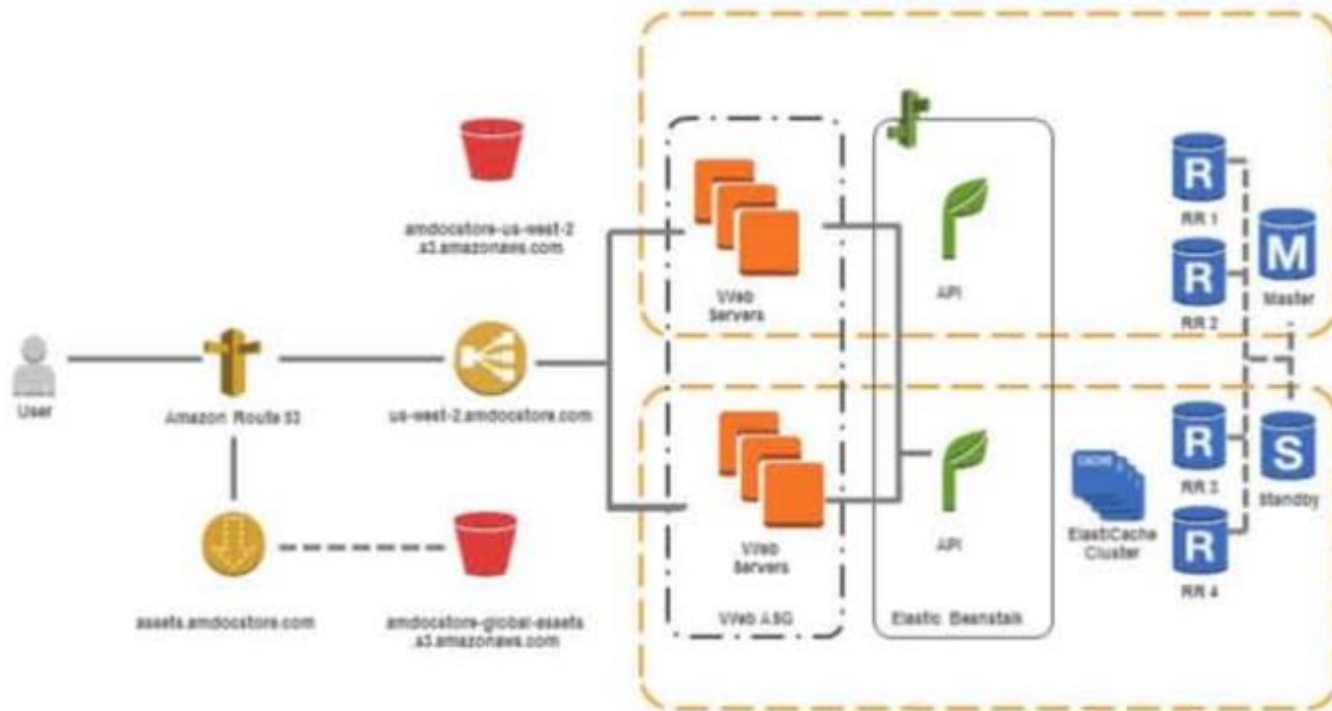
View and Download Docs: Let's look at options for security

DOCSTORE



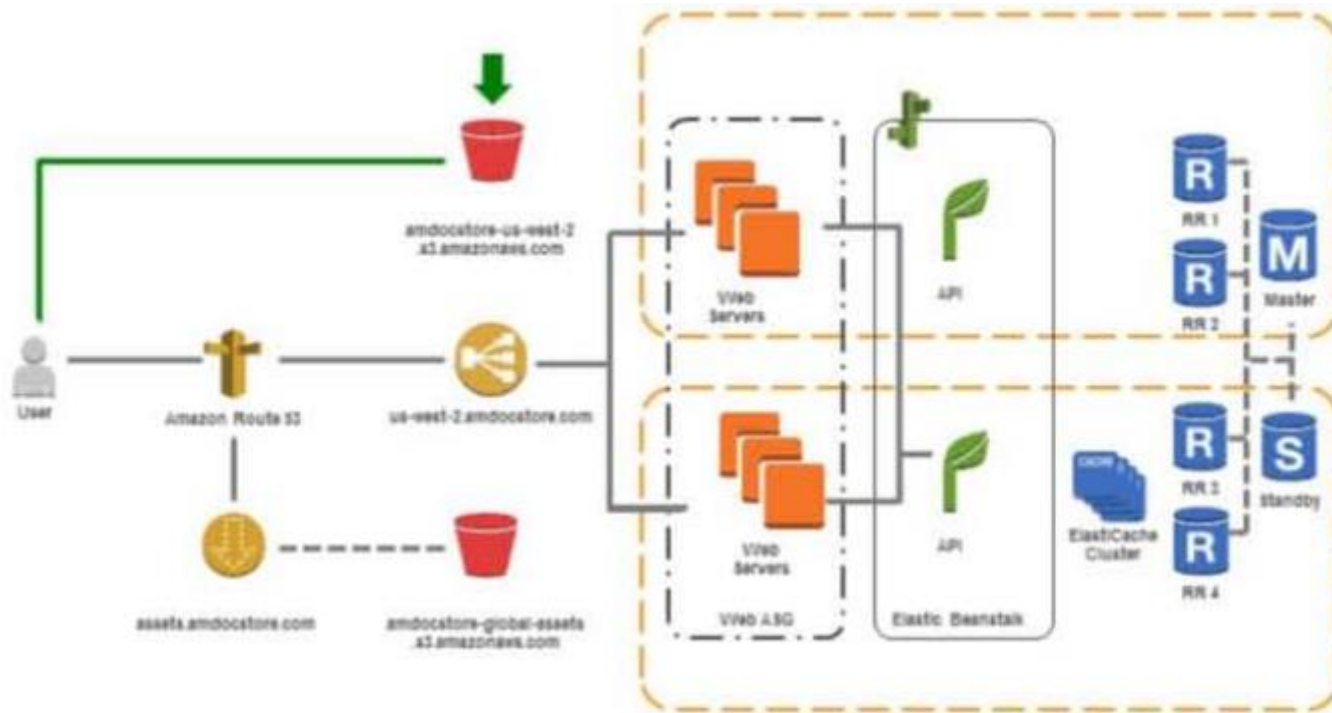
View and Download Docs: Where can security measures and controls be implemented in this architecture?

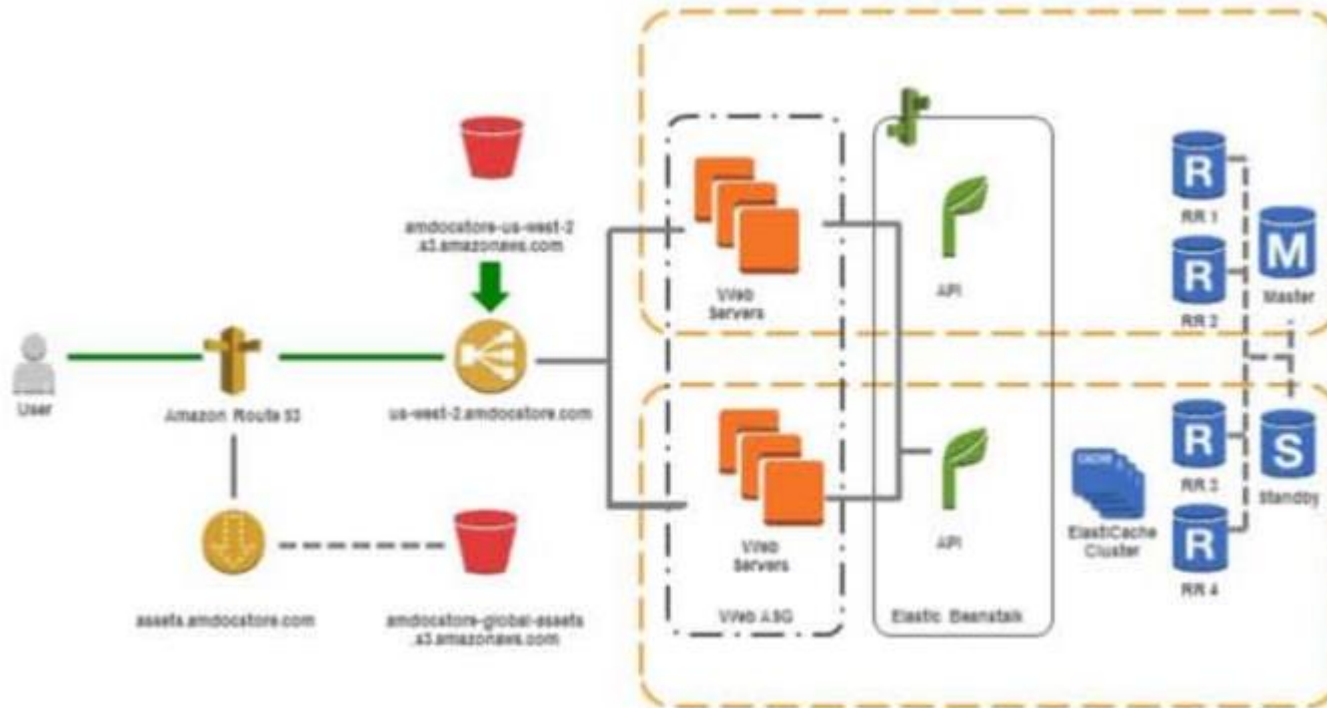
# DOCSTORE

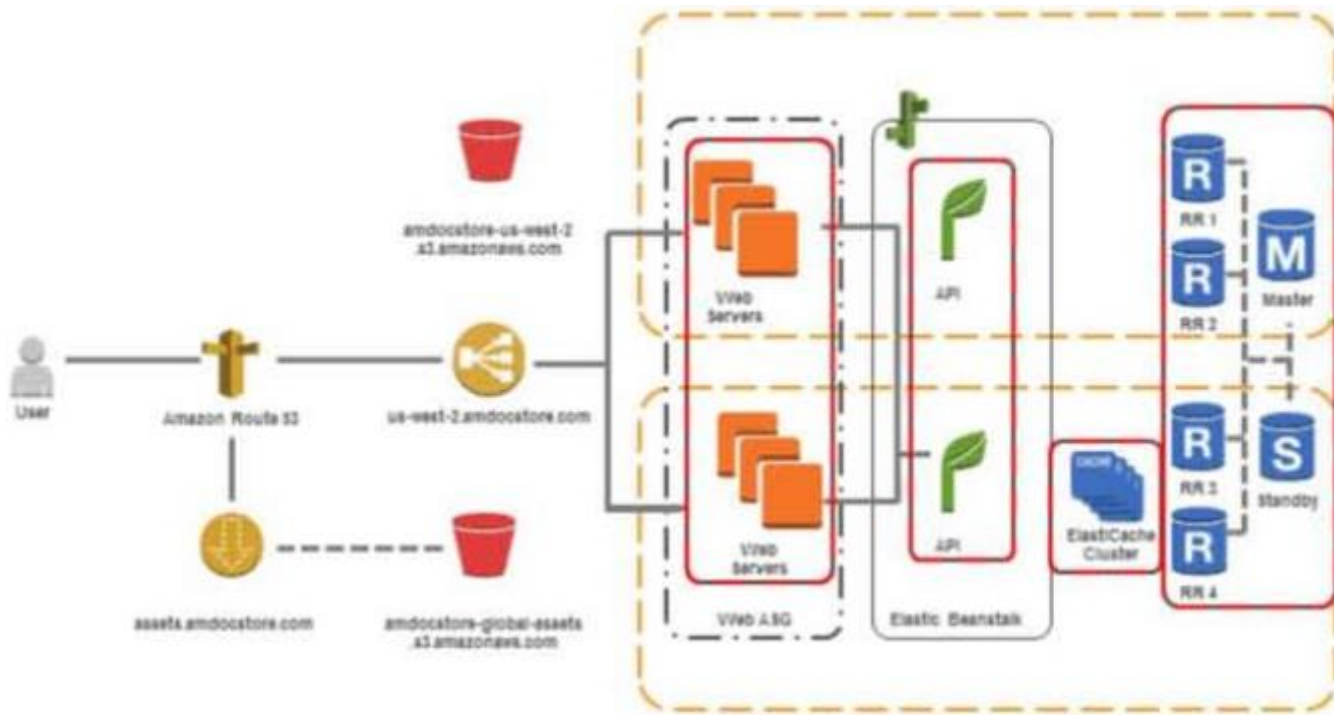


View and Download Docs: User downloads documents directly from S3 using a secure, signed URL and over SSL

# DOCSTORE

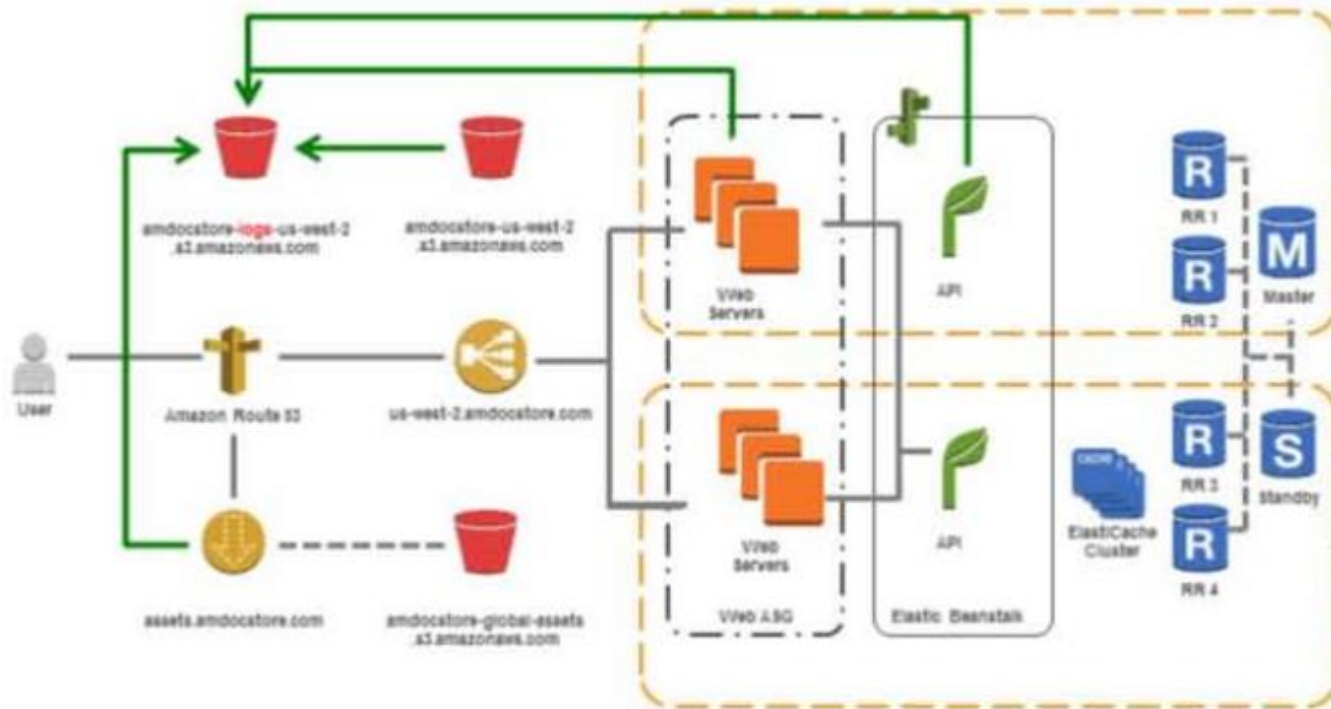






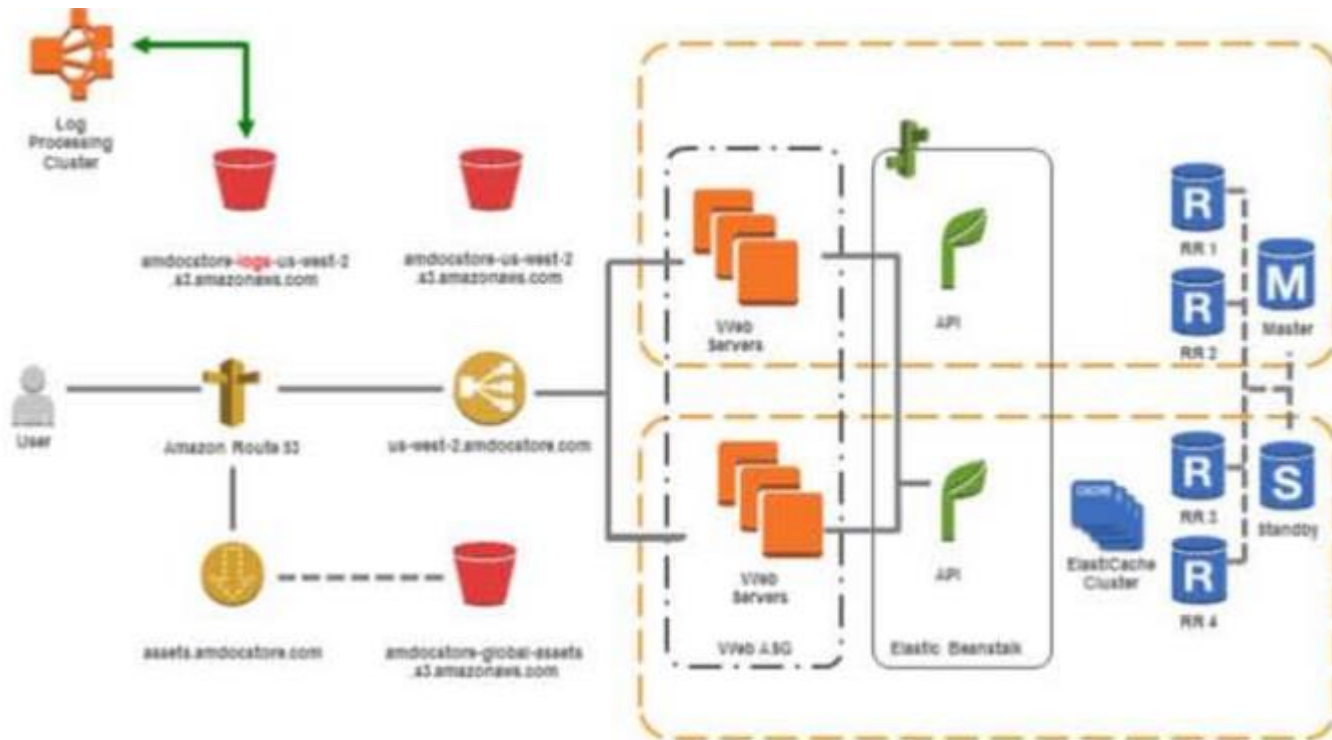
View and Download Docs: Access logs from CloudFront, EC2, and S3 are pushed to S3

# DOCSTORE



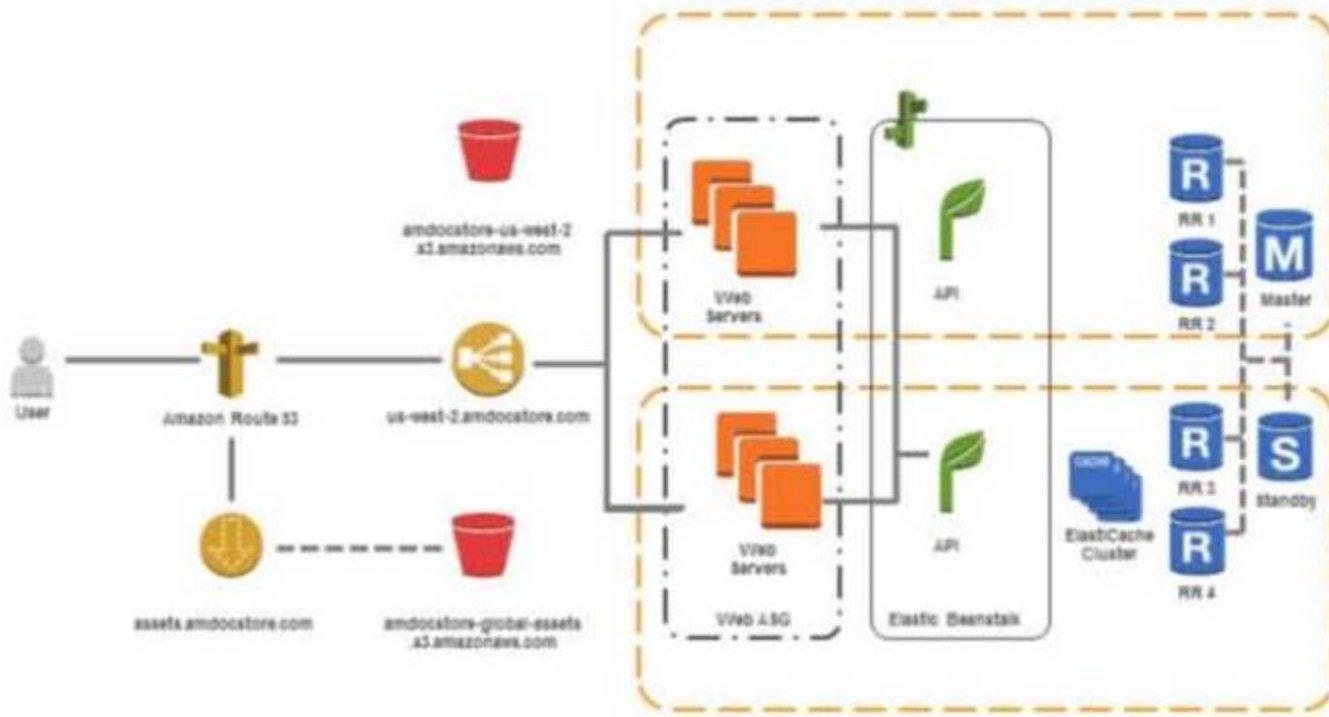
View and Download Docs: Logs are processed on a schedule using Elastic MapReduce cluster

# DOCSTORE



Let's focus on the networking environment

# DOCSTORE





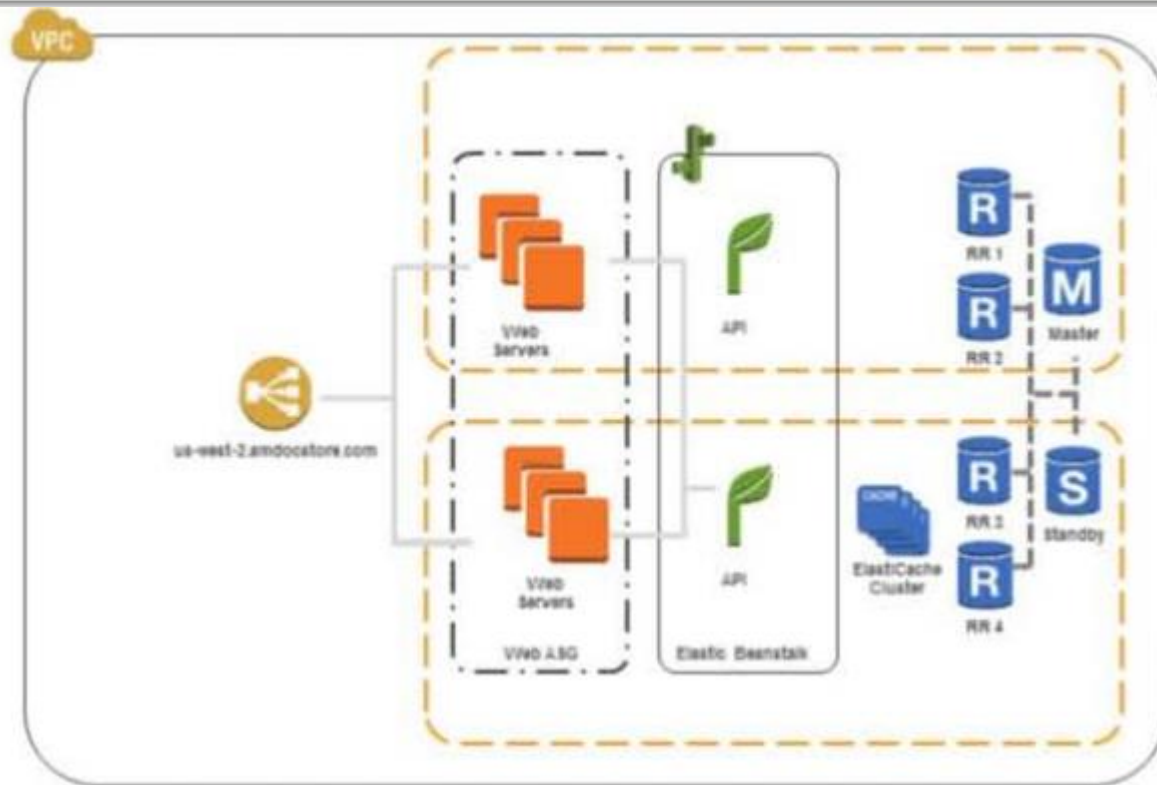
Let's focus on the networking environment

# DOCSTORE



VPC: this VPC is in us-west-2 and has a 10.0.0.0/16 CIDR block

# DOCSTORE

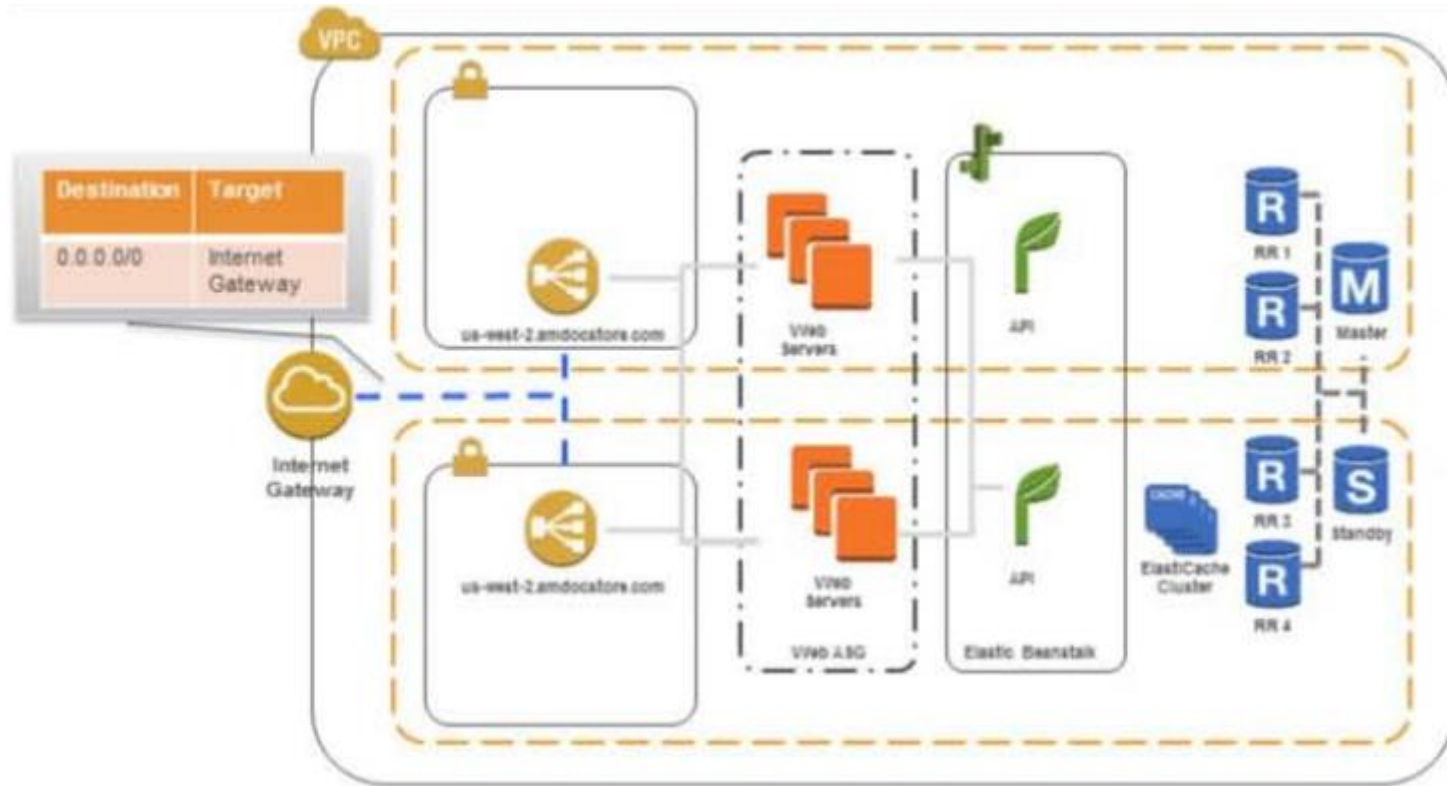


# DOCSTORE



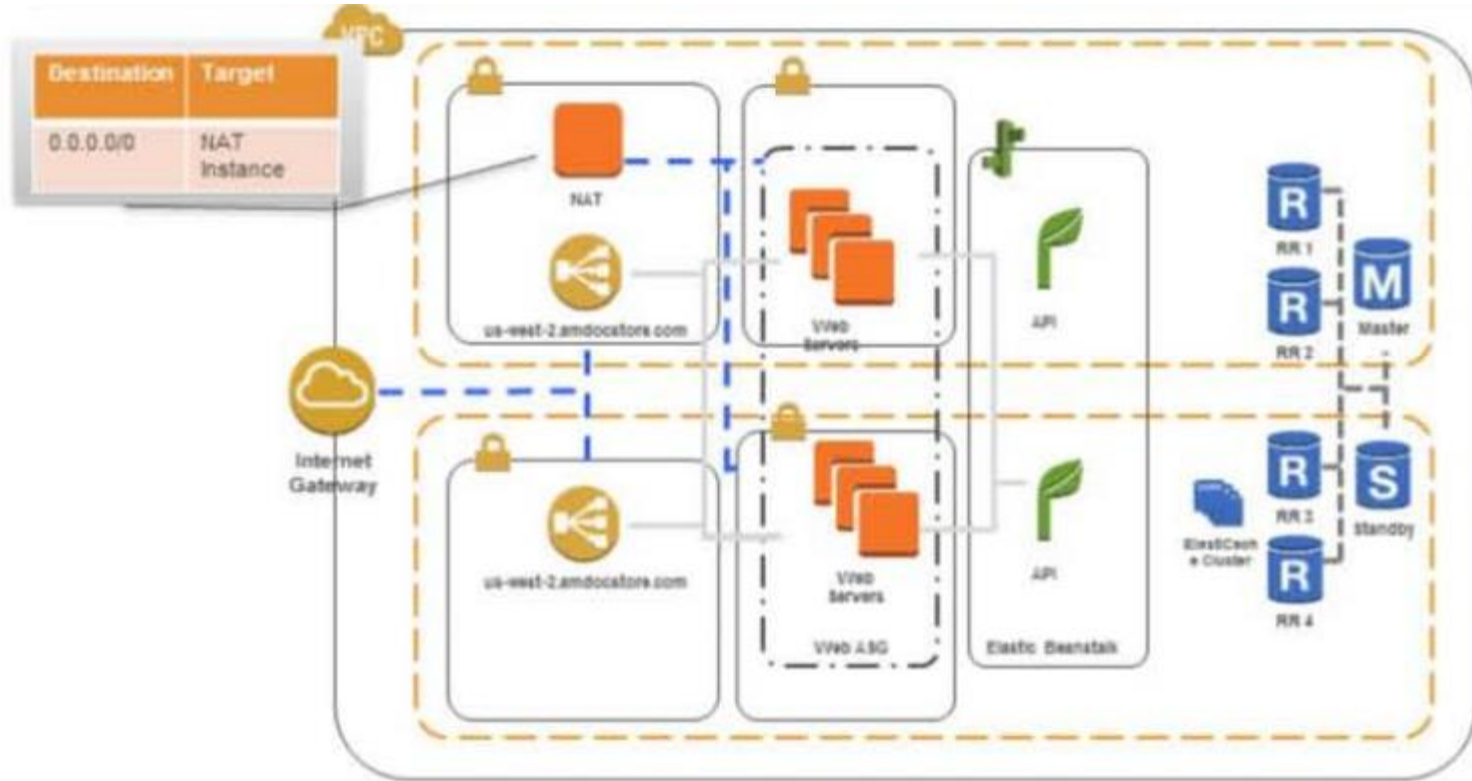
VPC: These subnets should be public. We will attach an Internet Gateway and define a Route Table entry to make them public

# DOCSTORE



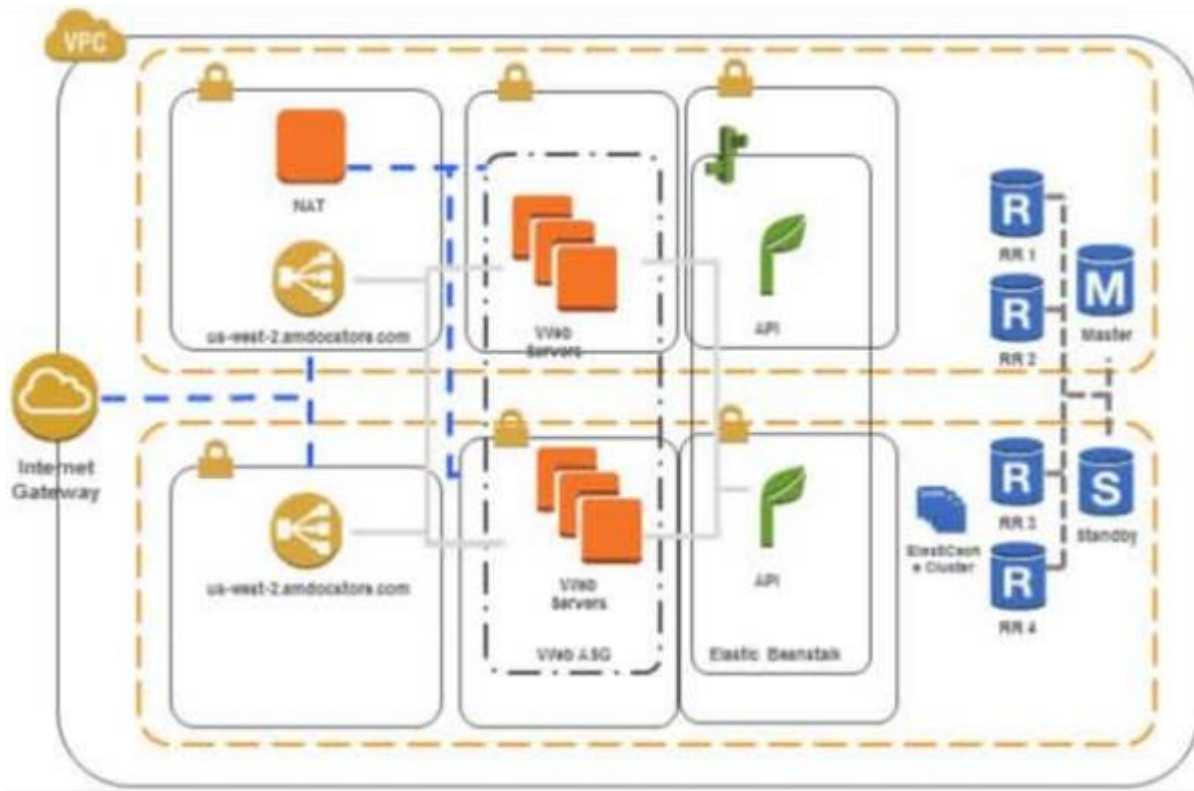
VPC: Web servers run in a private subnet, but have outbound internet access via an EC2 instance in the public subnet running NAT

# DOCSTORE



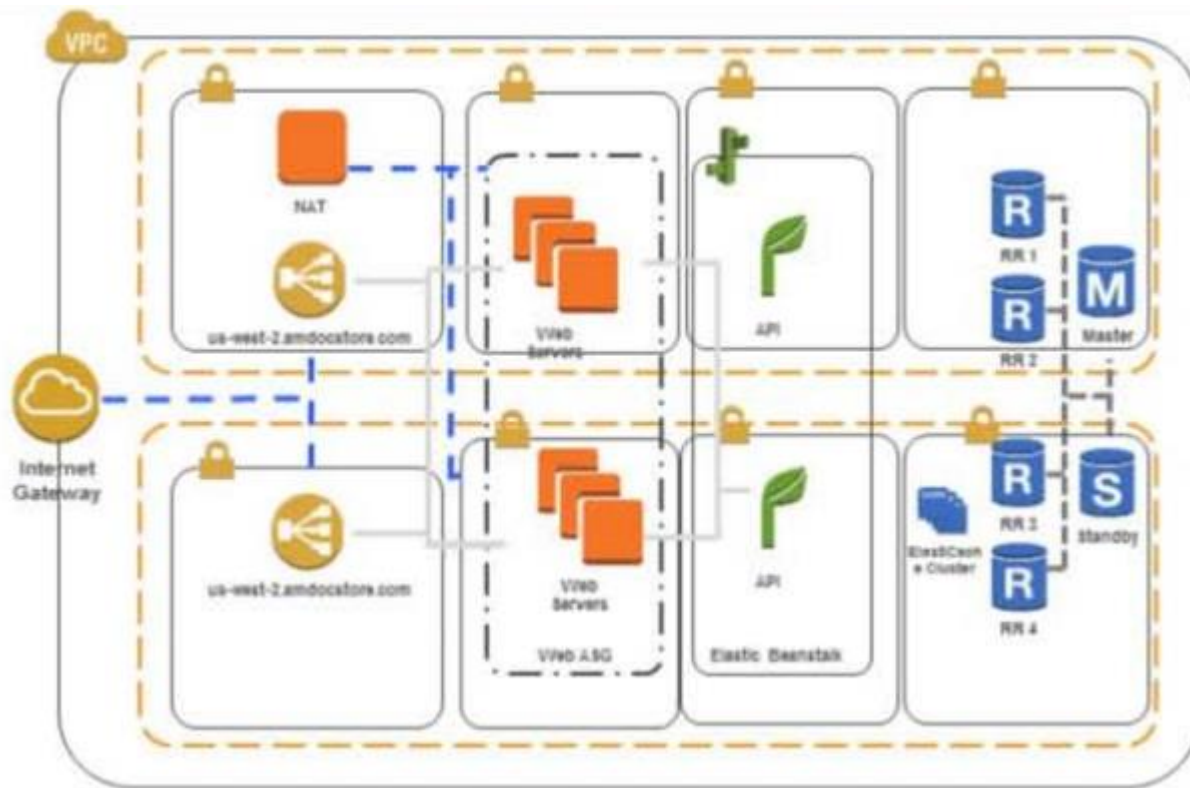
VPC: The DocStore API running on Elastic Beanstalk is deployed in both a public and private subnet

# DOCSTORE



VPC: Finally, the RDS and ElastiCache instances are deployed in a private subnet

# DOCSTORE



CloudFormation: Entire application is deployed as a stack from a template file

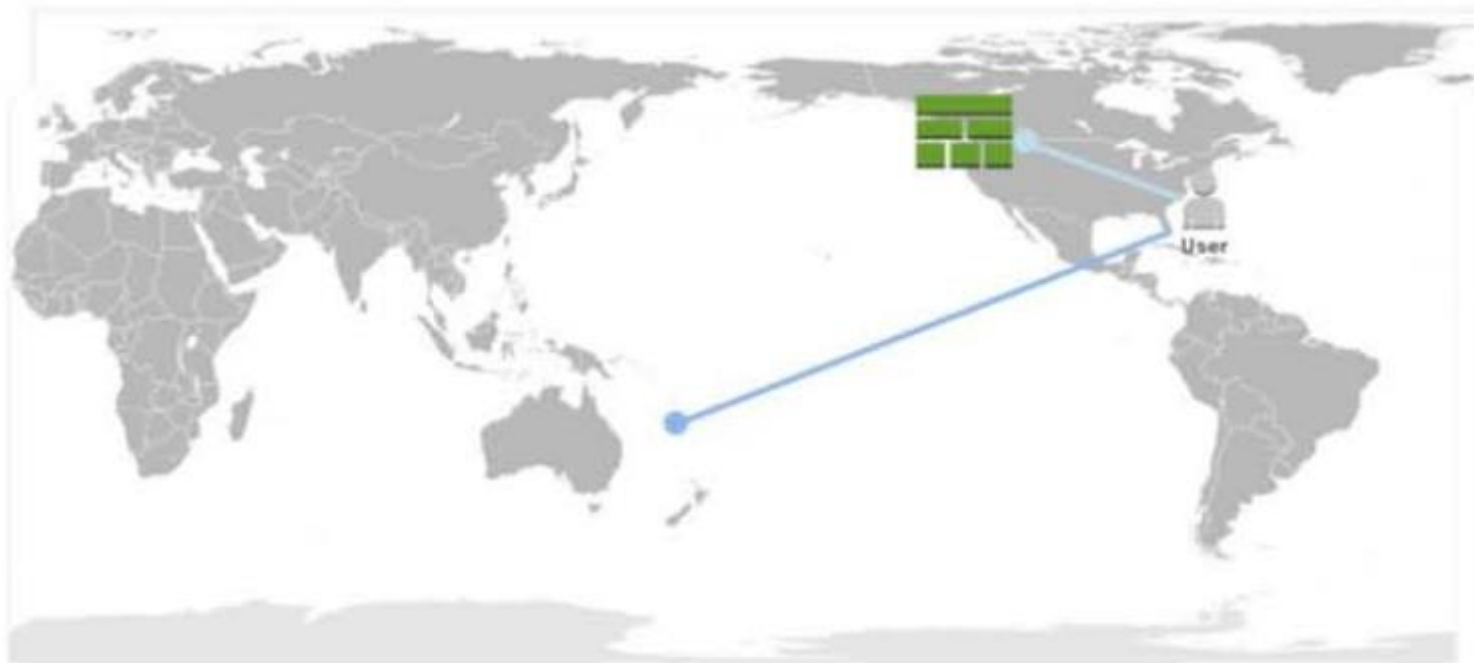
DOCSTORE






CloudFormation: The sample template file is used to deployed to us-west-2

# DOCSTORE





## #2. Batch processing Back-End

## #2. Batch processing Back-End

- Sync user account
- Enforce max storage limits
- Extract and index based on document type
- Detailed dashboard with system-wide totals:
  - # docs, avg, doc size, total storage

## #2. Batch processing Back-End

- Responds to events triggered by Web interface

## #2. Batch processing Back-End

- Responds to events triggered by Web interface



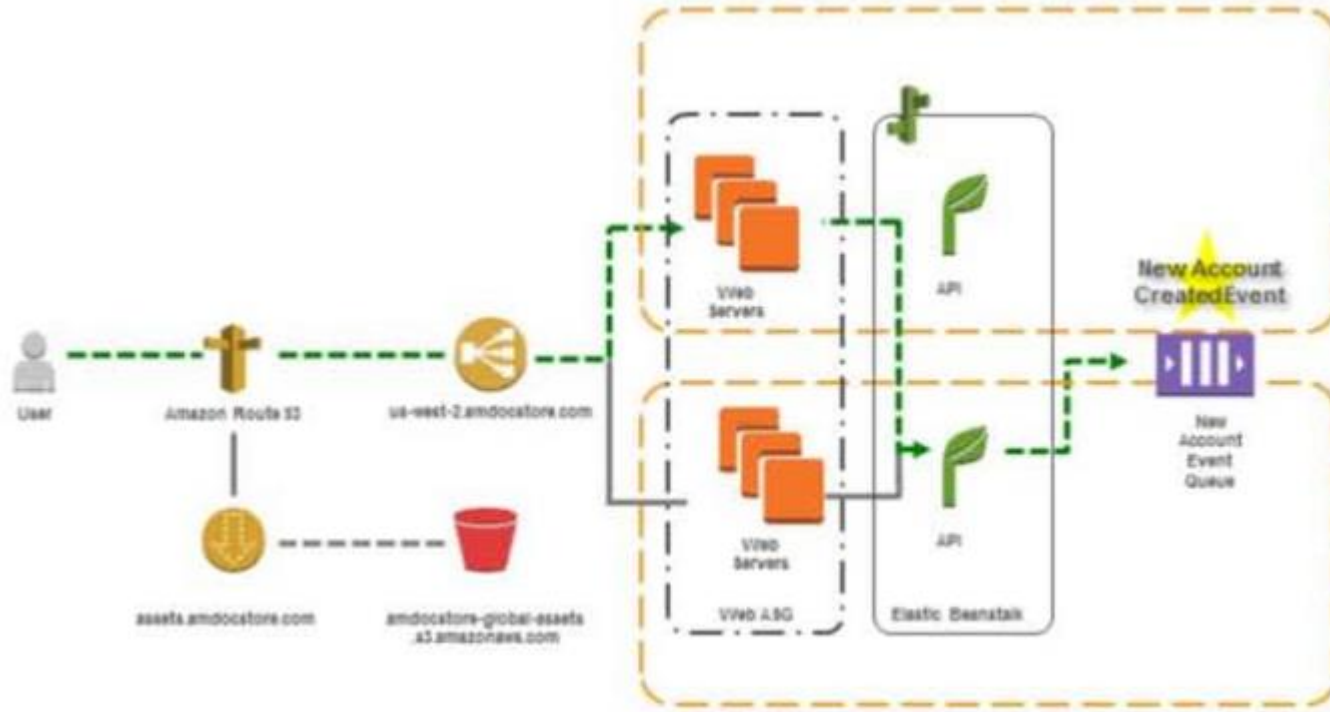
## #2. Batch processing Back-End

- Responds to events triggered by Web and Mobile interfaces



New Account Create Event: Triggered when a user creates a new account. Event message in queue includes username, password and home region


# DOCSTORE






New Account Created Event: Triggered when a user creates a new account, event message in queue includes username, password and home region

DOCSTORE




New Account Created Event: An instance running the Account Synchronization service polls the New Account Event Queue

DOCSTORE



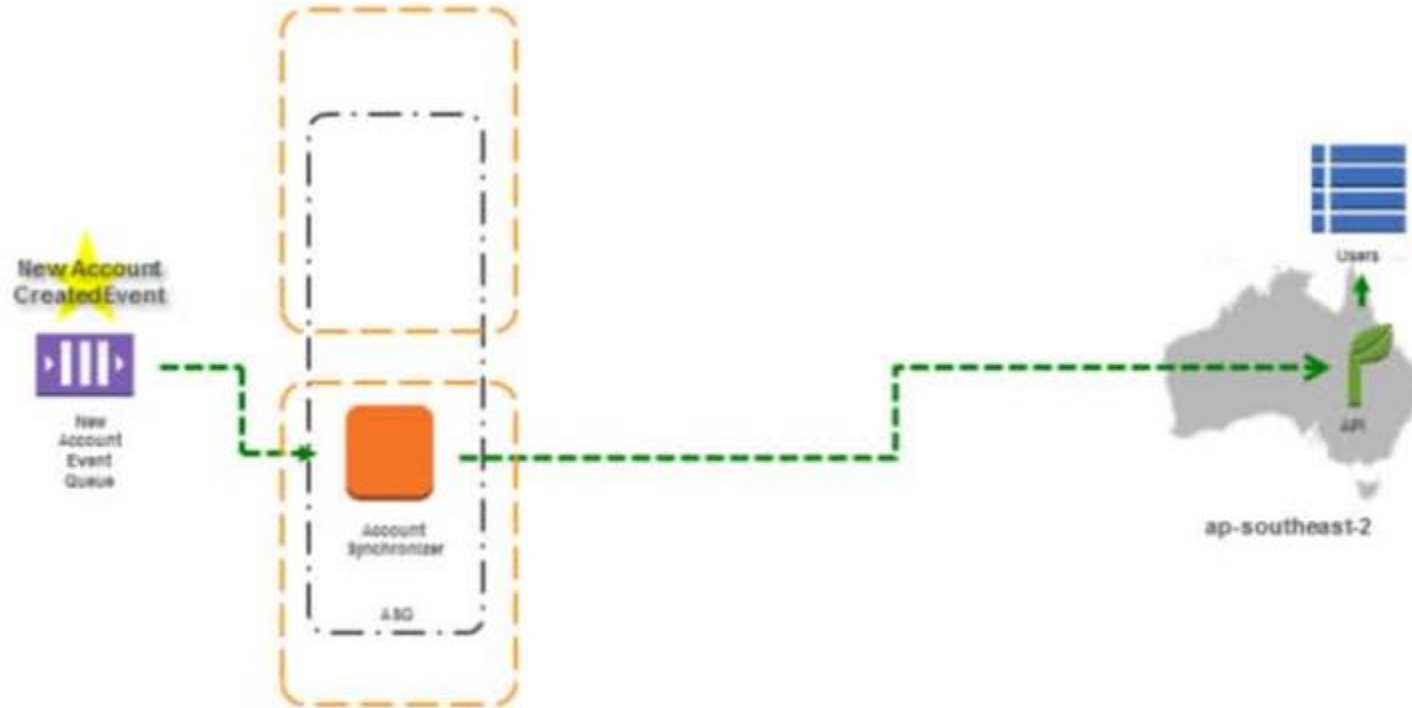
New Account Created Event: If the instance fails, Auto Scaling will launch a replacement

# DOCSTORE



New Account Created Event: Account Synchronizer invokes API in ap-southeast-2 to synchronize new user's account data

DOCSTORE




New Account Created Event : When a user begins uploading a document to S3, an event is triggered so the upload can be tracked via workflow

# DOCSTORE




New Account Created Event : When an event is received from the queue, the workflow initializer kicks off a SWF workflow execution

# DOCSTORE




Doc Uploaded Initiated Event: When a event is received from the queue, the workflow initializer kicks off a SWF workflow execution

# DOCSTORE



Doc Upload Initiated Event: A an instance running Decider code receives the workflow execution started task, and schedules a Check Upload Status activity

# DOCSTORE





Doc Upload Initiated Event: A an instance running the Activity Worker code receives the Check Upload Status task

# DOCSTORE




Doc Uploaded Initiated Event: The Activity Worker periodically checks S3 to determine if the doc upload has completed

# DOCSTORE



Doc Uploaded Initiated Event: Once the document is deleted in S3, a final is scheduled in the near future to ensure that the document has been property processed by the Doc Upload Complete workflow

# DOCSTORE




Doc Uploaded Initiated Event: When a document upload has been complete, it needs to be processed

DOCSTORE




Doc Uploaded Initiated Event: A workflow initializer determines the appropriate workflow to initiate

# DOCSTORE



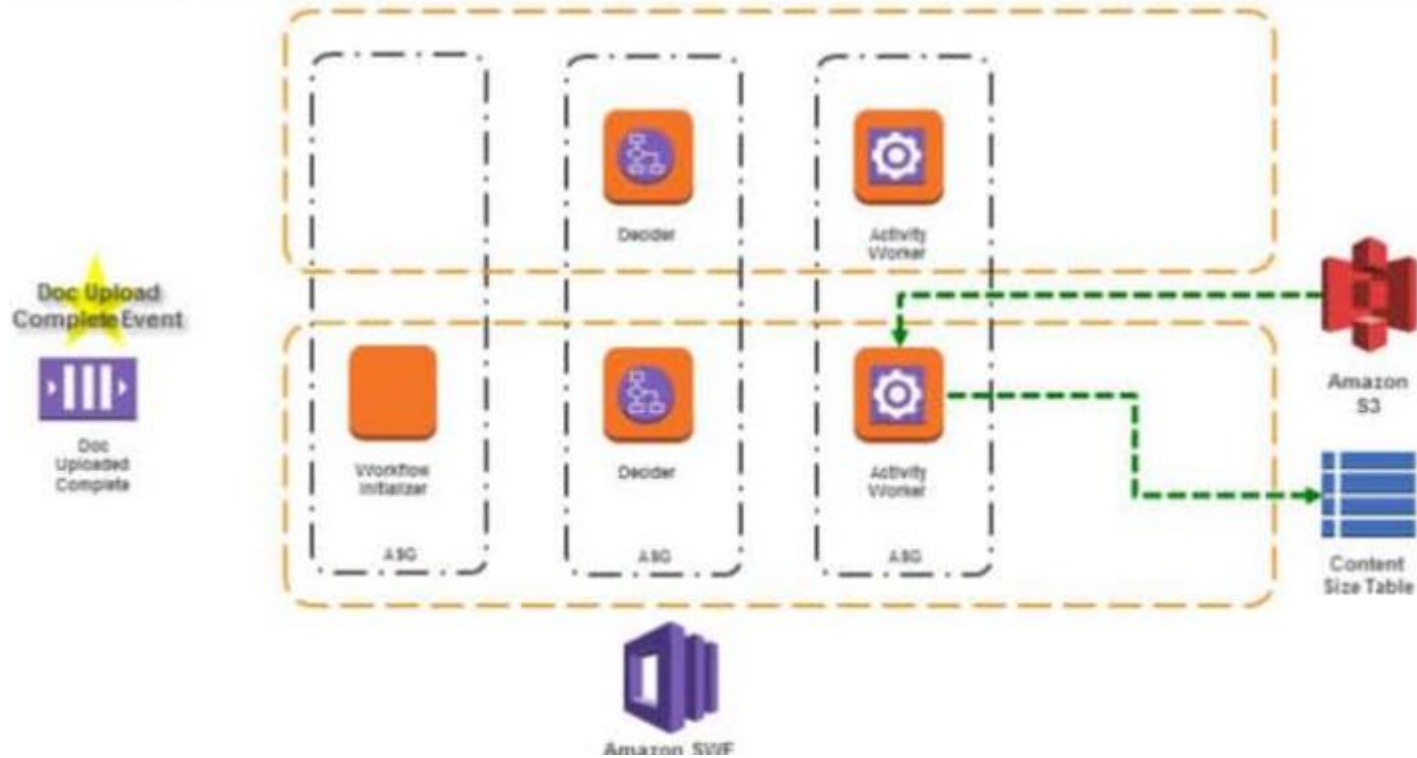
Doc Uploaded Initiated Event: A decider schedules the appropriate activity based on the document type

# DOCSTORE




Doc Uploaded Initiated Event: the content-length metadata property of every object should be retrieved from S3 and used to increment a counter tracking the total size of all uploaded objects

# DOCSTORE



Doc Uploaded Initiated Event: a counter tracking the total content size uploaded by each users is also incremented. If the new doc causes the user to exceed his allocated amount, a flag is set and the user will receive an email via SES


# DOCSTORE





Doc Uploaded Initiated Event: if the document was uploaded by a premium user and it has text context, the full-text content will be indexed in CloudSearch

# DOCSTORE



Doc Uploaded Initiated Event: Finally, several custom metrics are also put to CloudWatch, allowing us to monitor how many new documents are being uploaded by users type, how big the documents are, etc.

# DOCSTORE

